

FACIAL RECOGNITION FROM VIDEO SEQUENCES

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Martyn W Booth

Computing (Industry)

Session (2006/2007)

The candidate confirms that the work submitted is their own and the appropriate credit has been given where reference has been made to the work of others.

I understand that failure to attribute material which is obtained from other sources may be considered as plagiarism.

(Signature of student).....

FACIAL RECOGNITION FROM VIDEO SEQUENCES

ABSTRACT

Face recognition has proven to be a difficult process to automate, even though humans are able to detect and identify faces with little or no effort. As a result of this, the research into automated face recognition has remained a focal point in the field of computer vision and artificial intelligence for over twenty five years. There have been many different attempts at solving the problem of automatic facial recognition (from static images and from video). In order to address this, a number of recent developments in the field of computer vision are examined. Some of these techniques are then applied to the design, implementation and test of a prototype facial recognition system utilising a live video feed. Furthermore, the system will use a number of video frames in order to attempt recognition. This is an attempt to improve recognition rates.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

ACKNOWLEDEMENT

I would like to thank Professor David Hogg for his help, assistance and patience in helping me with some of the more complex areas of this project. Without his help, this report would be in an extremely poor state.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

TABLE OF EQUATIONS	X
CHAPTER 1 – REPORT INTRODUCTION	1
1.1 INTRODUCTION.....	1
1.2 PROBLEM DEFINITION	1
1.3 STUDY SCOPE.....	3
1.4 STUDY OBJECTIVES	3
1.5 PROJECT PLAN.....	4
1.5.1 <i>Project Plan Revisions (Shown in italics above Table 1.5-1)</i>	4
CHAPTER 2 – FACIAL RECOGNITION	5
2.1 ABOUT FACIAL RECOGNITION.....	5
2.1.1 <i>Difficulties with current face recognition systems</i>	6
2.2 HUMAN INTERPRETATION OF FACES.....	8
2.3 CONCLUSION	9
CHAPTER 3 – LITERATURE SURVEY ON FACIAL RECOGNITION TECHNIQUES.....	10
3.1 INTRODUCTION.....	10
3.2 INTRODUCTION TO AUTOMATED FACE RECOGNITION SYSTEMS	10
3.3 FEATURE EXTRACTION	11
3.4 FEATURE MATCHING	13
3.5 IMAGE NORMALISATION	14
3.6 FACE DETECTION TECHNIQUES.....	14
3.6.1 <i>Viola and Jones Face Detector</i>	15
3.6.2 <i>Sparse Network of Winnows (SNoW) Face Detection</i>	17
3.7 FEATURE EXTRACTION TECHNIQUES.....	18
3.7.1 <i>Bayesian approach to feature localisation</i>	19
3.7.2 <i>Two-dimensional PCA-based feature localisation</i>	20
3.7.3 <i>Edge-based shape detection</i>	22
3.7.4 <i>Kernel Eigenfaces and Fisherfaces approach comparison</i>	23
3.7.5 <i>Support Vector Machines (SVMs)</i>	26
3.8 CONCLUSION	27

FACIAL RECOGNITION FROM VIDEO SEQUENCES

CHAPTER 4 – SYSTEM DESIGN.....	27
4.1 INTRODUCTION.....	27
4.2 SYSTEM REQUIREMENTS & SPECIFICATION	27
4.2.1 <i>System Requirements.....</i>	<i>28</i>
4.2.2 <i>User Requirements.....</i>	<i>28</i>
4.2.3 <i>System Specification.....</i>	<i>28</i>
4.3 FACE DETECTOR SPECIFICATION	29
4.4 FEATURE LOCALISATION SPECIFICATION	29
4.4.1 <i>Data Training Specification.....</i>	<i>29</i>
4.4.2 <i>Data Testing Specification.....</i>	<i>30</i>
4.5 DESCRIPTION OF DESIGN METHODOLOGIES	30
4.5.1 <i>Description of Information Flow.....</i>	<i>31</i>
4.5.2 <i>Description of State Transition.....</i>	<i>32</i>
4.5.3 <i>Object Orientated Analysis (OOA).....</i>	<i>32</i>
4.5.4 <i>Justification of class choices.....</i>	<i>33</i>
4.5.5 <i>Explanation of message passing (Figure 0-1).....</i>	<i>34</i>
4.5.6 <i>Graphical User Interface</i>	<i>34</i>
CHAPTER 5 – SYSTEM IMPLEMENTATION	34
5.1 INTRODUCTION.....	35
5.2 SYSTEM IMPLEMENTATION	35
5.2.1 <i>Feature Detection Training (Eye).....</i>	<i>35</i>
5.2.2 <i>Data Storage/Manipulation Classes</i>	<i>36</i>
5.2.3 <i>Database Manipulation Classes.....</i>	<i>37</i>
5.2.4 <i>Video handling class.....</i>	<i>37</i>
5.2.5 <i>Face Detection class.....</i>	<i>37</i>
5.2.6 <i>Overview of main program functionality</i>	<i>38</i>
5.2.7 <i>Individual frame decoding.....</i>	<i>40</i>
5.2.8 <i>Frame handling functionality</i>	<i>41</i>
CHAPTER 6 – SYSTEM EVALUATION	41
6.1 INTRODUCTION.....	41
6.2 SYSTEM TESTING.....	42
6.2.1 <i>System Black Box/ Functionality Testing.....</i>	<i>42</i>
6.2.2 <i>Training system testing</i>	<i>42</i>
6.2.3 <i>System Accuracy Testing</i>	<i>43</i>

FACIAL RECOGNITION FROM VIDEO SEQUENCES

6.2.4	System Requirements Testing	43
6.2.5	Test Summary.....	43
6.3	SYSTEM EVALUATION.....	44
6.3.1	Adopted approach.....	45
6.3.2	Outcome	45
CHAPTER 7 – STUDY CONCLUSION AND AREAS OF FURTHER INVESTIGATION		48
7.1	INTRODUCTION.....	48
7.2	STUDY CONCLUSIONS.....	48
7.3	FURTHER AREAS OF INVESTIGATION	49
CHAPTER 8 – APPENDICES		50
8.1	APPENDIX A – DESIGN DOCUMENTATION	50
8.1.1	High Level System Data Flow	50
8.1.2	Level 1.3 High Level Face Detection flow	50
8.1.3	Level 1 feature training flow	51
8.1.4	Level 1.1 feature training flow	51
8.1.5	Level 1.1.1 detailed feature training flow	52
8.1.6	Level 1.2 feature extraction testing (comparison) flow	53
8.1.7	Level 1.2.1 feature matching scheme flow.....	53
8.1.8	Level 1.2.1.1 low level Euclidean distance comparison flow.....	54
8.1.9	State Transition Diagram	54
8.1.10	Object Orientated Analysis (Class) Diagram	55
8.1.11	Threads of execution Diagram	56
8.1.12	High Level Implementation Diagram	57
8.1.13	Final class diagram.....	58
8.1.14	Simple Single-table Database Schema	59
8.2	APPENDIX B – INTERFACE WINDOWS.....	59
8.2.1	Main Interface Window	59
8.2.2	Add Reference Image to database window	60
8.2.3	Remove Reference Image from database window	60
8.3	APPENDIX C – CODE LISTING.....	60
8.3.1	Training Eye Feature Detector (MATLAB)	60
8.3.2	Face.java class source code.....	62
8.3.3	FaceCrop.java class source code	64
8.3.4	EyePatch.java class source code	65
8.3.5	Matrix.java class source code	66

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.6	<i>OpenFileAction.java class source code</i>	68
8.3.7	<i>FaceLocation.java class source code</i>	69
8.3.8	<i>FaceLocations.java class source code</i>	70
8.3.9	<i>FaceDataAccessLayer.java class source code</i>	71
8.3.10	<i>FaceAddToDatabase.java class source code</i>	73
8.3.11	<i>FaceRemoveFromDatabase.java class source code</i>	77
8.3.12	<i>FrameAccess.java class source code</i>	79
8.3.13	<i>FaceDetection.java class source code</i>	86
8.3.14	<i>FeatureLocalisation.java original class source code</i>	88
8.3.15	<i>FrameHandler.java class source code</i>	95
8.3.16	<i>WorkQueue.java class</i>	101
8.3.17	<i>FaceRecogGUI.java class (Main program entry point)</i>	102
8.3.18	<i>Java test classes:</i>	106
8.4	APPENDIX D – SYSTEM TESTING	113
8.4.1	<i>System Testing Results</i>	113
8.4.2	<i>System Requirements Testing Results</i>	116
8.4.3	<i>Evaluation Framework</i>	117
8.5	APPENDIX E – PERSONAL REFLECTION	118
8.6	BIBLIOGRAPHY	120

FACIAL RECOGNITION FROM VIDEO SEQUENCES

TABLE OF FIGURES

Figure 2.1-1 – Face detection using Viola & Jones face detector (Bornet, 2005).....	6
Figure 2.1-2 - Changes in facial expressions (Dyer, 2004).....	7
Figure 2.2-1 - Representation of the child face space. Inferred from.....	9
Figure 2.3-1 - False acceptance, false rejection rate example (Efford, 2006).....	10
Figure 3.3-1 – Average eye image (acting as template).....	12
Figure 3.3-2 – Eyes localised on test image (Zhao & Grigat, 2006).....	12
Figure 3.3-3 - Test patch	13
Figure 3.3-4 - Sliding Window test image	13
Figure 3.3-5 - Results of blob detection	13
Figure 3.6-1 - The integral image example	16
Figure 3.6-2 - Weak classifiers used by Face Detector	17
Figure 3.7-1- Results of errors of classification and regression eye localisation	19
Figure 3.7-2 - Principal Component Analysis example	20
Figure 3.7-3 - Comparison of no. of PCs in PCA & 2DPCA (Yang, Zhang, Frangi, & Yang, 2004)	21
Figure 3.7-4 - Feature detection using shape matching.....	23
Figure 3.7-5 - Decision Plane example	26
Figure 3.7-6 - Hyperplane example.....	26
Figure 3.7-7 - Hyperplane mapped (using kernels) to Decision plane.....	26
Figure 5.2-1- Training Data calculations pseudocode.....	35
Figure 5.2-2 - Pseudocode for sliding window approach.....	39
Figure 5.2-3 - Sum-of-squares difference java source code.....	40
Figure 8.1-1 - High Level Data Flow Diagram (Whole System)	50
Figure 8.1-2 - High Level Face Recognition System.....	50
Figure 8.1-3 - Level 1 System Data Flow Diagram	51
Figure 8.1-4 - Level 1.1 Feature Training Flow.....	51
Figure 8.1-5 - Level 1.1.1 Low Level Feature Training.....	52
Figure 8.1-6 - Level 1.2 Feature Testing Data Flow	53
Figure 8.1-7 - Level 1.2.1 Feature Matching Scheme Data Flow	53
Figure 8.1-8 - Level 1.2.1.1 Low Level Matching Data Flow	54
Figure 8.1-9 - System State Transition Diagram.....	54
Figure 0-1- Class Diagram	55
Figure 0-2 - Thread of execution (main program).....	56
Figure 0-3 - HLD Diagram.....	57
Figure 0-1 - Final Class Diagram	58
Figure 0-2 - System Database Schema	59
Figure 8.2-1 - Main Interface GUI display.....	59
Figure 8.2-2 - Add new reference image to database GUI.....	60
Figure 8.2-3 - Remove existing image from reference database GUI.....	60
Figure 8.3-1- getMeanEyeRight(Left).m Training Module	61
Figure 8.3-2 - Face.java class.....	63
Figure 8.3-3 - FaceCrop.java class.....	64

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Figure 8.3-4 - EyePatch.java class	65
Figure 8.3-5 - Matrix.java class.....	67
Figure 8.3-6 - OpenFileAction.java class.....	68
Figure 8.3-7 - FaceLocation.java class.....	69
Figure 8.3-8 - FaceLocations.java class	70
Figure 8.3-9 - FaceDataAccessLayer.java class.....	72
Figure 8.3-10 - FaceAddToDatabase.java class.....	76
Figure 8.3-11 - FaceRemoveFromDatabase.java class	78
Figure 8.3-12 - FrameAccess.java class	85
Figure 8.3-13 - FaceDetection.java class	87
Figure 8.3-14 - FeatureLocalisation.java Source code (original).....	94
Figure 8.3-15 - FrameHandler.java source.....	100
Figure 8.3-16 - WorkQueue.java source	101
Figure 8.3-17 - FaceRecogGUI.java source code, main program entry point.....	105
Figure 8.3-18 - ImageDisplay.java test class source code.....	108
Figure 8.3-19 - ImageDisTest.java source code.....	110
Figure 8.3-20 - DBTestClass.java test class source code.....	112
Figure 8.4-1 - Database unavailable window	114
Figure 8.4-2 - Reference Image missing warning	114
Figure 8.4-3 - Mean .dat file missing error message.....	114
Figure 8.4-4 - Covariance .dat file missing error message.....	114
Figure 8.4-5 - Video source unavailable error message.....	115
Figure 8.4-6 - Final running main interface	115
Figure 8.4-7 - Results Graph	116
Figure 8.4-8 - Evaluation Framework Diagram	118

TABLE OF TABLES

Table 1.5-1- Project Plan Table.....	4
Table 3.6-1 - Face Detection system uses (Zhao, Chellappa, Rosenfeld, & Phillips, 2007).....	15
Table 3.7-1 - Comparison of face recognition methods,.....	21
Table 4.2-1 - System requirements table.....	28
Table 4.2-2 - User Requirements Table	28
Table 4.2-3 - System Specifications Table.....	29
Table 4.3-1 - Face Detection Specification	29
Table 4.4-1 - Feature extraction training specification	30
Table 4.4-2 - Data Testing Specification table.....	30
Table 8.4-1 - System tests	114
Table 8.4-2 - System requirements results ref. (Table 4.2-1).....	116
Table 8.4-3 - User requirements results ref. (Table 4.2-2).....	117
Table 8.4-4 - System specification results ref. (Table 4.2-3)	117
Table 8.4-5 - Face Detector Specification results ref. (Table 4.3-1).....	117

FACIAL RECOGNITION FROM VIDEO SEQUENCES

TABLE OF EQUATIONS

Equation 3.6-1- SNoW-based face detection model building.....	18
Equation 3.7-1 - 2DPCA Optimal Projection Vectors	21
Equation 3.7-2 - Eigenfaces Transformation equation.....	24
Equation 3.7-3 - Eigenvalue problem equation.....	24
Equation 3.7-4 - Eigenfaces coefficient equation.....	24
Equation 3.7-5 - Matrix equation derived from Equation 3.7-4 above	24
Equation 3.7-6 - Kernel problem equation	24
Equation 3.7-7 - Eigenfaces matching scheme equation.....	24
Equation 3.7-8 - Eigenvalues and Eigenvectors derivation equation.....	25
Equation 3.7-9 - Kernal function equation	25
Equation 3.7-10 - Main kernel equation.....	25
Equation 5.2-1 - Log-likelihood comparison	39
Equation 5.2-2 - Bayesian mean image building equation.....	39
Equation 5.2-3 - Bayesian mean image.....	39
Equation 5.2-4 - Bayesian Covariance building equation	39
Equation 5.2-5 - Bayesian addition of regularisation parameter equation	40

Chapter 1 – REPORT INTRODUCTION

1.1 Introduction

The purpose of this chapter is to demonstrate the motivation for the study being undertaken. It will begin with a problem definition that will include how the problem has emerged followed by the scope of the study. The scope will include a section on the limitations and constraints of this project. The final part of this section will list the minimum requirements expected by the project and what exactly is expected to be achieved by undertaking this work.

Although the concept of recognizing someone from facial features is intuitive and automated in the human brain from an early age, facial recognition, as a biometric, makes human recognition a more automated, computerised process. What sets apart facial recognition from other biometrics is that it can be used for surveillance purposes. For example, public safety authorities want to locate certain individuals such as wanted criminals, suspected terrorists, and missing children. They also want to automate such areas as border checkpoints and automate access to secure areas. Facial recognition has obvious advantages in this area, even though commercial uses have already been made of this software; Newham (Meek, 2002) borough Council has installed 140 cameras on their high street and have linked these cameras to a complex facial recognition database that has resulted in a ‘marked reduction in crime in the area.’

1.2 Problem Definition

Over recent years, Computer Vision has become a more interesting topic as its uses in practical scenarios becomes more apparent. Essentially, it is now more than just a theoretical exercise and instead offers real value if implemented correctly. This leap forward in possibility is due to the developments of technology and theory that now make it possible to implement Computer Vision systems effectively and efficiently.

There are numerous definitions for Computer Vision and Facial Recognition;

It has been suggested that Computer Vision is a broad field, which covers topics such as “image restoration, image enhancement, automated visual inspection, computer based image understanding of three dimensional scenes, and visual perception and cognition.” (Vernon, 1991).

Whereas Computer Vision is different to Image processing because “image processing is the processing of an image, typically by a computer, to produce another image, while Computer Vision is about image acquisition, processing, classification, recognition, and, to be all embracing, decision making subsequent to recognition.” (Low, 1991).

Facial Recognition is a subset area of Computer Vision and deals specifically with the processing of an image, the localisation of the face, the extraction of the features from the face and then the comparison of those features against some already held in some secondary storage medium to ascertain the likelihood of a match.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

This report considers facial recognition in a very particular context. It considers identifying a particular person by storing attributes from the features on their faces and comparing these features to those extracted from a live video feed (in this context, the video source could be considered to be “pseudo-live” as it may not be possible to run the necessary algorithms in real time). This identification will be made based on locations of the eyes in particular and will utilise several frames of video to obtain a mean distance between features. This is to minimise the error detected between the various features in each individual frame and to increase the accuracy of the overall system.

A current problem in the area of Computer Vision is the successful implementation of a working facial recognition system. As mentioned above, this area has excellent commercial potential for areas such as border control and could replace key lock mechanisms. In this scenario, the door could automatically be opened as a person approaches it given that they have a face held on the database.

However, there are few facial recognition systems that work reliably and even fewer that are cost effective. Essentially, the field of facial recognition currently remains firmly in the realm of research. This will remain the case until certain problems have been overcome. This is not to say that considerable progress has not been made in recent years in some of the problem areas; a robust face detection system now exists that accurately locates a face in a complex background and in a variety of sizes (Viola & Jones, 2004). Further problems that are still under investigation exist and some of them are addressed in this study:

- Facial characteristics; Glasses, hair styles and partial facial occlusion can all affect the ability of the face detector to accurately localise the face. Some improvements have been made in this area and the Viola & Jones face detector (mentioned above) are now fairly difficult to fool.
- Facial Expressions; a person that is smiling or frowning may lead to incorrect measurements being taken from around the subject's mouth. This is due to the fact that many facial recognition systems use training data to ‘train’ a system to be able to recognise a particular face. During this training period, the faces used are normalised in a single position and are therefore susceptible to changes in facial expressions.
- Lighting conditions; Variation in lighting conditions can throw off a feature extraction algorithm because the ‘templates’ are essentially incorrect. It then becomes very difficult to locate features and a possible match may be missed.
- Face Position; for many legacy facial recognition systems, it was expected that the face would be offered to the face detector as a ‘straight-on’ representation of the face. Even now, variations on this are extremely difficult to deal with. Even though the Viola & Jones face detector has the ability to locate faces from a ‘side-on’ view it is difficult to do this in an automated fashion. From a technical perspective, the face may be partially occluded if not straight on and the features will appear slightly closer together (or further apart) even if the face is accurately detected but the pose is partially off-centre.

1.3 Study Scope

The study scope, in short, is to gain an understanding of the techniques used in current facial recognition systems and then use a number of those methods to create a basic facial recognition system prototype that can extract certain features from a face and use those to compare the ‘live’ face to ‘test’ faces stored in a database.

The techniques compared include:

- Elastic two-dimensional shape matching
- Bayesian learning approach
- Two-dimensional PCA
- Kernel-based Eigenfaces
- Support Vector Machines (SVMs)

The prototype system should reliably match faces in a database based on facial features under conditions that are not strictly controlled. For example; with people having different facial expressions, slight variations in pose and wildly varying lighting conditions. The main evaluation of the system will revolve around its ability to deal with the aforementioned conditions and still accurately locate the face and the necessary features as part of that detection.

For each of the areas being considered; face detection, ‘individual’ feature localisation and comparison methods. Only one technique from the literature survey will be utilised, however, the reasons why that technique has been decided upon will be clearly stated. The following assumptions are made about the subjects in the images:

- The subject’s head will not rotate more than 10 degrees in any directions from a straight-on camera view for the duration of the video capture sequence. This constraint is due to limitations of the face detection algorithm.
- Lighting variations should be minimised to ‘normal’ video noise levels rather than drastic changes for the duration of the video sequence.
- Facial expression should be changed only minimally and extremely slowly over the course of the video capture sequence.
- Background is assumed and handled by the face detector but multiple faces are not at this stage. (The solution itself will support detection and processing of multiple faces but due to real-time video processing concerns, it is assumed that a single face is present. This also helps the application to ‘track’ the face over multiple frames because no separate facial tracking software is being used.)

1.4 Study Objectives

The primary objectives of this study are to produce a comprehensive report, specifically aimed at people of mathematically and computing backgrounds as is required by the BCS (British Computing Society.) This includes gaining a good understanding of computer vision, its application in this area and sound reasoning for the methods chosen to be used in the prototype.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

As this study deals with some complex subject matter and programming paradigms, a second objective is to gain a much better understanding of the high level language used to implement the prototype system (Java or C#).

The prototype system itself is a key objective of this study. It will accurately locate faces and facial features from a video sequence and will use these to attempt to classify the face being viewed. In addition, a graphical user interface will be provided to add new test data to the database and show the results of the video sequence. Finally, the system will use a ‘fusion’ of video frames in order to obtain a ‘mean’ face of the person being studied in an attempt to reduce the impact of erroneous feature detection and improve overall effectiveness of the application.

Finally, the system will be evaluated to ascertain its success or failure and then suggest improvements based on the findings of this appraisal.

1.5 Project Plan

<i>Date</i>	<i>Milestone</i>
October – December 2007	<ul style="list-style-type: none">• Background Research conducted• Programming mock-ups conducted for suitably
January 2007	<ul style="list-style-type: none">• Implement feature localisation in Matlab and in Java• Establish user requirements
January – Early March 2007	<ul style="list-style-type: none">• Full Requirements gathering and specification• <i>Further feature extraction research</i>
March – April 2007	<ul style="list-style-type: none">• Completion of prototype system• <i>Graphical interface design work</i>
April 2007	<ul style="list-style-type: none">• <i>Changes to video handling to take advantage of threaded processing</i>• Full project report write-up

Table 1.5-1- Project Plan Table

1.5.1 Project Plan Revisions (*Shown in italics above Table 1.5-1*)

As the project progressed, it became necessary to make revisions to the schedule as unanticipated work presented itself and as deadlines became tighter. The main revisions are detailed; more research was conducted late on for other feature extraction techniques as it became apparent that the original ones (those summarised in the mid-project report) were inadequate in real-time. It was also decided that more work should be focussed on the design of a usable graphical interface as this was essentially the only part of the system the user would see.

Finally, the completion of development presented problems with the flow of execution. Considerable effort was put into correcting this and is therefore cited as a change from the planned schedule.

Chapter 2 – FACIAL RECOGNITION

2.1 About Facial Recognition

Over the past few decades there has been a research focus in the area of computerised identification of the human face. Rapid and prolonged technological advances over a more recent period result in the ability to move this area of research from theoretical, into practical. Technological advances include; rapid advancements to the speed and efficiency of computer hardware (processors, memory and even webcams etc.), advanced support for video and image processing in high-level languages, the realisation of robust real-time face detection software (Viola & Jones, 2004) and increased focus on feature localisation techniques such as the Bayesian approach to eye localisation (Everingham & Zisserman, 2006).

Specific uses of automated facial recognition and where it would be particularly useful are as follows:

The identification of people as they approach, and wish to enter, buildings that have restricted access; currently, security guards are responsible for the labour intensive task of checking people against identification cards and confirming that they are indeed entitled to the access that they seek. This requires a large investment of staff but because humans are susceptible to error, it is likely that a large number of mistakes will be made over the course of the checks. These errors are the result of human tiredness, fatigue and general lapses in concentration. An automated face recognition system could check access to building (or even specific rooms) twenty-four hours a day, seven days a week and could alert a single guard if something suspicious occurred. Additionally, the system would be able to keep detailed logs of person movement with little extra effort. Finally, the system is not biased and does not become familiar with subjects, as happens in human-based systems, so the rules do not become relaxed.

Many more tedious and difficult jobs could be removed by using face recognition systems; a key area for this in the current political climate is terrorist identification and alert. The same principal applies to the search for missing children, this could be handled by face recognition software as an alternative to the classic American ‘milk carton’ approach. These applications and many others are proposed in a recent report to the state of Virginia crime commission (Horn, Gatune, Thomas, & Woodward, 2003) along with several constraints. The main issues revolve around the difficulty of recognising faces in real-time especially when comparing to an ever increasing database of reference faces. The notion seems to imply that processing of the faces may need to be handled ‘off line’ rather than in real time. However, this paper does fail to realise that the most recent improvements in computer hardware technology has the potential to process the images in near real time speeds. Running on ordinary computers (Pentium 3 1 GHz) the face detector software can accurately locate a face at fifteen frames per second. Despite this, many of the problems

FACIAL RECOGNITION FROM VIDEO SEQUENCES

present at the end of the last decade have gone some way to being alleviated; face detection in real-time, fast objection matching and image processing.

The reduction in cost of computer equipment mentioned above could render the password obsolete. A simple webcam could instead be used to properly identify a person when sitting in front of their computer. This also acts as a step to reduce the increasing problem of identity theft, as it becomes much more difficult to steal a person's face rather than just their name.

2.1.1 *Difficulties with current face recognition systems*

As mentioned in Chapter 1, the problems with facial recognition are slowly being eroded away. Recent advancements in many areas have made facial detection and recognition a reality and thus much work has begun in this field.

Many obstacles still exist to creating a working, reliable facial recognition system and are summarised as follows (Torres, 2004):

Pose variation (facial expression and facial position), illumination conditions, scale variability, images taken years apart, glasses, moustaches, beards, low quality image acquisition, partially occluded faces etc. A separate set of problems revolves around the fact that facial images (once acquired by the face detection system) are remarkably similar; they all have a mouth, a pair of eyes, a pair of ears and a nose etc.

The problem of image condition is being dealt with slowly however, the face detector mentioned previously (Viola & Jones, 2004) is robust under many different conditions. An existing problem with face detection was the 'noise' (areas around the face that may throw off the detector,) this has been almost alleviated by the Viola and Jones face detector. It is now possible to locate a number of faces against complex backgrounds. In Figure 2.1-1 below, the faces are accurately localised, with a single false positive result. All of the below faces have regular facial characteristics (no beards, moustaches etc.) and this may have influenced such positive results.




Figure 2.1-1 – Face detection using Viola & Jones face detector (Bornet, 2005)

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Another problem linked to the condition of the acquired image is the lighting of that image. Changes in the lighting around the person being identified can throw the matching process off. This problem has been to locate the facial features. This is especially true when using a learning algorithm in order to do the comparison because the 'mean feature' will have a particular lighting variation.

Facial expression is an existing problem in recognition systems. Variations in facial recognition; a face frowning as opposed to a face smiling for example essentially changes the proportions of the face and can cause the recognition system to make an incorrect match:



Figure 2.1-2 - Changes in facial expressions (Dyer, 2004)

Significant advances have been made to remove this problem; a system is proposed to accurately map a face even though the expression may be in a state of constant change (Kittler & Nixon, 2003), their system builds a three dimensional model of the face and utilises Principal Component Analysis (PCA) to allow the face to be measured even though its basic dimensions may have changed. This work is based on further studies into using PCA to recognise faces which is invariant to facial expressions (Turk & Pentland, 1991). Eigenfaces refer to this PCA analysis and are the collection of Eigenvectors measuring the distance between various facial components. They are now considered to be a defacto standard in facial recognition and have been used successfully in the aforementioned cases to track faces in an expression invariant manner.

In summary, the following problems have existed for computer vision and specifically facial recognition since it came into research focus (Torres, 2004):

- Pose variation (facial expression and facial position)
- Illumination conditions,
- Scale variability,
- Images taken years apart,
- Glasses,
- Moustaches,
- Beards,

FACIAL RECOGNITION FROM VIDEO SEQUENCES

- Low quality image acquisition,
- Partially occluded faces
- Facial images are remarkably similar; they all have a mouth, a pair of eyes, a pair of ears and a nose etc.

Many of these problems have slowly been addressed over recent years; Lighting, Scale, Pose, Low quality images and partial occlusion are now dealt with effectively by a number of face detectors and PCA-based feature extraction algorithms.

2.2 Human interpretation of Faces

Humans can identify a large number of faces with ease. In contrast, computer systems find it extremely difficult to process a face containing an image and accurately match it to another face.

It is reasonable to assume that humans do not simply use pattern matching techniques in order to recognise a person's face but that they use their innate instincts in order to do it successfully. Humans can even recognise people long after they last saw them because of other non-visual queues, such as gestures or the way they talk. It is also not unreasonable to assume that if a person only had a short glimpse of another person (similar to facial recognition systems) they would not be able to differentiate when presented with thousands of subjects. This is one of the reasons that make automated facial recognition so hard, especially given that a system must usually assume that the face is unfamiliar and look for features which could indicate that they are the same person. However, given time to learn the face, humans can know up to seven hundred faces on a personal level and even potentially be able to recognise thousands of others (Samal & Iyengar, 1992); computers too should have this ability but a database that large invariably incurs a performance problem.

Faces in human memory can be thought of as being viewed in multidimensional space, the size of this space increases through increased exposure to the faces in the space (Pezdek, Blandon-Gitlin, & Moore, 2003).




Figure 2.2-1 - Representation of the child face space. Inferred from (Pezdek, Blandon-Gitlin, & Moore, 2003)

Figure 2.2-1 above shows how children (and adults) perceive facial images. Faces are essentially split between distinctive and common. Common faces lie close to the centre of the space where there is less variation and therefore they are more difficult to distinguish from one another. More distinctive faces lie closer to the boundaries of the area and are therefore easier to differentiate. Given this information and the fact that, by definition, common faces are more frequent than distinctive ones, it is perhaps understandable that computers struggle to differentiate between the faces because they are so similar.

The following chapter examines some of the more prominent methods of face acquisition and feature extraction.

2.3 Conclusion

Automated facial recognition systems would have many uses in the real world amongst public and private sectors users; law enforcement and building security are a few of the more obvious ones. In order to realise the use of facial recognition as a reliable tool, a number of problems need to be overcome and researchers in this area are working hard to overcome the issues (a number of which, as mentioned, have already been overcome to some extent.)

In addition to the technical advances being made in this field, it is important that the legalities are not overlooked and that suitable decisions are made about the objects that these systems guard. It is important that no facial recognition system impede any individuals' rights with regard to data protection and personal liberties. Extreme caution should be taken when deploying systems and when classifying reference individuals. In addition, a key threshold should be established to ensure that the system has a higher false positive rate or false negative rate depending on the type of application. For instance, access to a high security facility by means of facial identification

FACIAL RECOGNITION FROM VIDEO SEQUENCES

should have a higher false negative rate to prevent erroneous access. If a legitimate subject were refused entry, this could be easily overcome with less severe consequences than the alternative.




Figure 2.3-1 - False acceptance, false rejection rate example (Efford, 2006)

Finally, when designing a system to achieve facial recognition, serious thought should be attributed to the difficulties identified above (Torres, 2004) and methods chosen which are invariant, or at least partially invariant, to as many of them as possible.

Chapter 3 – LITERATURE SURVEY ON FACIAL RECOGNITION TECHNIQUES

3.1 Introduction

This chapter is designed to give a detailed overview of the different methods currently being used for facial acquisition and feature extraction that can be used for automated facial recognition systems. The chapter will give a brief overview of each method, how it works and the disadvantages and advantages that it offers. The discussion will cover popular methods and recent important research focuses in the field that may be of interest to future systems. The chapter will conclude with a comparison of some of the various methods.

3.2 Introduction to automated Face Recognition systems

Research into the field of computing is increasing as more efficient and effective methods become available to the public domain. This section looks at some of these methods and how successful they have been.

The actual process of recognising a face can be split into a number of different processes; the detection of a face amongst a cluttered background, the extraction of a number of features from the face area collected above and finally the comparison of those features against a number of reference faces held in secondary storage in order to produce a match. In this report, the location of the face in the video frame is not predetermined and must be established by the program itself using one of the mentioned methods. The

FACIAL RECOGNITION FROM VIDEO SEQUENCES

process following this could be to take the face as a single component and attempt to match it to a reference face in that way, however the failings of this method are well documented (Heisele, 2006). Therefore, the process to be used instead is to extract an as yet undetermined set of features and then use those for comparison.

A common set of steps is usually undertaken by most systems in order to attempt to recognise a face (Samal & Iyengar, 1992):

1. Determine a set of features to represent the subject's face
2. Transfer the features identified in step 1 to a database or other secondary storage
3. Determine a set of features (the same features as step 1) for a previously unseen face
4. Use an algorithm to match the features in steps 1 and 3 held in step 2 and system memory respectively

Some background information is needed about the steps above and is presented in sections 3.7.1 through to 3.7.5.

3.3 Feature Extraction

Although many traditional facial identification systems work by matching whole faces using templates or some other geometrics, it has become much more common for systems to attempt to locate any given number of facial features and compare the measurements of those features in relation to one another. Once located, performing the necessary functions on the features becomes relatively straight forward. Therefore, most design effort should be placed on the effective localisation of the features.

The parts of the face vary in importance with regard to how significant they are to the distinctiveness of any particular human face. In no particular order, the following parts of the human face are often described as being most significant to the identification of that face (Zhao, Chellappa, Rosenfeld, & Phillips, 2007):

- Hair
- Face Outline
- Eyes
- Mouth

Extracting the aforementioned features efficiently and effectively is crucial to creating a reliable facial recognition system. The nose plays a very insignificant role in the recognition of a face and as its shape does not lend itself to pattern matching, so should typically be ignored. As a side point; the upper portion of the face is more important than the lower portion and this fact could be utilised to increase performance (or at least reassign processing from lower to upper accordingly.) Studies have also shown that the attractiveness of faces affects their ability to be recognised easily. However this, at least at our current technology level, is impossible to translate into a computer algorithm.

Because of the difficulty of extracting features from a face it is important to focus efforts on those features than can be gathered most easily and which are most invariant to the conditions mentioned in section 2.1.1. The following prerequisites are inferred from a recent computer vision conference for the effective use of facial features (Jiao, Gao, Chen, Cui, & Shan, 2002):

FACIAL RECOGNITION FROM VIDEO SEQUENCES

- Estimating the area where the features are most likely to occur must be as simple as possible
- Changes in facial expression should have a small effect on the overall location of the feature
- The information posed by the feature should be crucial to the overall recognition of the face
- Light conditions should have as small an effect on the feature as possible

It is critical that the features used are quick to locate but that they also conform to the above standards where possible. If the feature is sensitive to some of the problems cited in section 2.1.1 then the feature may not be useful in an unnormalised environment. This calls into question the viability of an expression variant mouth in the above stated scenario.

The main methods used for the duration of this report are the Bayesian learning approaches and edge detection approaches. These are discussed in the following paragraphs.

The Bayesian learning approach refers, in essence, to an algorithm that uses training data to create a 'template' which can then be compared to every part of the test image to find a match. The match should be the centre of the patch that most closely resembles the feature being searched for.




Figure 3.3-1 – Average eye image (acting as template)




Figure 3.3-2 – Eyes localised on test image (Zhao & Grigat, 2006)

The above images show the uses of template based matching that can be used in facial recognition systems. The red line in the image on the left, above (Figure 3.3-1) is the visual representation of an 'average' eye. This is created by looking at a number of eyes and creating a mean model over the entire set. That model is then compared to every section of the test image and the most similar two results, in theory, would be the eyes. The results are shown above right (Figure 3.3-2). A similar technique can be used to localise the other features.

The template image is compared to every part of the test image in a sliding window fashion and an algorithm is applied to ascertain how similar the template is to the patch being investigated.

FACIAL RECOGNITION FROM VIDEO SEQUENCES



Figure 3.3-3 - Test patch



Figure 3.3-4 - Sliding Window test image

The above figures show how the sliding window approach works. The red square (size x by y Figure 3.3-3) represents a subset of the test image, a window within a picture. That subset is slid across the test image (Figure 3.3-4) so that every set of pixels (size x by y) is processed in turn.

The second main method used to localise features is the edge detection method; either by detection of edges or by detection of blobs and then using some sort of principal component analysis to establish which blobs are collection of detected edges correspond to each individual facial features. The image below illustrates the use of edge detection and blob detection (Kawaguchi, Rizon, & Hidaka, 2005) algorithm:



Figure 3.3-5 - Results of blob detection

As you can see from the above image, it is fairly easy to locate the eyes, eye brows and hairline once the other information has been eliminated. From there it is straightforward to estimate the location of the nose and mouth from the remaining blobs, as it is always located directly below the mouth.

3.4 Feature matching

Matching the selected feature to the feature on the test face is a critical part of facial recognition and if done incorrectly, invalidates the entire recognition process.

Four commonly used methods are cited (Zeilek & Ersi, 2006):

- Correlation: Templates of a known face and of an unknown face are compared to ascertained differences.
- Feature cluster: The degree of similarity between two faces is ascertained by ‘clustering’ certain areas of the face.
- Set feature/aspect partitioning: certain features on the test face will have vectors that are very different to those on the training images. Therefore, using a process of elimination, those features can be removed and faces categorised as non-matching.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

- **Euclidean Distance:** A comparison of the feature vectors on a training set of images against a test image reveals how closely related the faces are and therefore the smallest number should yield a facial recognition match. This is the most commonly used matching scheme.

The approach to be used depends on the data being utilised but some of the methods are more common than others.

3.5 Image Normalisation

When an image is captured from the capture device, whether it is a static photograph, a video segment or a live video feed, the image is in an unnormalised state. This means that the images could have any pose, any lighting conditions, and the background may have unwanted clutter. All of these conditions as mentioned above could have a detrimental effect on the accuracy of an automated face recognition system.

Normalisation was an important process to legacy facial recognition systems because of their brittleness. If a face was not exactly in the same position as the reference faces used to create the model, then a match would be almost impossible to obtain.

The process of normalisation, in a traditional sense, refers to a sequence of steps that are completed prior to any matching taking place. The objective is to remove variations in position, scale, and rotation of the facial image (Brunelli & Poggio, 1993). Normalisation takes place during the learning phase and on test data before any attempted match is made. This makes it possible to match static faces that are at different angles to the camera, under different lighting conditions or if the face was closer to or farther away from the camera than when the reference face was acquired. Traditionally, two points on the face may be chosen in order to ascertain the scale and position of the face prior to matching. This process relies on accurate feature extraction. Obviously, this does not take into account the effect of lighting; in order to standardise the lighting of an image so that it may be used for recognition purposes a series of filters may need to be used to standardise the lighting at each pixel location (Heusch, Cardinaux, & Marcel, 2003).

In the area of automated facial recognition from video, normalisation is difficult to perform in real-time. Serious thought must be given to the way in which images are normalised. It is assumed for the purposes of this report that the training data (if used) will be fully normalised but that the test data will be input from an unpredictable live video source and as such will not be normalised. Although, the pose of the head will be within certain limits due to the fact that a 'front on' face recognition system is being developed, not a 'profile' system.

3.6 Face Detection Techniques

The face detection itself is quite a distinct activity from the recognition of a face and is performed before any feature extraction can take place. The process is important because it reduces the search space from an entire image to just a facial area and reduces the likelihood of errors during feature extraction. The process needs to be fast, so that it can be performed in near real time and ideally should be invariant to many of the obstacles mentioned in section 2.1.1.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

3.6.1 Viola and Jones Face Detector

The Viola and Jones face detector (Viola & Jones, 2004) represents a leap forward in facial detection techniques. It can perform accurate facial detection at a rate of fifteen frames per second on standard home architecture (Pentium 3). The detection process is split into 3 main stages; each stage is based upon work previously done in the field. The uses for a well designed, efficient and robust face detection system are extensive. Some of these uses are summarised in Table 3.6-1 below:

<i>Areas</i>	<i>Specific Applications</i>
Biometrics	Drivers' licenses
	Immigration, National ID, Passports, Voter Registration
	Welfare Fraud
Information Security	Desktop Logon
	Application Security Database Security, File Encryption
	Intranet Security, Internet Access, Medical Records
	Secure Trading Terminals
Law Enforcement and Surveillance	Advanced Video Surveillance, CCTV Control
	Portal Control, Post Event Analysis
	Shoplifting; suspect tracking and investigation
Smart Cards	Stored Value Security, User Authentication
Access Control	Facility Access, Vehicular Access

Table 3.6-1 - Face Detection system uses (Zhao, Chellappa, Rosenfeld, & Phillips, 2007)

The detection system uses an AdaBoost learning algorithm:

In an AdaBoost algorithm a number of weak classifiers are combined to create a smaller number of strong classifiers.

The simple learning algorithm is a weak learner; this is because even the best classifiers do not classify the data particularly well. In order to boost the classifier, it must first be made to solve a series of learning problems. After the first round of learning, the training set is re-assessed with weights according to the accuracy of the classification. Hence, this is where it learns the information it needs to better classify the data in subsequent training cycles.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

The final strong classifier is a weighted combination of weak classifiers followed by a threshold. AdaBoost essentially finds a small number of good features, which have significant variety.

In the Viola and Jones' face detector, the AdaBoost algorithm is used to minimise the number of Haar-like features being computed. It does this to increase the computational efficiency of the detector by eliminating features that are unlikely to be of any help to the overall classification. Each weak classifier is constrained to a single feature for the purposes of this classification.

Viola and Jones used several other important methods to build their state-of-the-art facial detection system:

- Focus of attention approach:
This approach allows the system to locate areas of an image where the face is most likely to occur in order to avoid over-processing an area where a face definitely will not occur.
- The integral image:
The integral image is an intermediate representation of the original image which contains a subset of that image; at point (x,y) the integral image is all of the points above and to the left of point (x,y) .




Figure 3.6-1 - The integral image example

This provides the basis for the cascading detection used to quickly filter out unnecessary data.

In the Viola and Jones implementation, the integral image is an image that ignores the intensity of the original. Integral images can be computed extremely quickly and the results have the benefit of being able to compute haar-like features at any scale or location in constant time. The system does fully realise that the rectangular feature matching system that it uses is imperfect for matching to variable shapes.

However, the extreme computational efficiency more than makes up for this short fall and fully explains why alternatives like Steerable filters are not used.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Using classifiers, the focus of attention operator discards 50% of an image at a cost of approximately 1% of genuine faces. The classifiers are evaluated in 20 simple operations per location.

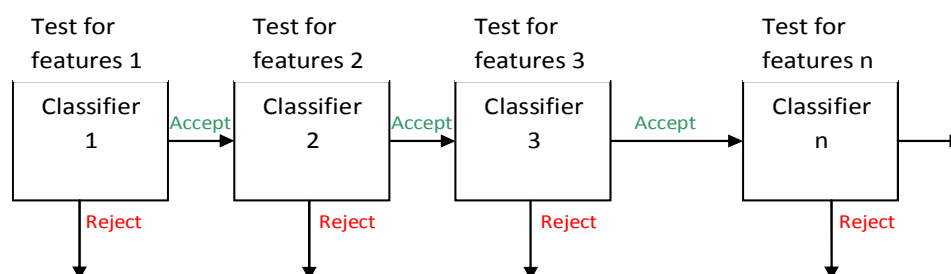


Figure 3.6-2 - Weak classifiers used by Face Detector

The classifier cascade described above removes sub windows that do not contain faces from the original image. Each classifier eliminates more non-faces whilst retaining nearly 100% of the actual faces.

The Viola and Jones face detection system performs fifteen times faster than the next most efficient face detector constructed by (Rowley, Baluja, & Kanade, Neural Network-Based Face Detection, 1998).

Using features for classification, as opposed to pixels saves a great deal of wasted computational efficiency, the system is therefore much more effective than its predecessors given the same time frame. This speed is proportional to the number of features evaluated. A balance has to be found where the minimal number of features is evaluated in order to obtain a ‘good’ rate of successful facial detections.

The Viola and Jones system was tested on a well-studied facial test set. The results against this were excellent and the fact that the set was extremely complex proves the reliability of the system.

The system itself is slightly less accurate than previous best efforts because the weak classifiers, by definition, can cause important information to be lost without being properly evaluated. This is an accepted risk of the system and its accuracy is well within acceptable limits. In addition, this slight loss in accuracy vastly improves the speed of the algorithm making it a much more viable choice for real-time face detection.

3.6.2 Sparse Network of Winnows (SNoW) Face Detection

Another type of face detector has been created in the past utilising an approach known as the Sparse Network of Winnows (SNoW) approach (Roth, Yang, & Ahuja, 2000).

The sparse network of winnows approach is a learning architecture consisting of a sparse network of linear units over an incrementally learned feature space. Each linear unit represents a face or a non-face pattern. The face image (or non face image for learning purposes) is offered as an input layer to the linear units (the input layer consists of a number of values representing the features contained within the input image¹) and these are used to augment the linear units accordingly.

¹ The values representing the features can be Boolean (indicating presence or lack thereof) or real (indicating the strength of the given feature)

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Once the linear units have received all of their input data, they are linked via weighted edges to input features to build a model of the face. The presence of a face is dependent on the presence (or lack) of active a number of active features. Each individual feature is active if its weighted edge value is greater than a threshold value. The threshold value is set during the learning phase. Also during the learning phase, the algorithm must be offered positive (face) examples and negative (non-face) examples to learn from, these must be labelled at input time.

The algorithm can be shown as:

$$\sum_{i \in At} w_i^t > \theta_t$$

Equation 3.6-1- SNoW-based face detection model building

Where $\sum_{i \in At}$ represents the set of features in the image being tested is, w_i^t is the weight of the edge connecting the i^{th} feature (in the set) to t (the linear unit) and θ_t is the threshold to determine which values should be taken into consideration.

The learning algorithm is mistake driven. It therefore utilises an update rule, in this scenario, this is the Winnow update rule. The linear units are thus updated as the algorithm learns; $\alpha > 1$ is used as a promotion parameter when the algorithm works correctly and $0 < \beta < 1$ is used as a demotion parameter when the algorithm is incorrect. This is how the learning feature works and weights correct responses accordingly. The key to the Winnow update rule (the only rule applied in this detector) is that the number of examples it requires to learn a linear feature grows linearly with the number of relevant features and grows logarithmically with the total number of features being used.

The SNoW-based face detection algorithm is extremely accurate in detecting faces given appropriate learning material. In addition it can accurately locate a face in a number of poses, scales and under varying lighting conditions. It is therefore extremely robust. However, when compared to the algorithm above (Viola & Jones, 2004) it has a slight advantage in terms of accuracy but cannot compete in terms of speed of processing. This is due to the fact that the algorithm uses more costly processing and evaluates all areas of the image, not saving any processing power by utilising a weak classifier prior to any other processing. Indeed, this is not suitable for this algorithm. Also, the SNoW-based face detector cannot currently detect rotated faces. This, therefore, creates the need to impose difficult to enforce normalisation rules on the video image².

Finally, no information is available with regard to this approaches ability with processing video frames. Although, every indication is that the algorithm could process multiple frames per second, it is likely to perform much less effectively than the aforementioned method.

3.7 Feature Extraction Techniques

Following the successful acquisition of a human face from a cluttered background, it is necessary to extract several of the features from that face quickly and accurately in order to process those features and

² A solution does exist to this problem (Rowley, Baluja, & Kanade, Rotation Invariant Neural Network-based Face Detection, 1998) but offers even further reduced performance.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

attempt to recognise the face. There are many proposed methods for the successful acquisition of facial features; some of the most popular and more recent advancements are described in the following sections.

3.7.1 Bayesian approach to feature localisation

This is an approach, in part, attributed to a paper prepared on eye localisation (Everingham & Zisserman, 2006). According to which, is more successful and affective than some other approaches mentioned in the same paper. It is described below.

The accuracy of facial recognition using Principle Components Analysis (PCA) or Linear Discriminant Analysis (LDA) is known to degrade with poor eye localisation. This approach offers a robust method to locate the eyes reliably, thus improving the overall accuracy of facial recognition systems using the above stated methods. Essentially, the eyes are considered to be amongst the most important features of the face and it is important that they are located using a reliable method.

The system attempts to classify patches as either being not an eye or a patch containing an eye.

Three approaches to the successful localisation of the eyes are explored. Each one is different and has varying degrees of success:

- Regression method:
 - Parameters are learnt by minimising the Euclidean distance between true and predicted eye positions. This type of learning approach should offer promising results over a large dataset.
- Bayesian Method:
 - The Bayesian method is essentially an algorithm that gives a positive output if the patch does contain an eye or a negative output if the patch does not contain an eye. This makes it easy to classify if any particular patch does, or does not, contain an eye. A number of algorithms are proposed in order to implement this method.
- Sliding Windows:
 - Bayesian and discriminative approaches are used to locate the eyes. Square patches are extracted from the test image along a line in a 'sliding window' fashion. The classifier is then applied to a number of pixels, using each pixel in the square area as the centre of the constant value γ has the effect of expanding the distribution to make the matrix useable.

These approaches are compared in Figure 3.7-1 below:

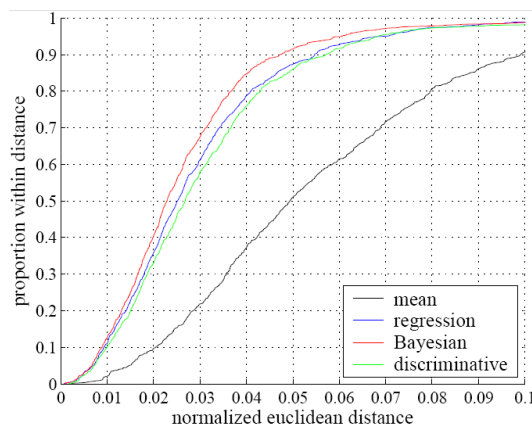


Figure 3.7-1- Results of errors of classification and regression eye localisation methods performed on FERET database (Everingham & Zisserman, 2006)

FACIAL RECOGNITION FROM VIDEO SEQUENCES

According to the published results for these methods (Everingham & Zisserman, 2006) and as briefly mentioned above the Bayesian approach offers the most reliable results amongst the popular methods offered by this paper. Figure 3.7-1 above shows how successful the Bayesian (classification) approach actually is. Given that it is also the most simple and computationally efficient algorithm to implement, it is should be very effective effective.

In 90% of images, the eye was localised to within 2.04 pixels during extensive tests. The total difference between the two eyes is therefore up to 4.08 pixels per pair of eyes. These are excellent results. However, there is potentially an identification problem with this type of detection error because it could fool the matching scheme into believing the subject is somebody else with similar eye spacing distance. Conversely, this is likely to be the best algorithm available for accuracy and computational efficiency.

Because this method was decided upon for use in the final prototype, the exact details about its implementation are presented in section 5.2. This section is meant as an overview to the method for consistency.

3.7.2 Two-dimensional PCA-based feature localisation

Principal Components Analysis (PCA) has been used extensively and successfully in the area of facial recognition. The process has been in use since the early nineties with many examples of its use recorded (Kirby & Sirovich, 1990).

PCA is used commonly in the field of computer vision as a cluster analysis tool. It is designed to capture the variance in a dataset in terms of principle components. The dataset refers to the images being used for comparison (in this instance the face training images.) The PCA process attempts to reduce the amount of data in the images whilst retaining the defining parts and simultaneously filtering out noisy data, such as background and other non-facial data. This model is illustrated in Figure 3.7-2. Normalisation can remove some of the variance in the data making the PCA process less effective.




Figure 3.7-2 - Principal Component Analysis example
(University College London; Oncology, 2006)

The blue lines above represent two separate principal components (PC-1 and PC-2). In reality, a face would have many more principal components. PCA is based upon covariances and the information they give about multiple dimensions of data. The key to principal components are that they work better with increasing dimensions of data as they take advantage of compression by removing noisy data that is existence in covariance. The principal components are selected from the single-dimensional eigenvalues with the highest values (Smith, 2002).

FACIAL RECOGNITION FROM VIDEO SEQUENCES

The above paragraph should explain why PCA is so useful to use with complex facial data.

PCA has many advantages but is becoming an old method and improvements have been made. The single biggest problem with PCA is that it is single dimensional. Improvements have been suggested (Yang, Zhang, Frangi, & Yang, 2004).

The optimal projection vectors of two-dimensional PCA x_1, \dots, x_d are used for feature extraction in a similar way to traditional PCA. Given a sample image **A**:

$$Y_k = \mathbf{A} \mathbf{X}_k k = 1, 2, \dots, d$$

Equation 3.7-1 - 2DPCA Optimal Projection Vectors

Then, we have a family of projected feature vectors, $\mathbf{Y}_1, \dots, \mathbf{Y}_d$ which then form the principal components vectors³ of the given sample image (in this case **A**). From these vectors, a matrix is created called the feature matrix. The feature matrix is then compared to its neighbours to assess their similarities in much the same way as PCA does.

Method	Recognition Rate
Fisherfaces	94.5%
ICA	85.0%
Kernel Eigenfaces	94.0%
2DPCA	96.0%

Table 3.7-1 - Comparison of face recognition methods, adapted from (Yang, Zhang, Frangi, & Yang, 2004)

Perhaps unsurprisingly, the results of this approach offer better recognition rates than standard PCA (and in fact many other results too, see Table 3.7-1 above.) However, something which is surprising about this method; it uses less principal components than standard PCA offering a performance boost (see Figure 3.7-3 below.) Despite this, PCA is still a lengthy process and the processing time involved in 2DPCA prior to testing is extensive.

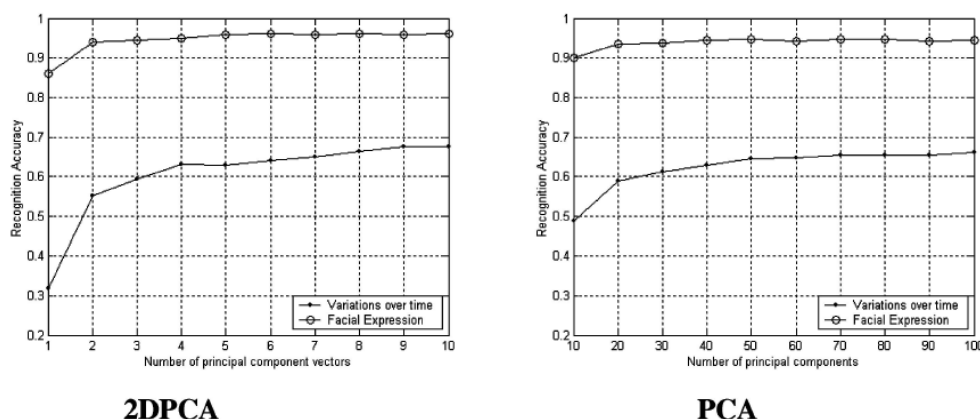


Figure 3.7-3 - Comparison of no. of PCs in PCA & 2DPCA (Yang, Zhang, Frangi, & Yang, 2004)

³ In 2DPCA these values are vectors, in PCA they are scalars

In addition, because the PCA process attempts to extract a subspace where the variance is maximised (and the reconstruction error is minimised) changes in lighting and facial expression can cause unreliable results. Especially given that these changes in lighting and pose cause greater variations in the images than a change in facial identity, making identification extremely difficult.

3.7.3 *Edge-based shape detection*

The method now being considered is an edge-based shape detection algorithm. It is likely that whilst attempting to recognise objects, humans (at least partially) look at the edges of that object to ascertain its shape and whether that edge belongs to the object at all. Essentially the approach looks to extract full shapes rather than parts of an edge map. This has several advantages over the edge map approach, outlines below.

A recent edge-based shape detector is proposed (Moon, Chellappa, & Rosenfeld, 2002):

The approach considers the generalised Hough transform as a means of detecting features in test images. However, it also recognises the approaches limitations; it is rotation invariant and relies on salient points in order to work properly. Therefore, this particular approach leads to poor feature localisation. Fitting a polynomial curve to the outline of a face has also been considered in much the same way as an active-shape model, but the shape is too approximate to be of any use. Traditional edge detection algorithms suffer from too much emphasis on localised information and can discard edges without evaluating all of the relevant information.

In order to ascertain the presence of a particular shape, the pixels around the test pixel are evaluated for changes to intensity rather than examining the pixel itself. In this method, we must define the shape matched operator by expanding the most favourable step edge operator (which minimises the noise and mean squared error) along the speculative object boundary contour so that responses from intensity differences along the boundary are collected at the same time in order to improve performance. The responses to this process are gathered at the centre of the operator where a log-likelihood ratio determines whether a shape is present or not.

Because facial features are easy to model as two-dimensional shapes, they can be located using this method; circular irises, intensity differences at the corners of the eyes, and the shade between the upper and lower lips. To locate an eye, a circular shape can match to the eyes whilst a line could approximate the eyelids directly above. A straight-line segment can be used to approximate the position of the mouth and this shape can be bent by moving a pixel in the middle of that line to approximate facial expression, see Figure 3.7-4 for a visual example.



Figure 3.7-4 - Feature detection using shape matching

This method offers an extremely simple method for the accurate location of facial features by matching basic shapes. This model could truly work in a real-time environment and offers some resilience to facial expression changes. However, the model is susceptible to changes in pose (outside the bounds of regular movement) and the fact that the algorithm relies on intensity changes in certain places, significant changes in image illumination could have a negative impact on this process.

The method proposed here utilises edge detection and model fitting to basic shapes. The likelihood of a match is a probability function extremely similar to the one cited in section 3.7.1 (Everingham & Zisserman, 2006).

3.7.4 *Kernel Eigenfaces and Fisherfaces approach comparison*

Principal Components Analysis (3.7.2) and Fisher linear Discriminant methods have been successful in the area of facial recognition. However, they only look at lower level statistical dependencies between pixels and do not address higher levels such as the relationships between a number of pixels (three or more).

In order to address this problem the use of Kernel Eigenfaces and Kernel Fisherfaces is proposed here (Yang M.-H. , 2002).

The idea of this method is to expand PCA, Eigenfaces and Fisherfaces approaches so that they may take account of higher order statistics and make use of a greater order of surrounding pixels because it is believed that this extra information is the future of facial recognition. Kernel Eigenfaces and Fisherfaces as presented here, not only attempt to extract higher order statistical information but also maximise the class separation when these features are projected onto lower dimensional space for easier recognition.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

3.7.4.1 Kernel Eigenfaces (Kernel PCA)

Basic PCA is explained in detail in section 3.7.2, as part of the explanation of two-dimensional PCA. The explanation here explains how the Kernel method attempts to make the algorithm more efficient. In simple terms, each vector is projected from the input space to a high dimensional feature space, in order to take advantage of higher order statistical data. This transformation is handled by a non-linear mapping:

$$\varphi: R^n \rightarrow R^f, f > n$$

Equation 3.7-2 - Eigenfaces Transformation equation

Where R^n is representative of the input space, R^f represents a high dimensional feature space. The feature space can be very large. For R^f , the eigenvalue problem is defined as:

$$\gamma \mathbf{w}^\varphi = C^\varphi \mathbf{w}^\varphi$$

Equation 3.7-3 - Eigenvalue problem equation

C^φ is a covariance matrix acting as a regularisation term in this case. All solutions \mathbf{w}^φ lie in the span of $\varphi(x_1), \dots, \varphi(x_m)$ where x is a vector projected from the original PCA algorithm, m is given here as the number of components. There exist coefficients (α_i) for these such that:

$$\mathbf{w}^\varphi = \sum_{i=1}^m \alpha_i \varphi(x_i)$$

Equation 3.7-4 - Eigenfaces coefficient equation

The above equation denotes an $m \times m$ matrix K by:

$$k_{ij} = k(x_i x_j) = \varphi(x_i) \cdot \varphi(x_j)$$

Equation 3.7-5 - Matrix equation derived from Equation 3.7-4 above

The above equations form standard PCA results, the preceding equation becomes a Kernel problem:

$$m\gamma K\alpha = K^2\alpha \equiv m\gamma\alpha = K\alpha$$

Equation 3.7-6 - Kernel problem equation

Above, α denotes a column vector with entries 1 through m (form the matrix). All projected samples are centred in the high-dimensional feature space. The vectors in that high-dimensional feature space must be projected onto a lower space spanned by eigenvectors \mathbf{w}^φ for matching schemes to attempt recognition. \mathbf{X} is a test sample whose projection is $\varphi(x)$ in R^f :

$$\mathbf{w}^\varphi \cdot \varphi(x) = \sum_{i=1}^m \alpha_i (\varphi(x_i) \cdot (\varphi(x))) = \sum_{i=1}^m \alpha_i k(x_i, x)$$

Equation 3.7-7 - Eigenfaces matching scheme equation

3.7.4.2 Kernel Fisherfaces

The test vectors x are again centred in high-dimensional space R^f . Within-class scatter matrix S_w^φ and between-class scatter matrix S_B^φ , it is necessary to find eigenvalues γ and eigenvectors \mathbf{w}^φ :

FACIAL RECOGNITION FROM VIDEO SEQUENCES

$$\gamma S_w^\varphi w^\varphi = S_B^\varphi w^\varphi$$

Equation 3.7-8 - Eigenvalues and Eigenvectors derivation equation

These can be obtained by some matrix calculations, normalising the matrices and giving a vector as output.

The kernel function is then described as:

$$(k_{rs})_{tu} = k(x_{tr}, x_{us}) = \varphi(x_{tr}) \cdot \varphi(x_{us})$$

Equation 3.7-9 - Kernel function equation

Let K be an m x m matrix defined by the elements $(K_{tu})_{u=1, \dots, c}^{t=1, \dots, c}$ where K_{tu} is a matrix composed of dot products in the high-dimensional feature space.

Without covering the rest of the proof of the equation, the main equation to be used by this method can be shown below; more details are available in a recent paper comparing these methodologies (Yang M.-H. , 2002):

$$\begin{aligned} \mathbf{w}_{OPT}^\varphi &= \arg \max \mathbf{w}^\varphi \frac{|(\mathbf{w}^\varphi)^T S_B^\varphi \mathbf{w}^\varphi|}{|(\mathbf{w}^\varphi)^T S_w^\varphi \mathbf{w}^\varphi|} \\ &= \arg \max \mathbf{w}^\varphi \frac{|\alpha K Z K \alpha|}{|\alpha K K \alpha|} = [\mathbf{w}_1^\varphi \dots \mathbf{w}_m^\varphi] \end{aligned}$$

Equation 3.7-10 - Main kernel equation

Projecting this high-dimensional representation into lower dimensional space for comparison can be achieved in much the same way as in Kernel PCA in section 3.7.4.1. \mathbf{w}_{OPT}^φ above is the set of optimal eigenvectors gathered from the process by means of evaluating the class scatter matrices S_w^φ and S_B^φ respectively. The extracted eigenvectors gathered from the equation directly above are Kernel Fisherfaces; utilising the standard Fisherfaces method described by many sources (Belhumeur, Hespanha, & Kriegman, 1997) and (Choi, Lee, Lee, & Yi, 2001), but extrapolating the results to high dimensional space for analysis before reducing them to lower-level orders for comparison and matching.

Unfortunately for this method, even during the training phase, each image was down sampled to 23 x 28 pixels because of the complex mathematics involved in multiple matrix calculations. According to results (Yang M.-H. , 2002) done on these methods (amongst others) the Kernel Fisherfaces results were the most successful. In addition, the kernel methods for Eigenfaces and Fisherfaces performed better than their original counterparts. The methods both perform better than their traditional counterparts when considering them in the context of pose and expression variations. Although, it is worth noting that the SVM (Support Vector Machine) method performed better than any of the methods under these conditions.

In short, under normalised conditions and with no expectations of processing efficiency, the kernel augmented methods perform better than traditional Eigenfaces/Fisherfaces. The error rates are markedly lower and the accuracy rates are high.

However, the methods are computationally inefficient and as mentioned, have serious problems dealing with facial expressions and pose (as do conventional PCA-based methods). Therefore, this approach may

FACIAL RECOGNITION FROM VIDEO SEQUENCES

be suitable in a situation where normalised images are used in a static environment and processing can be done offline. They are not particularly well suited to live (unnormalised) video sequences.

3.7.5 Support Vector Machines (SVMs)

Support Vector Machines is a pattern classification algorithm providing a means of training information used by polynomial classifiers and neural networks (Vapnik, *The Nature of Statistical Learning Theory*, 1995), (Vapnik & Cortes, *Support-Vector Networks*, 1995).

Support Vector Machines are based on the concept of decision planes that define decision boundaries. A decision plane separates classes of objects (Statsoft, 2004). For example; if a decision was being made between dogs (red in Figure 3.7-5) and men (green in Figure 3.7-5), a Golden Retriever being added to the decision plane should be classified as a dog (and represented as red in Figure 3.7-5), Figure 3.7-5 illustrates:




Figure 3.7-5 - Decision Plane example

The above problem separates linear classes (those that can be split by a line) most classifications are more complex than that:




Figure 3.7-6 - Hyperplane example

Figure 3.7-6 depicts a more complex case of classification; the circular dot in white represents a test case requiring classification. It can therefore be classified as red. The process of SVM is the process of mapping the representation in Figure 3.7-6 to that shown in Figure 3.7-5:




Figure 3.7-7 - Hyperplane mapped (using kernels) to Decision plane

Training these classifiers for facial detection requires an enormous amount of training data. The system as described above has been implemented previously for facial detection (Osuna, Freund, & Girosit, 1997). There is no reason why the same approach cannot be applied to the detection of individual features. In the

FACIAL RECOGNITION FROM VIDEO SEQUENCES

face pattern detection scenario, the system is trained using literally thousands of examples (50,000 in the example cited) so that the model can be as accurate as possible. The test image is scanned exhaustively at numerous scales and orientations to locate and match the facial shapes. Extending this to features is trivial; a separate model needs to be built for each of the features and these can be compared to the test patch as it is extracted as part of the search space.

There are obvious issues with this algorithm; firstly, the ‘exhaustive search’ suffers from serious computational efficiency problems, the addition of some weak classifiers would be beneficial but currently, this algorithm is far too inefficient to be used in real-time.

Secondly, the process of actually obtaining the training images is by no means simple. 50,000 noiseless images for each feature being detected is extremely difficult to find and utilise effectively (memory requirements alone would be extreme for this).

Finally, the algorithm itself is unsuitable in some situations; extreme variations in lighting may affect the greyscale images so much that feature matching may become unreliable anyway. In an unnormalised environment (such as the one in acquisition and processing of video sequences) lighting variations are real and difficult to remove effectively.

Conversely, in a normalised environment, SVMs may offer a very accurate means of detecting and recognising faces. However, the process is still likely to be far too slow for anything other than static image analysis.

3.8 Conclusion

A number of current approaches to the extraction of facial features have been proposed. The efficiency of these algorithms will vary. However, a decision had to be made based on the expected efficiencies. The Bayesian approach in section 3.7.1 is simple, supposedly efficient and extremely reliable. It has therefore been chosen as the main feature extraction algorithm to be used.

In addition, the face detector approach also had to be decided upon. Although, the SNoW-based approach discussed in section 3.6.2 offers certain accuracy advantages, the performance gains offered by the weak classifier approach in section 3.6.1 mean that it is the obvious choice.

Chapter 4 – SYSTEM DESIGN

4.1 Introduction

The purpose of this chapter is to illustrate the design process that took place during the development stages of this project. The chapter begins by looking at design specifications and some requirements capturing. It then moves on to covering design methodologies, relevant design diagrams and finally it looks briefly at the way the graphical user interface has been designed.

4.2 System Requirements & Specification

The system specifications are created by considering the user’s and system requirements respectively.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

4.2.1 System Requirements

Number	Requirement
1	The face recognition system should work directly from a live video feed (or buffered feed if the hardware is incapable of supporting real-time recognition)
2	Reference images are normalised before being uploaded to the database
3	System should utilise one or more of the methods mentioned in section 3.7 for feature extraction
4	System should utilise one or more of the methods mentioned in section 3.6 for face detection
5	Reference users should be easily removable from the database
6	GUI should show matched face over set interval (i.e. every 10 seconds the image is updated)
7	Information entered at the GUI should be validated properly
8	Error handling should be elegant and extensive

Table 4.2-1 - System requirements table

4.2.2 User Requirements

Number	Requirement
1	The Graphical User Interface (GUI) should be simplistic to use and intuitive
2	It should allow the user to upload reference images by specifying locations and subject's name
3	The system should respond in real-time when identifying a person (where hardware permits)
4	The GUI should display the live video feed and the closest person match in the same window
5	Errors should be clearly displayed when necessary

Table 4.2-2 - User Requirements Table

4.2.3 System Specification

Number	Specification
1	The Graphical user interface will be written in a high-level programming language that has well founded support for interface components. The Java Swing library or C# visual components would be suitable. Panels nested within frames will be used for the visual layout (JPanel or Panel respectively.)
2	All progress messages and errors will be printed to the console by default. Serious errors will be displayed in a dialog box where necessary
3	Separate java frames with relevant information boxes will be displayed allowing the addition or removal of reference faces. These frames will be accessible from a menu system on the main GUI window.
4	The images available for upload (reference images) can be in any standard format

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Number	Specification
	(jpg, png, gif, bmp)
5	Text boxes will be made available for the entry of the subject's name
6	A list component will be used to display the users in the database for removal

Table 4.2-3 - System Specifications Table

4.3 Face Detector Specification

The choice of face detector is important; it needs to operate in real-time and must accurately locate a face in unnormalised conditions. The choice of face detectors is derived from the comparison in section 3.6.

It must exhibit the following specifications in order for it to be viable in this scenario:

Number	Specification
1	The face detector must run in real-time (hardware dependant)
2	It must accurately locate faces at a variety of scales, rotations and under various lighting conditions
3	The system must be compatible with the high-level programming language of choice (Java or C#)

Table 4.3-1 - Face Detection Specification

4.4 Feature Localisation Specification

The feature localisation algorithms are chosen from the comparisons made in section 3.7, the approach being investigated here for the purposes of specification gathering is detailed in section 3.7.1, a Bayesian approach to feature localisation (Everingham & Zisserman, 2006).

The equations being used here are derived entirely from a recent paper in this field (Everingham & Zisserman, 2006). The specification is split between training and testing.

4.4.1 Data Training Specification

It is essential to the successful operation of a (learning-based) feature detection system, that the learning phase be executed correctly. This is to ensure the appropriate acquisition of the mean feature, covariance matrix and to guarantee that this data is as noiseless as possible. For this particular method, a log-likelihood ratio will be collected. The largest value is then deemed to be the closest match.

Number	Specification
1	Acquire training images from pre-existing database (FERET or OWL)
2	Create mean image from entire collection of training data for each feature (left eye and right eye)
3	Extract covariance matrices and standard deviation from data obtained in step three
4	The Gaussian distribution can be formulated in the following way: $p(x c) = \frac{1}{(2\pi)^{d/2} \Sigma ^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Number	Specification
	(This will be modelled in a programming language for implementation, full details on how this algorithm is implemented are available in section 3.7.1)
5	The distribution in step 4 is compared to potential eye patch and a log-likelihood ratio is used to assign confidence as to whether the patch is an eye or not: $llr(x) = \log p(x e) - \log p(x \bar{e})$
6	The resulting values will be compared to those stored in a database in order to obtain a match. The method will also be used to obtain the features from reference images and those values will then be uploaded into a database for long-term storage.

Table 4.4-1 - Feature extraction training specification

4.4.2 Data Testing Specification

It is the responsibility of the system to compare the features obtained from the live video source and compare the features to those within the reference images⁴.

Number	Specification
1	The test image is to be obtained from a live video source
2	The face detector is run on the image to ascertain the position of the face
3	The facial area is extracted and its size normalised depending on the size of the test image, for this it will be rescaled using language image processing tools
4	A 'sliding window' approach is applied to the test image so that every grouping of patches (size k, where k is the size of the patch of the image obtained from the training section)
5	The log-likelihood ratio mentioned above is used on every patch to ascertain which is most likely to contain an eye image

Table 4.4-2 - Data Testing Specification table

At this point it is expected that the proposed system will utilise face detector (Viola & Jones, 2004) and at least one of the feature recognising techniques, probably the Bayesian approach (Everingham & Zisserman, 2006). The graphical interface will be built separately from these other components and the above specifications give a good indication of how the system is expected to work.

At this point, java was selected as language to be used for the proposed system. This was due to the fact that it was easier to create a wrapper for the OpenCV version of the face detector system being used.

4.5 Description of Design Methodologies

The design work carried out during the development stages of this project consist of a few main areas, namely; the design of the face detection system, the ability to process live video footage and the implementation of at least one of the feature localisation methods discussed in section 3.7. The work is fully structured and is modular in conception; therefore a similar approach is adopted for the design

⁴ It is assumed that the same features have already been extracted from the reference image when it is added to the system as mentioned in earlier chapters.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

methodology. In effect, this means the system is being developed in a structured but modular fashion. Use of DFD's (Data Flow Diagrams), use case diagrams and other visual design methods are used to reduce complexity and illustrate decisions made.

There is an essential step between the illustration of information flows and the inception of classes to support these flows. Therefore, the industry standard methodology for object orientated design is used. This is the Object Orientated Approach (OOA) and it allows better modelling of class interaction.

The final part of this chapter briefly looks at the design of the visual components of the graphical interface so that they conform to the requirements of the users as specified in section 4.2.2.

The design of the system was originally going to be implemented in C# (pronounced c-sharp) because of its close relation to C, the speed at which it processes and the author's familiarity with that language. However, because of the lack of support available for this, relatively new, language and the fact that there have been virtually no attempts made in the computer vision community using C# an alternative was decided upon instead. The language to be used now is Java; it operates independent of platform, has huge online support and has been used extensively in the field of computer vision. It is generally considered to be slower than C# because it is essentially an interpretive language but performance is not expected to be a constraining factor in this implementation.

4.5.1 *Description of Information Flow*

The designs mentioned in this section can be found in section Appendix A. The data flows are depicted at various levels in order to appropriately model the information flows without providing too much information in any single diagram.

There exist two diagrams showing the overall information flow of the system. Figure 8.1-1 shows the overall system process for the face recognition system. It shows that a set of training images are compared to an image acquired from the face detector in an attempt to obtain a match for the face. Figure 8.1-2 forms a sublevel of the above high-level figure. It shows that the face detection system accepts a buffered image as input and returns a cropped picture as a result containing the subject's face only.

Figure 8.1-3 depicts a level 1 representation of the feature training process. It shows that the covariance, mean image and scaling information is derived from a database of training features and stored in a java data store for use in the main program. These values are then applied to the main program along with the live face detector in order to attempt a facial match. Sub processes exist for these steps in two separate forms; training and testing.

The training stage is shown in more detail on the level 1.1 diagram in Figure 8.1-4. This diagram shows that the greyscale images are derived from the chosen database (most likely FERET) and stored in .dat files for the java program to read. The exact process for the derivation of these values is shown in Figure 8.1-5. This diagram is a level 1.1.1 data flow and shows the exact process undertaken for the extraction of the mean, scaling and covariance values needed for the matching process. In short; the patches are taken one at a time and the feature (eye, mouth etc) is extracted. That feature is then normalised and reshaped to a vector (rather than a matrix). The covariance is determined and then inverted to create the inverse

FACIAL RECOGNITION FROM VIDEO SEQUENCES

covariance matrix, which in-turn, is saved to a .dat file. The scaling value is determined and saved and the mean feature is extracted. The mean feature is centred and reshaped as a vector before being saved to secondary storage.

As mentioned above, the second main process is the testing phase of the feature comparison. The level 1 diagram depicted in Figure 8.1-3 and mentioned above also describes the process 'feature detection testing'. This sub process is extrapolated in the level 1.2 diagram shown in Figure 8.1-6. This information flow diagram shows how the test face (as acquired from the face detector output) is systematically searched for the feature as trained during the training phase. The algorithm should take the size of the patch (size k , as determined during the training phase) as input and then systematically examine every set of pixels (size k) in that test image and compare the subset of pixels against the mean values obtained during the training stages to find a match. The actual matching process, once the measurements have been taken, is described in more depth in the diagram in Figure 8.1-7. This is a level 1.2.1 diagram and shows how the location of the features in the test image is compared to those in the reference images. The Euclidean calculation is shown in Figure 8.1-8, a level 1.2.1.1 diagram subtracting the test image feature location from each of the reference images. If the result of this calculation is zero the people match. The further from zero the result, the less likely that the two faces match.

4.5.2 *Description of State Transition*

The state transition can be found in section 8.1.9; Appendix A, it depicts the different states that the face recognition system can enter and aids in the design of the 'multiple frames' detection filter.

The system begins in an initial state, i.e. before the system has started. On initial execution, the system will begin to train the images (if it can) and begin the facial detection process. The training state on the diagram is distinct from the other states and is set off manually. It is usually run once and the data retained but can be retrained if necessary.

The actual face recognition process is executed in a loop. The system waits for the frames from the video source and extracts them one-at-a-time in a loop, prior to this, the system must have the data from the feature training in order to continue. The system must then await extracted faces and process them one-by-one in order to locate the features for matching. Failure to find a face or recognise a face accurately simply sends the system back to the beginning where it begins the process again. Final state here refers to the successful localisation of features prior to resetting the system.

4.5.3 *Object Orientated Analysis (OOA)*

A complete class diagram for the face detection system (written in java) is supplied in section 8.1.10, Appendix A. This diagram is intended to show the relationship between the various classes implemented within the system.

It was decided, using a combination of techniques (the design methodologies in section 8.1) to implement the system according to object orientated best practices. In order to do this it was necessary to create 14 classes excluding test classes and inner classes (those classes that are defined privately within other classes.)

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Justification for the use of the aforementioned classes is in the next section:

4.5.4 Justification of class choices

FaceRecogGUI class: This class represents the main interface for the face recognition system. All messages are passed via this class and this class collects panels from the other classes and displays them in the correct order in a frame.

ImageDisplay class: This class provides the main ‘back-end’ class for the GUI mentioned above. It liaises with the face detection and feature detection classes and does the actual matching before passing the result to the GUI for display.

FaceLocations class: This class creates an arraylist object and stores the locations of faces. It takes the face from the face detector class and processes the result so that the x and y coordinates of the face is stored here.

FrameAccess class: Provides the face locations class with individual faces from individual frames.

FaceDetection class: The class provides the *FrameAccess* class (above) with the face locations on each frame (image) collected from that class.

FaceCrop class: Takes the individual frame from the *FrameAccess* class and combines it with the face location from the face detector class in order to crop the image leaving only the face to return.

FaceLocation class: Basic storage class containing accessors and mutators for the locations of the face; x, y and width.

Face class: stores individual information about the face. The subject’s name, x and y locations of the eyes and cropped facial image are held here for the purposes of the reference face and the comparison are all stored here.

FeatureLocalisation class: This class attempts to localise the facial features using the methods investigated by this report.

EyePatch class: describes an individual eye patch so that the eye patches can be stored in an array and easily compared to find the relevant features.

Matrix class: A class providing useful Matrix functions that can be used by the feature finding process.

FaceDataAccessLayer class: This class provides all the program’s database reading and writing functions in one place. All the relevant information is sent to this class for writing to the database.

FaceAddToDatabase class: This class is an interface; it provides a separate GUI to upload a reference picture containing a face and some attributes describing that image to a database for storage.

FaceDeleteFromDatabase class: This is similar to the above *FaceAddToDatabase* class, it provides a separate GUI and a list collected to the database for easy removal of reference faces from the database.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

4.5.5 Explanation of message passing (Figure 0-1)

- (1) FaceDataAccessLayer → FaceAddToDatabase/FaceDeleteFromDatabase → FaceRecogGUI
This message flow represents the error message flow from the database back to the main GUI in the event of a failed call to the database
- (2) FaceRecogGUI → FaceAddToDatabase/FaceDeleteFromDatabase
This message represents the invocation of the classes at the receiving end of the call. A second GUI is created in response to allow editions to the database to be made.
- (3) FaceRecogGUI → ImageDisplay
This message is designed to invoke the imageDisplay class and begin the process of video frame retrieval and feature localisation.
- (4) ImageDisplay → FeatureLocalisation
Once the imageDisplay class has been invoked as per step number 4, the featureLocalisation class is called and the face image is passed to it so that the features can be found and saved to secondary storage.
- (5) ImageDisplay → FaceLocations → FrameAccess
This message call represents the request for individual frames from the video source. It then shows the the faces are stored in the FaceLocations class before being passed back to the ImageDisplay class for further processing

4.5.6 Graphical User Interface

The design of the user interface takes into account the user and system requirements mentioned in chapter 1. The GUI (Graphical user interface) must allow the easy matching of features and the easy addition and removal of reference faces from the program.

The main GUI window, when executed, begins the face recognition process automatically for whatever object is directly in front of the live video source (webcam). It can be seen in Figure 8.2-1 in **Error! Reference source not found.**

From the main window mentioned in the above paragraph, there is a menu system (file, in the top left of the screen) that allows the addition and removal of reference faces from the system database. The windows to support these functions can be seen in Figure 8.2-2 and Figure 8.2-3 respectively. The add screen allows the upload of a reference face image and a name to match that image to the database, using standard windows 'browse' functionality. The 'delete reference face from database' window provides a list of names from the database and the option to select a name and click ok to confirm deletion.

As you can see, the GUI windows have been designed to allow maximum functionality with minimal components. The Main GUI window has a live video feed source on the left hand side of the screen and a space on the right for an image to be displayed with a name to denote the matched reference face.

Chapter 5 – SYSTEM IMPLEMENTATION

5.1 Introduction

This chapter describes how the face recognition was implemented. It explains specifically how the equations for feature detection were implemented, how the video component was added and how the database interface was made. The chapter details differences from the intended design in Chapter .

5.2 System Implementation

The implementation was originally planned to follow the design as laid out in Chapter . As expected, some problems were encountered during this work and were dealt with in a way that varied from the design slightly. Those areas are explained in subsequent sections, whilst this section is reserved for the explanation of how the complex problems have been implemented.

5.2.1 Feature Detection Training (Eye)

The basis of this face recognition system lies, in part, on the ability of the feature detection algorithm to accurately localise facial features. Although this part of the study could conceivably have been written in Java, the built in functions available in Matlab make the process more streamlined. Therefore, a decision was made early in the development of this system to utilise Matlab for the training portion of the prototype. The algorithm developed here is based on the training part of the Bayesian approach to feature localisation examined in section 3.7.1 (Everingham & Zisserman, 2006).

The Matlab file prepared for this is shown in full in section 8.3.1 of Appendix C and is shown in pseudo-code below in Figure 5.2-1:

```
Create array (rightEye[]) to hold eye patches
For each eye (i) in FERET database
    Extract right eye patch
    Normalise eye (size and greyscale)
    Reshape eye from Matrix to column Vector
    Add eye patch to rightEye[] array
Extract the mean (M) eye image from the rightEye[] array
Centre the data obtained about the mean
Compute the covariance matrix (S)
Create Gaussian distribution
Save scaled mean (M) to file
Save inverse covariance matrix (Sinv) to file
Save standard deviation value (s) to file
```

Figure 5.2-1- Training Data calculations pseudocode

This function has therefore output the mean, covariance matrix and standard deviation to .dat files for use in the java program when it is run.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

5.2.2 Data Storage/Manipulation Classes

A number of classes have been written in order to store data in or do basic data manipulations. These are listed here with brief descriptions if appropriate, especially where there are differences from the original class diagram:

- Face.java class (Figure 8.3-2, Appendix C);
 - A class used to store information about the reference faces held within the database for matching purposes.
- FaceCrop.java class (Figure 8.3-3, Appendix C);
 - This class accepts a buffered image and the size of a rectangle as input and returns the sub-image of that image cropped according to the size and location of the rectangle. It can also accept a regular image object and crop that appropriately. This is designed for static-image testing of the face detection and feature localisation process.
- EyePatch.java class (Figure 8.3-4, Appendix C);
 - This class was not present in the original design. Its purpose is to store information about each of the patches investigated during the feature localising process. It stores the x and y position of the bottom-left corner of the patch, the size of the patch, an index and the likelihood that that patch contains the feature being searched for. This is for easy reference later.
- Matrix.java class (Figure 8.3-5, Appendix C);
 - This class is another class added since the initial design. It became apparent that a number of matrix calculations would need to be made by a number of classes (especially if this system was going to be extended). Therefore, this simple class provides a number of matrix multiplication, subtraction and column and vector calculations on behalf of other classes.
- OpenFileAction.java class (Figure 8.3-6, Appendix C);
 - The OpenFileAction class provides a GUI to open files from the file store and hence obtain the complete file path in a visual way. It offers this capability to other classes as they need it. Although this class is not shown in the original class diagram, its presence here simply reduces the time taken to write the code multiple times and does not reflect poor design technique.
- FaceLocation.java class (Figure 8.3-7, Appendix C);
 - Another class designed to store information about particular features. This class stores the location (x and y) and rectangle size of a face in relation to a particular image.
- FaceLocations.java class (Figure 8.3-8, Appendix C);
 - This class is a logical extension of the FaceLocation class mentioned above. It implements the face detector class and retrieves the faces (possibly multiple) from that class. It then creates a number of FaceLocation objects to represent those faces and stores them in an array for further processing.

There have been some deviations from the original design, as can be observed above. A new class diagram with all of the updated information will be made available at the end of this section.

5.2.3 Database Manipulation Classes

There exist a number of classes that are used to interface with the database. The purpose of these classes is to add and remove reference faces. The actual database integration is handled by a single data access layer class (in line with good programming practice (Aarsten, Brugali, & Menga, 1996).) This data access layer class, `FaceDataAccessLayer.java` handles all reading and writing to and from the database. Classes needing to use the database can simply call methods in the data access layer in order to facilitate that access. The class can be seen in Figure 8.3-9 of Appendix C.

The data access layer class above essentially offers database access to two of the other classes. Because the database is used in this context as persistent storage for the reference faces, the classes needing to access the database are those which add and remove reference faces to and from the database respectively. The classes in question here are the `FaceAddToDatabase.java` and `FaceRemoveFromDatabase.java` classes; both classes provide graphical interfaces for the addition and removal of reference faces and the database. The source code for these classes can be found in Figure 8.3-10 and Figure 8.3-11 of Appendix C respectively.

5.2.4 Video handling class

It became apparent during the design phase of this project that a class would be needed that would not only have the necessary codecs to be able to decode video files, but also have the ability to interact with live video sources and extract single frames of the videos and treat them as images. The full source code can be found in section 8.3.12 of Appendix C. In short, the class is responsible for decoding a video source so that it can be used and providing the individual frames as images for further processing.

Unfortunately, java offers no native video handling capabilities. The consequence of this is that a prerequisite condition is necessarily added to the entire program. The program must be made to use the JMF (Java Media Framework⁵), it is therefore necessary to install this prior to executing the program.

The class sets up a media player with individual frame access and uses the Windows native webcam as its source. The program utilises a number of codecs to successfully decode video sequences. The program uses JMF standard media control panels to support the video stream (pause, examine etc.) and uses a Media locator object to track the video source. The media stream is created, the frame accessor is assigned to listen to this stream, the controller components are added to the video display and that entire display is attached to a panel in order to be offered to the main GUI window.

5.2.5 Face Detection class

The face detection class has a simple job; it examines images and returns the locations of the faces (x and y position of the top-left hand corner of the face and the width of the square rectangle enclosing the face.) The class creates an object of class `JNIOpenCV`⁶; this class is stored in a .dll file and implements the OpenCV version of the face detector (Viola & Jones, 2004). It is, in effect, a java wrapper of a C program. The Face detection class acts as an interface for this wrapper and has been designed so that it

⁵ The Java Media Framework is an open sourced development project aimed at enabling media support for java applications.

⁶ Provided by Dr. Hu software organisation (Dr. Hu Software, 2002)

FACIAL RECOGNITION FROM VIDEO SEQUENCES

accepts an image and returns a list of face positions for processing. The next step is to pass these values to the face cropping algorithm to extract just the face. The full java source code can be found in Figure 8.3-13 of Appendix C.

5.2.6 Overview of main program functionality

The above classes provide storage, Graphical interfaces and low-level data manipulation tasks. The main program core is very much more complicated than that. From a high level perspective; the program accesses individual frames from a webcam and adds those frames to a queue. In the meantime, a separate thread of execution constantly processes the frames in the queue; detecting the faces and localising the features. The dual threads of execution are shown in pictorial form in Figure 0-2 of Appendix A. It is necessary that this functionality of the program be provided by a separate thread because accessing the work queue must be done on a constant basis for the duration of the overall program execution. In order to implement, there were two options:

- Create a timer for the program to periodically poll the queue for new frames to access
 - In theory this approach works quite well but in practice, without dynamically updating the time interval at which the queue is checked, the queue will quickly become overloaded. Especially given that the items in the queue are likely to be processed significantly slower than they are added. Thus, an unnecessary overhead is added to the queue.
- Create an infinite loop where the queue is constantly accessed and the frames are processed.
 - This invariably involves the creation of an infinite loop where the program continuously checks the queue for the presence of additional work to complete. In the event that the queue is empty, the loop will wait for an item to be added. As mentioned above, in order to prevent the program entering an infinite loop and then not completing the other processing (the extraction of individual frames), this loop would be implemented in a separate thread of execution and then both parts of the program (frame extraction and frame processing) would continue concurrently⁷.

Once the frames of video have been extracted, the feature localiser works by extracting sub-images from the test face (patches). Then comparing these features to mean features already obtained during training. Every possible patch is examined by iterating over all possible x and y locations in the image. The most likely patch location is returned to the main program for comparison; see Figure 0-3, Appendix A. The actual sliding window approach is therefore critical to the successful operation of the entire system. In pseudocode it performs the following:

```
for y = height/2 to y = height - patch size //Look at top section of face only
  for x = 0 to x = width - patch size // look at each row in turn
    Read sub-window as buffered image
    Convert sub-window to greyscale
    Convert sub window from matrix to vector
    Compare test vector to training mean feature image
```

⁷ It is possible with modern processors to set processor affinity to specific threads so that the threads would be executed in a truly concurrent manner.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Store log-likelihood result for later comparison

Figure 5.2-2 - Pseudocode for sliding window approach

In the Bayesian approach discussed in depth in section 3.7.1, probabilistic models of the appearance of a patch are built independently for positive and negative classes.

The log-likelihood output of these probability models is then used to ascertain whether a patch is an eye or is not, according to the aforementioned probability models. Simply, the patch with the greatest log-likelihood ratio is deemed to contain the eye. Obviously, there should be two high-valued log-likelihood ratios for each of the eyes in every image. The two highest values of log-likelihood should both be obvious and accurate detect the position of the two eyes. In practice, it may well prove to be better to split the training images between left and right eyes in order to obtain slightly better classification during the test phase.

The log-likelihood value can be created using the following formulae:

$$llr(x) = \log p(x|e) - \log p(x|\bar{e})$$

Equation 5.2-1 - Log-likelihood comparison

That is to say that the likelihood that patch x is an eye; llr (log-likelihood ratio) is given by the logarithm of x given that it is an eye minus the logarithm of x given that it is not an eye. This is the logarithm of the differences.

The complete Gaussian distribution for each eye and non-eye can be shown as such (where each class c (eye and non-eye) is modelled as such):

$$p(x|c) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$$

Equation 5.2-2 - Bayesian mean image building equation

d is the size of the patch ($k \times k$), x is the test patch to be classified (eye or non-eye), μ is the mean image represented as a vector and Σ is the covariance matrix. They are shown in the following way:

$$\tilde{\mu} = \frac{1}{N} \sum_{i=1}^N x_i$$

Equation 5.2-3 - Bayesian mean image

$$\tilde{\Sigma} = \frac{1}{N-1} \sum_{i=1}^N (x_i - \tilde{\mu})(x_i - \tilde{\mu})^T$$

Equation 5.2-4 - Bayesian Covariance building equation

There is no particular reason why this formula for the mean is used but it is as suitable as any other. The estimate of the covariance must be regularised since $d(d+1)/2$ parameters must be estimated. This is done by adding an identity matrix to the covariance matrix:

$$\Sigma = \tilde{\Sigma} + \gamma I$$

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Equation 5.2-5 - Bayesian addition of regularisation parameter equation

The above pseudo and mathematical equations have been implemented in java in the FeatureLocalisation.java class shown in section 8.3.14.

Unfortunately due to reasons discussed in the following chapter, this implementation failed in this particular instance. In short, because the complex matrix calculations take far too much time to complete, the queue being fed from the frame extractor class becomes overloaded very quickly. Therefore, although the training algorithm remained the same, Equation 5.2-1 was changed to be in the line with the pseudocode presented Figure 5.2-3. More details are given below.

As a work-around to this problem, the above implementation was changed. The mean feature was extracted in much the same way but is then compared directly to the test patch (extracted as part of the sliding window method) using a sum-of-squared difference approach. This has the unfortunate side-effect of being less accurate because the standard deviation and covariances are not being utilised. However, the process is fast enough to prevent system degradation within the first few iterations. The original class listing can be found in Figure 8.3-14. The change mentioned can be shown in java by replacing lines 370-402 with the following section of code:

```
double sumSquaredRight = Matrix.sumSquaredDiffVecs(meanEyeArrayRight, testEyePatchVector);
double sumSquaredLeft =
Matrix.sumSquaredDiffVecs(meanEyeArrayLeft, testEyePatchVector);
EyePatch thisPatchLeft = new
EyePatch(patchIndexCounter,x,y,sumSquaredRight,patchSize);
EyePatch thisPatchRight= new
EyePatch(patchIndexCounter,x,y,sumSquaredLeft,patchSize);
```

Figure 5.2-3 - Sum-of-squares difference java source code

As you can see, the calculation is actually performed by the Matrix class as before. It also handles this functionality.

5.2.7 Individual frame decoding

The frame decoding is handled by a specific class. This class serves two functions. It provides a ‘wrapper’ to the core functionality of the Java Media Framework⁸, used here to process the live video feed. It attaches the video as well as some ‘media player’ controls to a panel and prepares to return that to the main GUI in the form of a panel. In addition to showing the video, it also accesses each individual frame as it is decoded and converts these frames to BufferedImages, before adding sequentially to the cross-threaded work queue discussed earlier and shown in Figure 8.3-16, the frame decoding class is shown in Figure 8.3-15.

The interface used here supports a relatively high frame rate, one which is much higher than the other component (the feature localisation component) can possibly hope to support. Therefore, during the frame access process, a throttling process is added to prevent too many frames being passed to the queue and the program running out of memory. This is provided by simple sending the thread to sleep for a quarter of a second at every iteration. This has the effect of reducing the frame rate from 15 frames per second to 4

⁸ Code modified from that provided as an example interface to JMF by Sun Microsystems

FACIAL RECOGNITION FROM VIDEO SEQUENCES

frames per second. This unit may be changed during testing because it should, in effect, force the process of video frame extraction to take the same time as feature localisation and therefore ensures that the program will not run out of memory even if run for long periods of time.

5.2.8 *Frame handling functionality*

The frame handling class is different from that originally proposed during the system design phase. Its inception, as mentioned above was due to the fact that the part of the program it is responsible for (the feature localisation within decoded video frames) needed to be placed in its own thread so it didn't interfere with the other thread of execution. In short, therefore, its function is to access the work queue on a continuous basis, locate the face, and locate the features and return. The results of this 'matching scheme' are then placed on a frame within the main GUI window as results. The matching scheme referred to here is the process of comparing the located features on the test face with those features held in a database and obtaining the result. The faces in the database are loaded into an array at the beginning of the class as Face objects (section 5.2.2). The face object is updated statically by the class and is polled on a regular basis by the main GUI class to see if a different face has been matched and then update the GUI accordingly. The face coordinates, feature locations, subject's name and the reference image are held within the face object. This class, *FrameHandler.java*, replaces the functionality planned for use in *ImageDisplay.java*. *ImageDisplay.java* handled static images but was too awkward to change for use as a threaded process but still serves a useful purpose as a test class for the feature localisation process with regard to static images because debugging that process during the live video testing is extremely problematic. In order to fulfil the requirements of this project, inside the infinite loop mentioned above, where the frames are retrieved from the video source on a constant basis and processed for matching; a second loop exists inside that loop in order to take advantage of a number of successive frames for the matching scheme. In effect, a number of frames are processed and the features are located, the features are then loaded into arrays as part of this loop and at the end of the loop (after 5 or 10 frames have been processed) the values in the arrays are averaged (to get the average feature distance) in an attempt to remove 'noisy' data. The matching process then takes place and the main interface's matching scheme is updated to show the matched person.

The final class diagram for the implemented system can be found in Figure 8.2-1 of Appendix A. The final database schema for face storage is also shown in Figure 0-2, and the entire code listing can be found in Appendix C.

Chapter 6 – SYSTEM EVALUATION

6.1 Introduction

The chapter looks at the different types of testing that have been performed on the programme in order to ascertain its reliability. The purpose of the testing phase is to establish that the code is reliable and that all the functionality has been tested. This first part of the chapter is designed to give a brief overview of the testing undertaken. The chapter then looks at how well the solution was implemented and areas that require further work.

6.2 System Testing

The main method of testing throughout the design of the project was to establish that each individual class was performing as expected. In order to facilitate this, each class implements its own main method and can be executed independently of the rest of the system. This is to ensure that each component works before tying them together. In addition, a number of print statements have been prepared at different points throughout the program to help locate failures when the program had been integrated and the complex threads of execution were difficult to follow. Lines 61-66 of Figure 8.3-13 show how a main method can be used to ensure that the program works as expected, the print statements give the relevant information about whether the execution was successful or not. Some print statements showing useful information about program execution can be seen in lines 279-332 of Figure 8.3-12. This class also has a test main method which isn't utilised during the execution of the main program, shown on lines 226-257. The purpose of this test is to ensure that the video sequence is working correctly and that the frames are being accessed (shown by print statements on the console.)

In addition to main methods being created, a number of test classes were created to ensure correct operation prior to the final compilation of the code. ImageDisplay.java was originally planned for use in the main program; however, the problem of threads mentioned earlier meant that this class couldn't be easily modified for use with the rest of the system because its functionality was required for testing. Its function is to take a static image, run the face detector, outline the face with a rectangle and then apply the same process to feature localisation. It was therefore used to debug the face detection and feature localisation classes. The source code can be found in Figure 8.3-18 of section 8.3.18. Other similar classes were used for the testing of the database and updating the image display, shown in Figure 8.3-19 and Figure 8.3-20 respectively. These classes were created to ensure the database imaging was being created correctly by saving to the database and then retrieving it again and to display the image on a panel in a dynamic way before using it in the main entry class.

6.2.1 System Black Box/ Functionality Testing

In order to ensure the proper function of the main interface once running, black box testing was performed on the application. A number of test conditions were created to check abnormal behaviour. The test steps can be found in Table 8.4-1 of Appendix D. The screen shots associated with this testing can also be found in Appendix D and Appendix B.

The testing included; verification of error messages displayed to end users when reference information held in the filestore was missing and ensuring that the system met the requirements.

6.2.2 Training system testing

The testing of the training system was a fairly informal process. The solution was developed in Matlab because of its advanced computer graphics programming abilities. To ensure that the eyes were the features being extracted, a number of them were displayed at random as part of the loop. Because the solution was using images from the FERET database, this was simply to ensure that the coordinates within the eyes arrays represented the centre of those eyes. It was accepted that the data was noisy (the coordinates didn't necessarily represent the exact centre of the eye) but this was unlikely to cause

FACIAL RECOGNITION FROM VIDEO SEQUENCES

problems with the overall feature representation process because there were over three thousand features per feature.

The information was displayed on the console during each phase in the training process (see Figure 8.3-1 for code) to ensure that the matrices were the appropriate sizes and that the information contained within those matrices seemed appropriate.

The overall accuracy of this system, especially when the above data was used to localise and display eye positions in Figure 8.3-18 was used to check the integrity of the training algorithm.

6.2.3 *System Accuracy Testing*

In order to test the system, 10 reference faces were loaded into the system. This was deemed enough to test the accuracy without hindering performance. Serious degradation in performance is likely to be seen when comparing upwards of 100 faces because the time to iterate over them in an arraylist (java object array storage) is constant and therefore increases with increased load.

The system, as per the requirements, utilises a number of frames in order to remove erroneous feature detection from the matching scheme before that matching takes place. It is expected that the number of frames will have a positive effect on the matching performance of the system.

The results are in fact quite poor. They are shown in Figure 8.4-7 of Appendix D. As you can see from the graph, the correct identification of a face occurs in between 10 and 25 percent of recognition attempts. Although, this may seem poor, the results do show that accuracy is increased, on average, as more frames are used to identify the face. This, in part at least, proves that using multiple frames of video to identify a face is beneficial even though the localisation by the eyes alone is apparently inadequate.

6.2.4 *System Requirements Testing*

The system has successfully implemented all of the requirements gathered in sections 4.2, 4.3 and 4.4. The results of these tests can be seen in brief in section 8.4.2 of Appendix D.

6.2.5 *Test Summary*

The testing process in this section (6.2) is split in to three main parts:

The first part tests the user interface and checks that all of the expected functionality is in place and that all of the potential problems with the interface are dealt with elegantly.

Then the 'black box' testing checked the system to ensure that unexpected behaviour was dealt with effectively and that suitable responses from the system are issued to users so that they can attempt to fix the problems.

The final part of the testing centres on system accuracy testing. This highlights the poor results obtained from using the eyes alone for facial identification. However, the results also showed the added benefits obtained from using a fusion of video frames to obtain an average feature match.

6.3 System Evaluation

The report was written alongside the development of the system for the most part. It has been split into stages to mimic that of a real software development project; a background study of the approaches commonly employed, a formal system design and a system implementation of the approaches mentioned.

The design phase was based on work carried out by software development projects and took a number of weeks to complete fully.

The implementation phase was expected to be extremely difficult; a number of problems that could not be solved during the design phase were necessarily encountered here. These problems were likely to be centred on the performance of the Java virtual machine when dealing with facial detection and the application of the chosen feature detection algorithm (Everingham & Zisserman, 2006). However, other significant problems presented themselves; the necessity of deploying a multi-threaded application and Java's lack of support for any kind of video were the most significant. One of the aforementioned possible problems also presented itself when implemented. The feature detection algorithm proposed was extremely effective in its work, as determined during static-image accuracy testing. However, when this was combined with the video reading capability, it became apparent that the complex matrix calculations required to get the algorithm working would make the system impossibly slow to use in real-time. This was realised through the fact that the work queue quickly became completely overwhelmed. Instead, a much simpler approach was developed; the mean data extracted from the training process was retained but the covariance information and standard deviation had to be excluded. The test vectors were created using the sliding window approach as per the design but were then compared directly to the mean feature vectors using a sub of squared difference approach. This had the effect of making the system viable in real-time but at the expense of the accuracy. The margin of error already exhibited by the original algorithm was then exasperated and had an effect on the overall matching scheme.

It was clear from the design that it would not be possible to iterate over the database table for each and every feature match; in order to address this, all reference faces are retrieved from the database and stored in an 'in-memory' array for quick and easy access. However, iterating over this array itself was only possible in constant time and therefore it would not be possible to overload the database with reference faces without rethinking this approach. A weak classifier used in the face detection part of the program (Figure 3.6-2 on page 17), could be used here in an attempt to alleviate this issue.

An unforeseen problem was also present with the face detection .dll file used to detect facial patterns in images; it was assumed as part of the design that this library could take an image path as input for the purposes of testing the system with static images but then take a memory-buffered image directly when the approach was integrated with video. This, unfortunately, was not possible and a work-around had to be formulated; this involved temporarily storing the image on disk after being retrieved from the frame access class. The path would then be passed to the face detector, face detection performed and the image destroyed. This could potentially have slowed the system massively, but the actual effect is unknown.

The testing phase was completed in conjunction with the implementation phase. This involved a modular testing approach where each class was tested alone, then tested with the classes that use it directly. The

FACIAL RECOGNITION FROM VIDEO SEQUENCES

frame extraction class was built independently of the feature extraction classes so that those classes could be appropriately tested with static images before the video component was added to the main interface class. The testing, as mentioned in section 6.2, involved main methods for most of the classes and print statements to show the flow of execution.

After this was completed, black-box testing was executed on the interfaces to ensure appropriate (and elegant) operation in the event of failure. These results can be found in Table 8.4-1 of Appendix D and discussed in section 6.2.1.

6.3.1 *Adopted approach*

The approach adopted is based on the investigation work adopted in sections 3.6 and 3.7. The face detector employed is based on the current industry standard (Viola & Jones, 2004). Tests show this worked as expected, even though it appears to curtail some of the images more than others when detecting them and potentially caused matching problems when the image was rescaled; the rescaling involves changing the facial dimensions to ensure that they are the same size as the training images (100 x 100 pixels in this case.) The feature extraction process is based on a currently technologically advanced Bayesian method (Everingham & Zisserman, 2006), but heavily modified for performance purposes.

The language chosen for the implementation was Java. This was against preference of the author but because the support for images in Java is so extensive on the Internet and many other Vision systems have been developed in this language to prove its viability, it seemed like a good choice. Its performance limitations were noted but thought to be non-critical.

6.3.2 *Outcome*

The approach was adopted partially based on a recent paper citing eye localisation as both important, and the only feature necessary for fast, reliable face recognition (Hjelmås & Wroldsen, 1999). Given that the system needed to accurately recognise faces extremely quickly (real-time) and also that the system was designed as a prototype, this seemed an appropriate approach to follow.

The results shown in section 6.2.3 do not support this claim, although they do clearly show that using a ‘fusion’ of videos frames is extremely helpful to accuracy. This means that one of the main hypotheses of the report was accurate.

With regard to the initial report objectives, the following have been met:

- The initial objective was to gain a solid understanding of some of the more recent advancements in facial recognition techniques and employ some of those in the prototype; it is believed that this has been successfully met by the work carried out in Chapter 3 and the Feature localisation class implemented in Figure 8.3-14.
 - It is conceded that the feature localisation class, as mentioned above, was insufficient to meet the requirement of a reliable face recognition program. Another approach from section 3.7 could have been applied to another of the features; section 3.7.3 details a simple approach that uses a combination of two-dimensional shapes, in varying sizes and

FACIAL RECOGNITION FROM VIDEO SEQUENCES

line curvatures to find features (eyes, mouths, eyebrows, noses etc.) This approach is recommended because of its simplicity and speed of processing.

- In addition, the Bayesian approach used for eye localisation could be easily extended to locate other features without modifying the training process at all. In fact, the testing process (in the sliding window method) would only need to reference the new arrays created from the new features. Therefore, the addition of these features would require almost no extra processing in order to implement. And crucially, no extra mathematics or complex programming techniques would need to be developed.
- The second objective was to gain a better understanding of the chosen high-level language used to implement this system. Although there is no previous basis from which to compare the author's programming knowledge, it is believed that this has been demonstrated by the size and complexity of some of the techniques employed here; low-level graphical processing, library interfacing, database connectivity, video handling and multiple thread processing are some of the more complicated techniques not previously known.
 - The designs and implementation were undertaken according to best practices and therefore it is believed that any complexity encountered was simply overcome rather than avoided. The system was also designed according to object orientated design methodologies to produce the most appropriate object orientated solution to this problem. However, several improvements could be proposed; the use of processor affinity set for specific threads would be beneficial to performance and would effectively give the separate threads a truly concurrent programming environment. In addition, it would be worthwhile spending some time auditing the source code with regard to performance. Some of the methods could be overly complex and inefficient and these should be addressed.
- The final object was to develop a complex system for real-time facial recognition. The system implements the face detection, feature localisation and recognition from a 'fusion' of video frames. It does this in real-time or approximately 4 frames per second to be exact.
 - The Bayesian approach should be extended to cover more features as mentioned above. Also, some time should have been devoted to a more accurate face detection system because scaling the shape obtained is unpredictable if though the detector does usually locate most of the face.

Finally, the database component was not part of the original design but was thought to be a suitable medium to store the face features if they were subsequently accessed in an efficient manner. It is worth noting that the database schema is not in 'normalised' form and this should be addressed in a future version. The eyes, along with other features, should be split from the main table so multiple features can be stored without over complicating the main table.

Displaying the live video in the main interface window alongside the matched face and the menu structure offering access to database management windows was also not part of the original requirements.

Finally, the fact that the Bayesian eye localisation algorithm failed to perform well enough to be used in a real-time environment meant that it would have to be replaced by a more efficient approach. The simplest,

FACIAL RECOGNITION FROM VIDEO SEQUENCES

and therefore quickest, solution to this problem appeared to be a sum of squared difference calculation. Although, this meant a decrease in accuracy, the increase in performance countered this and made it possible to meet the real-time processing objective.

When comparing the project outcome to the user and system requirements, the following can be seen:

-

Overall, the system performs relatively well given the complex programming needed to get the system working. The feature localisation should be extrapolated to take account of more features and this should make the system more reliable. The face detection algorithm also requires some tuning as mentioned above. Finally, the system becomes unstable after a couple of minutes of running. This is because the frame access class constantly adds frames to the work queue, apparently, slightly faster than the frame handler can deal with them. Therefore, the java virtual machine slowly runs out of memory. Efficiency should be sought in the algorithms developed to make the frame handler faster or the frame rate should be throttled even more than it has been already to prevent this queuing problem. In the event that both of these fail, the system should be pursued in a faster language such as C or C++.

In order to apply a more formal evaluation approach (Boloix & Robillard, 1995), a software system evaluation framework has been applied. This is used to compare the objectives and expected use of the system with the outcome. The evaluation can be seen in Figure 8.4-8 of Appendix D (page 118). The evaluation highlights the following:

These areas of application importance can be categorised as ‘advanced’:

- Tools, Products, Technology and contribution;
 - The Tools, Products and Technology refer to the frameworks that the system runs on and how they are implemented within the application. The Java Media Framework and threaded execution can be considered advanced and fulfil the requirements of the project appropriately.
 - The contribution refers to how well the overall system reflects the objectives with regard to the end users. This is considered advanced because all of the components are available in a single place and offer complex functionality.

These areas of application importance can be categorised as ‘intermediate’:

- Process followed and system usability;
 - Although the process followed is based on solid experience and examples, the problems with real-time efficiency and lack of features used overall results in the system being relatively ineffective. It is categorised as intermediate because of the excellent basis for further work that is provided here.
 - System usability is simple; but no heuristic studies have taken place to confirm this and some components are missing (ability to change the file stores, change the .dat files used and general system clean-up activities etc.) The result is a system which is functional but no really ‘usable’.

These areas of application importance can be categorised as ‘basic’:

- The performance and the system compliance
 - The performance is sub-par; partially because of the ineffectiveness of Java and partially because of the inefficient algorithms being utilised. This needs some work in order to correct.
 - The compliance is considered basic because of the lack of accessibility features available through the interface and the complex frameworks that the system requires to even run. This coupled with the complexity of changing the video source and the requirement of an SQL Server 2005 database server results in poor compliance.

Chapter 7 – STUDY CONCLUSION AND AREAS OF FURTHER INVESTIGATION

7.1 Introduction

The purpose of this chapter is to briefly conclude the learning outcomes of the project and then address any relevant areas of further investigation that are recommended for future focus of study direction.

7.2 Study Conclusions

It has been the focus of this report to study current state-of-the-art facial recognition techniques and apply those techniques to a formal design, implementation and testing process to gain a better understanding of their relevance in the field.

The project was successful to a degree (discussed in section 6.3.2 above). It successfully implemented the entire user and system requirements deemed necessary and implemented a number of complex programming methods in order to support the objectives. In order to fully satisfy the requirement that the system should provide; that is robust, reliable face recognition in real-time, it is conceded that it would be necessary to extend the programming done to include other features for more accurate matching. It may also be beneficial to provide another feature localisation algorithm to provide faster localisation than the algorithm discounted as part of the system implementation (Everingham & Zisserman, 2006).

The Bayesian feature localisation solution; training and testing can be considered extremely successful. This algorithm is easy to train and relatively easy to apply. It also offers excellent results. If it weren't for the poor performance that this algorithm appears to offer (at least in Java), it would have made a valuable contribution to this project. However, it is worth pointing out that this approach has a margin for error of approximately ± 2 pixels for each localised eye position, which means a potential 8 pixels of error between both eyes. This, in turn, equates to approximately 8 pixels of total error with regards to the Euclidean distance between the eyes overall. Given this, it was naive to assume that the eyes could be used alone to recognise a face. Having said that, the eyes are still considered to be the most important feature when recognising a face (Zhao, Chellappa, Rosenfeld, & Phillips, 2007) and should be located with a suitably accurate algorithm.

The extension of the author's Java skills is quite extreme as this project was extremely difficult to build and debug.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

The major restriction on the project was the time available to be spent on it. It would have been very good to detect more features; firstly using the approach used already, then using another approach and comparing them. This was not possible however.

Finally, more research should be placed on efficient video frame processing and the order of importance of different facial features for the purposes of facial recognition. In addition, a face detector should be developed which is more suitable for this task (input and normalised face area output) or another one used which offers better consistency and performance.

7.3 Further areas of investigation

Areas of further investigation are covered mainly in the above section (7.2). They are summarised here for ease of reading:

- The key area of further work in preparing this application for commercial use lies in the extension of the feature extraction techniques already deployed. This means either:
 - Extending the Bayesian approach so that more features are used in order to match the face. As soon as these features are available, the Euclidean distances can be measured between them and a connected component analysis can result offering, potentially, a very accurate facial recognition system (Bing-Bing, Vass, & Zhuang, 1999).
 - Or alternatively, another feature detector could be used in conjunction with, or instead of, the Bayesian approach to better localise the features more quickly. This report suggests the technique covered in section 3.7.3 (Moon, Chellappa, & Rosenfeld, 2002) because of its ease of use and possible computational efficiency.
- Further to this, careful consideration should be given to the efficiency of the algorithms presented within the project. Ideally, some careful streamlining of these algorithms would make the application more effective and may allow the application to process a few more frames of video per second.
- Some investigation should be done into the newer queue systems available as part of the new Java 6.0 frameworks. It is likely that this will offer the functionality and efficiency required without the problem of synchronising the threads.
- The face detector software should be removed because of its incompatibility. It should either be replaced with another face detector available or one should be developed using the same techniques so that it performs better.
- If the methods used by the Bayesian approach could be made more efficient, the use of the Bayesian approach in its suggested state (rather than the modified version employed as part of this application) may become viable.
- Finally, a cost-benefit analysis should be conducted into translating this project into one created in C or an equally fast language to ascertain whether the extra number of frames processed could benefit greatly from doing this.

Chapter 8 – APPENDICES

8.1 Appendix A – Design Documentation

8.1.1 High Level System Data Flow

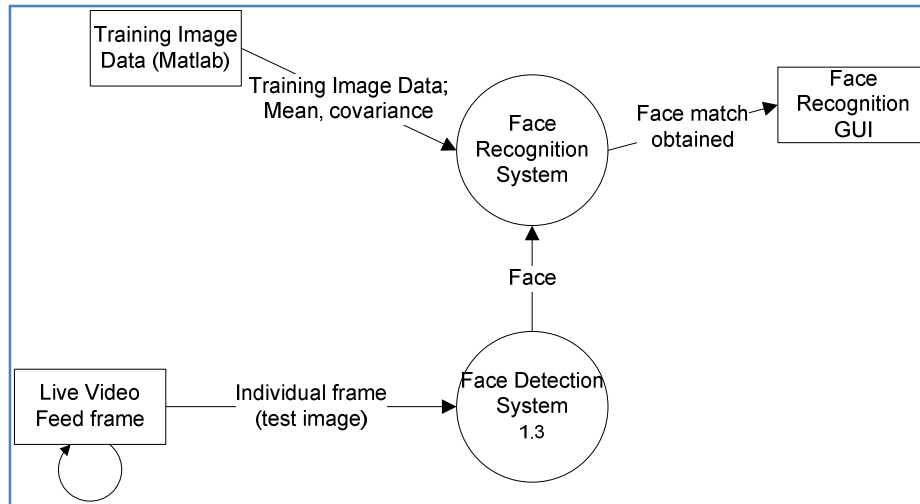


Figure 8.1-1 - High Level Data Flow Diagram (Whole System)

8.1.2 Level 1.3 High Level Face Detection flow

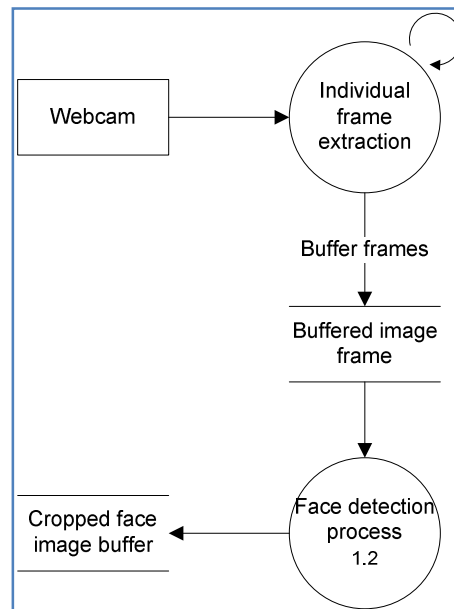


Figure 8.1-2 - High Level Face Recognition System

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.3 Level 1 feature training flow

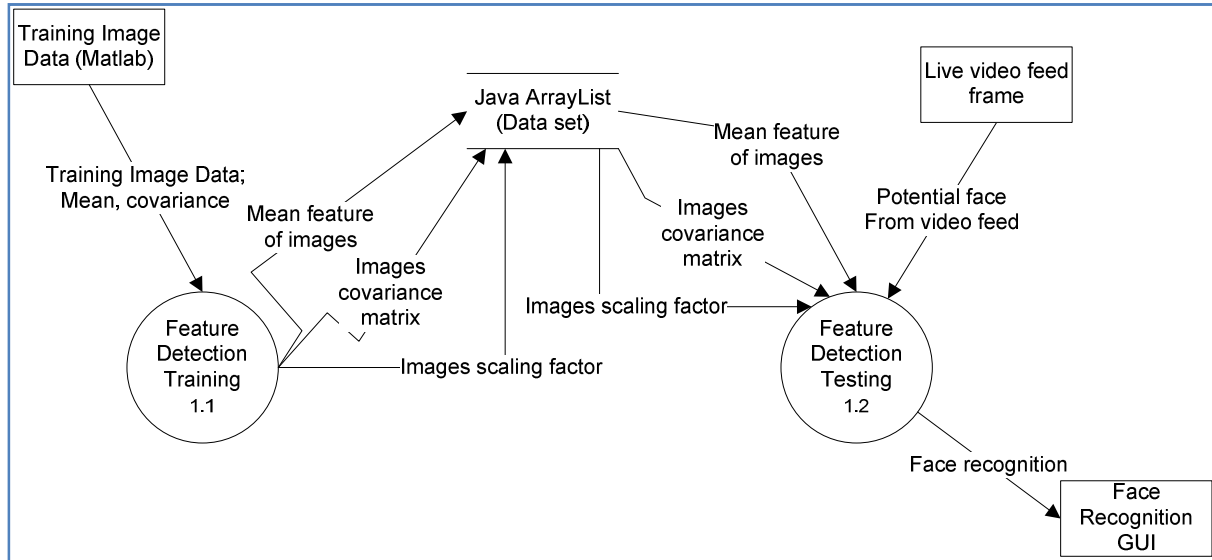


Figure 8.1-3 - Level 1 System Data Flow Diagram

8.1.4 Level 1.1 feature training flow

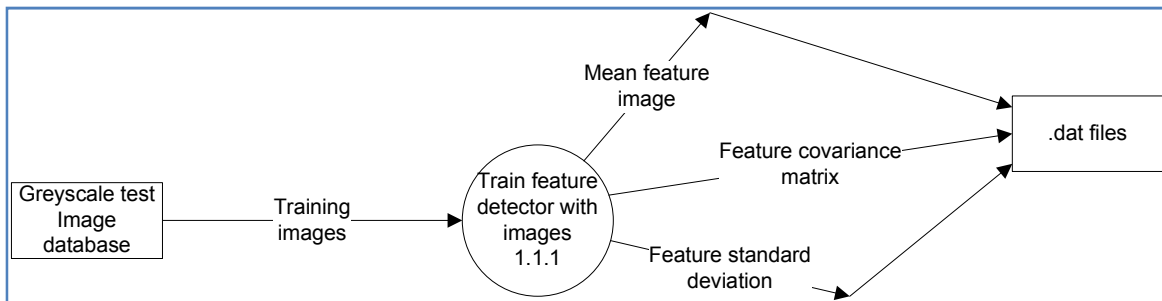


Figure 8.1-4 - Level 1.1 Feature Training Flow

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.5 Level 1.1.1 detailed feature training flow

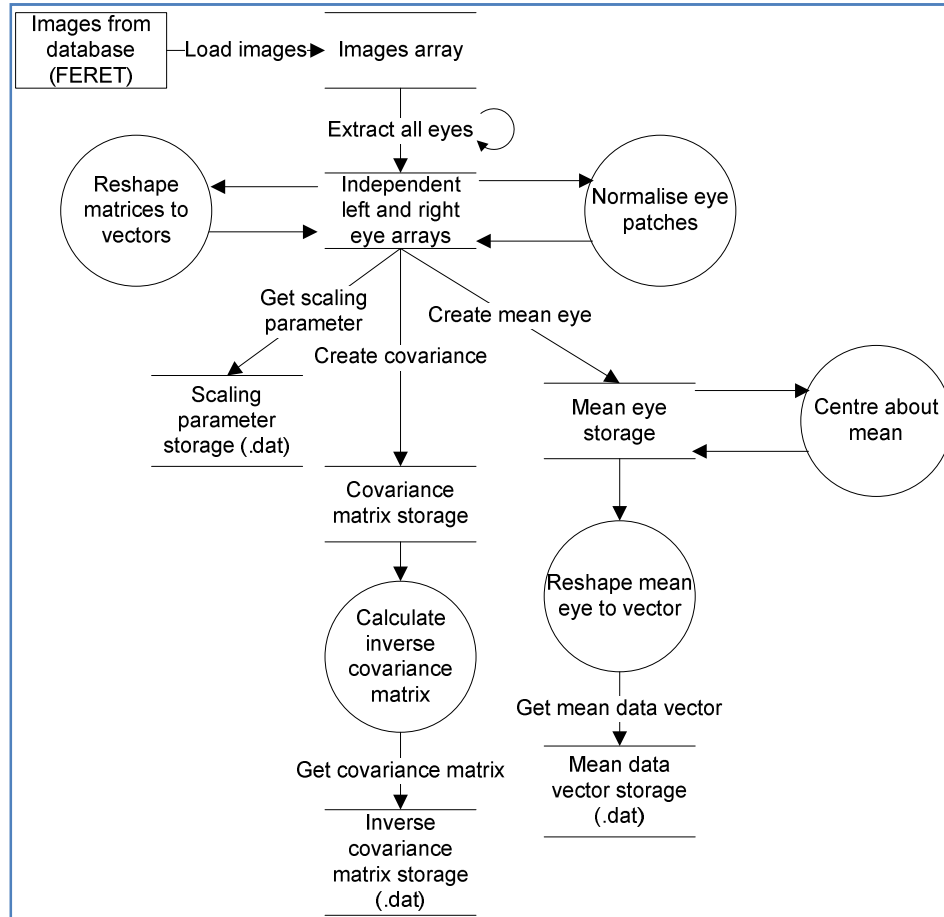


Figure 8.1-5 - Level 1.1.1 Low Level Feature Training

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.6 Level 1.2 feature extraction testing (comparison) flow

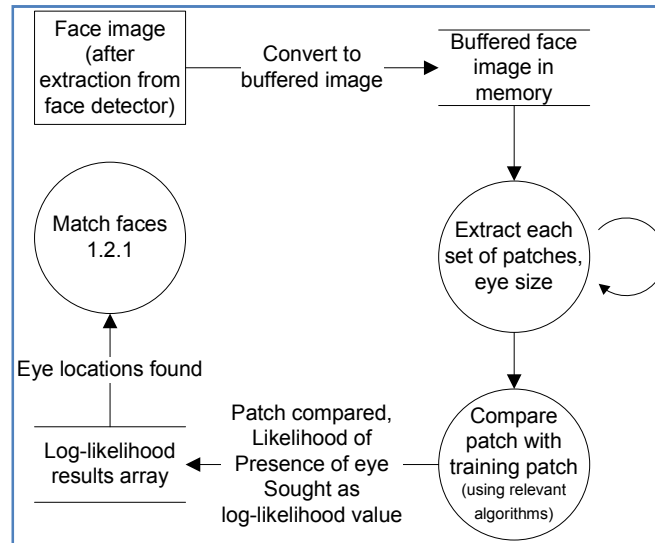


Figure 8.1-6 - Level 1.2 Feature Testing Data Flow

8.1.7 Level 1.2.1 feature matching scheme flow

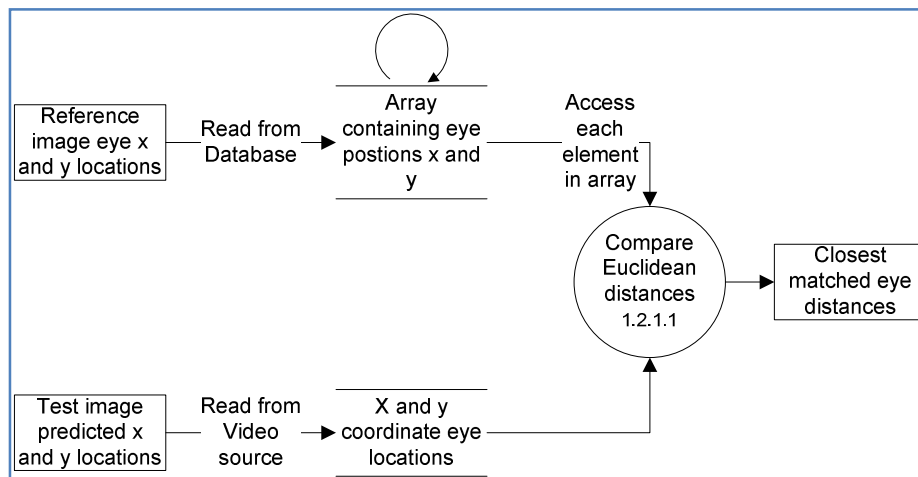


Figure 8.1-7 - Level 1.2.1 Feature Matching Scheme Data Flow

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.8 Level 1.2.1.1 low level Euclidean distance comparison flow

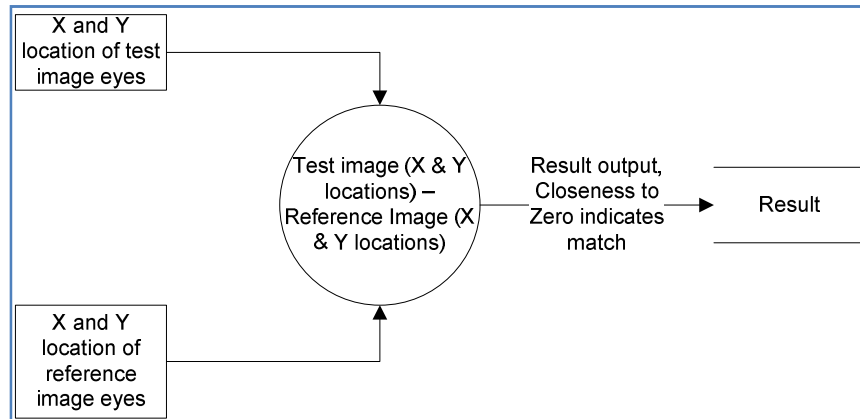


Figure 8.1-8 - Level 1.2.1.1 Low Level Matching Data Flow

8.1.9 State Transition Diagram

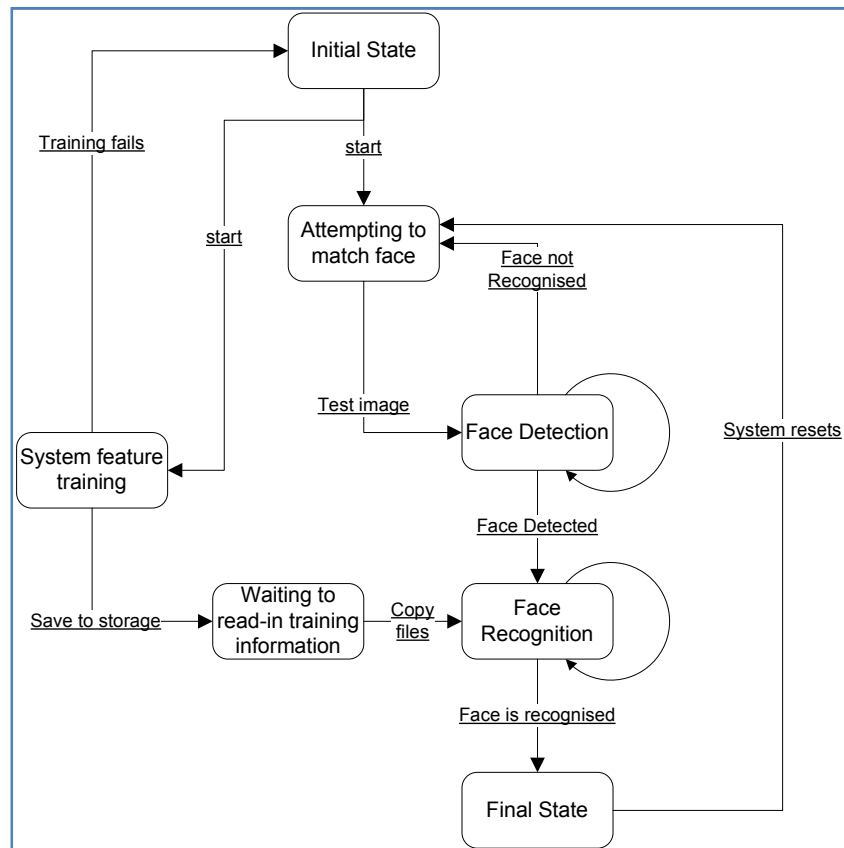


Figure 8.1-9 - System State Transition Diagram

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.10 Object Orientated Analysis (Class) Diagram

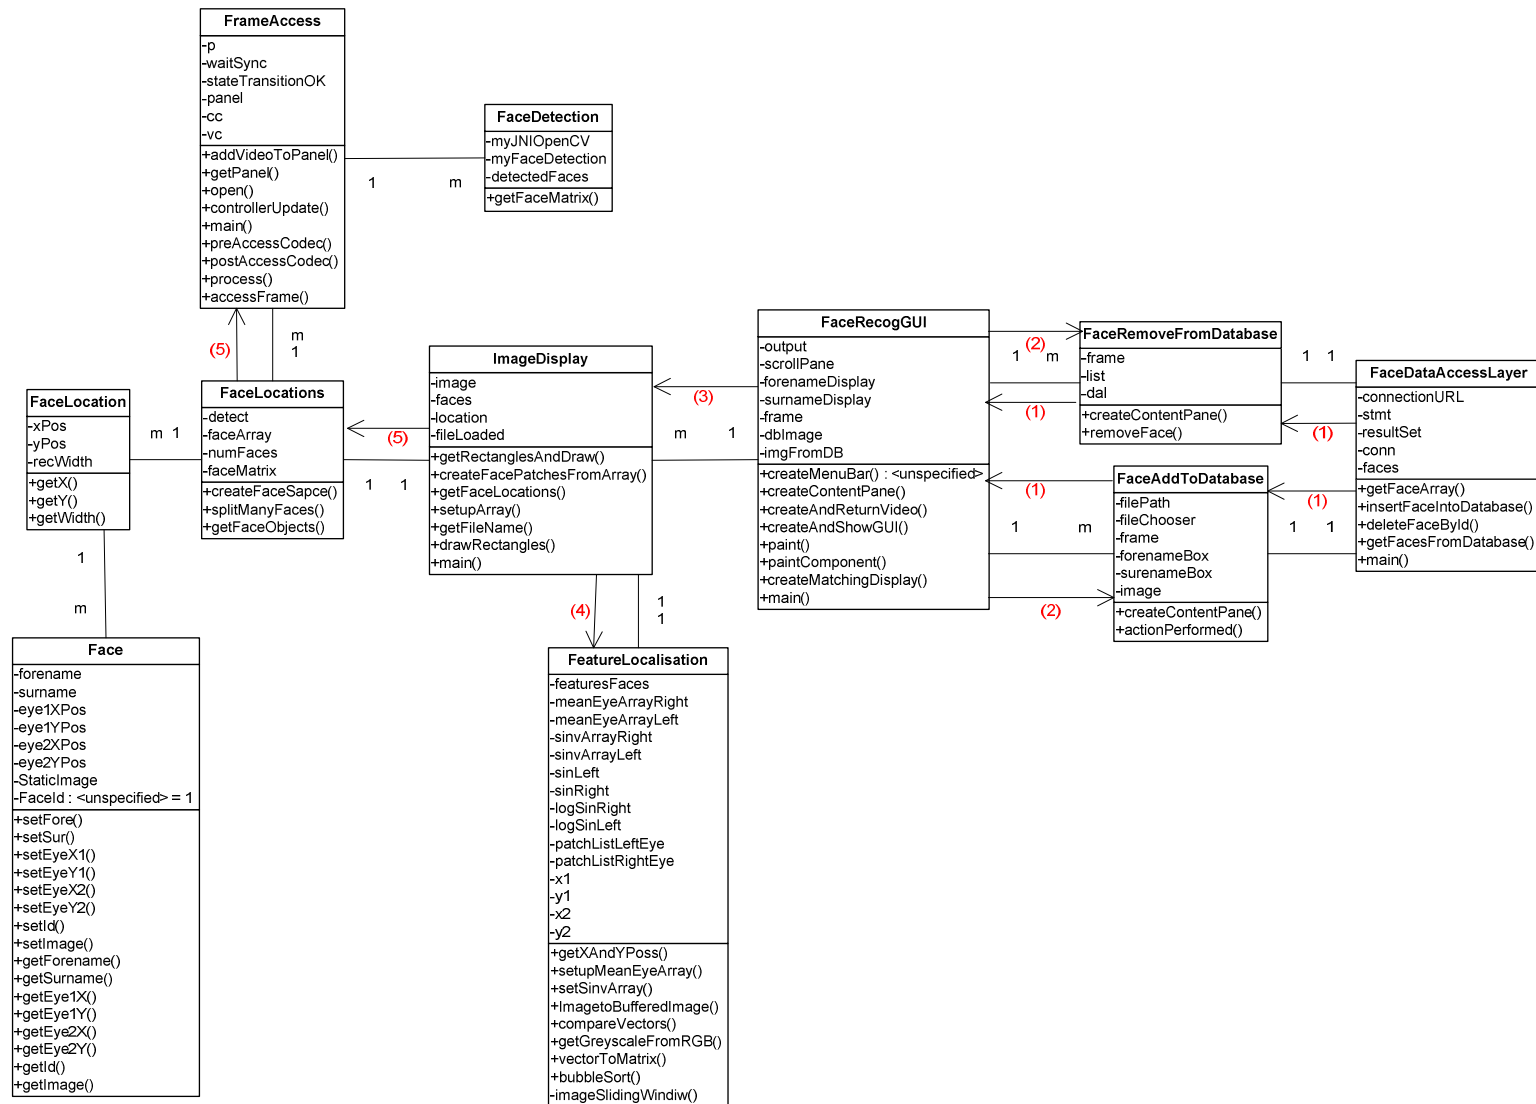


Figure 0-1- Class Diagram

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.11 Threads of execution Diagram

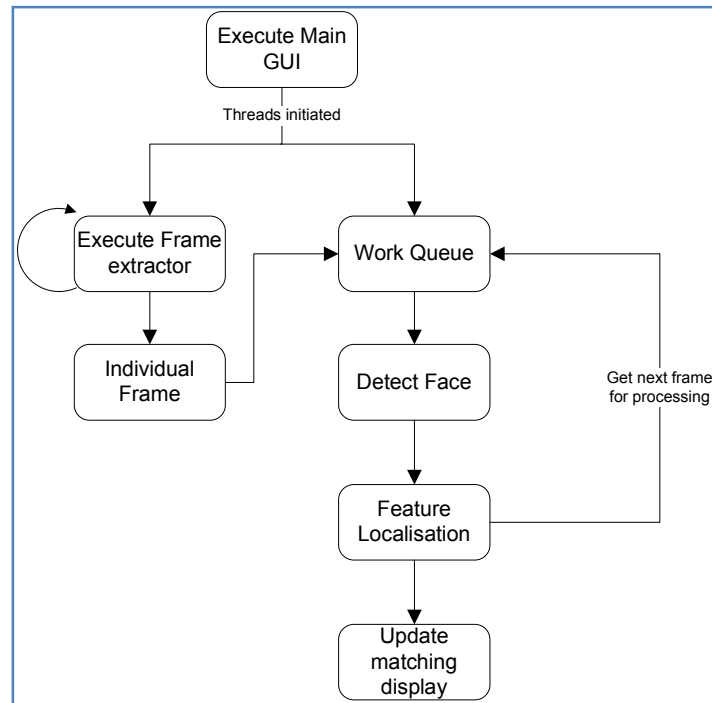


Figure 0-2 - Thread of execution (main program)

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.12 High Level Implementation Diagram

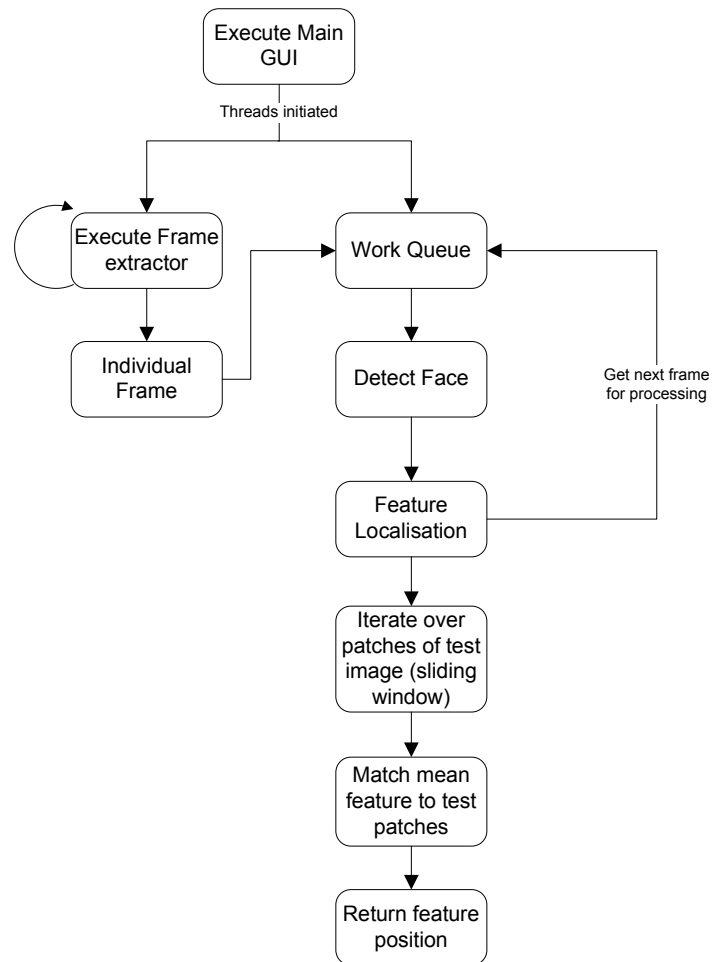


Figure 0-3 - HLD Diagram

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.13 Final class diagram

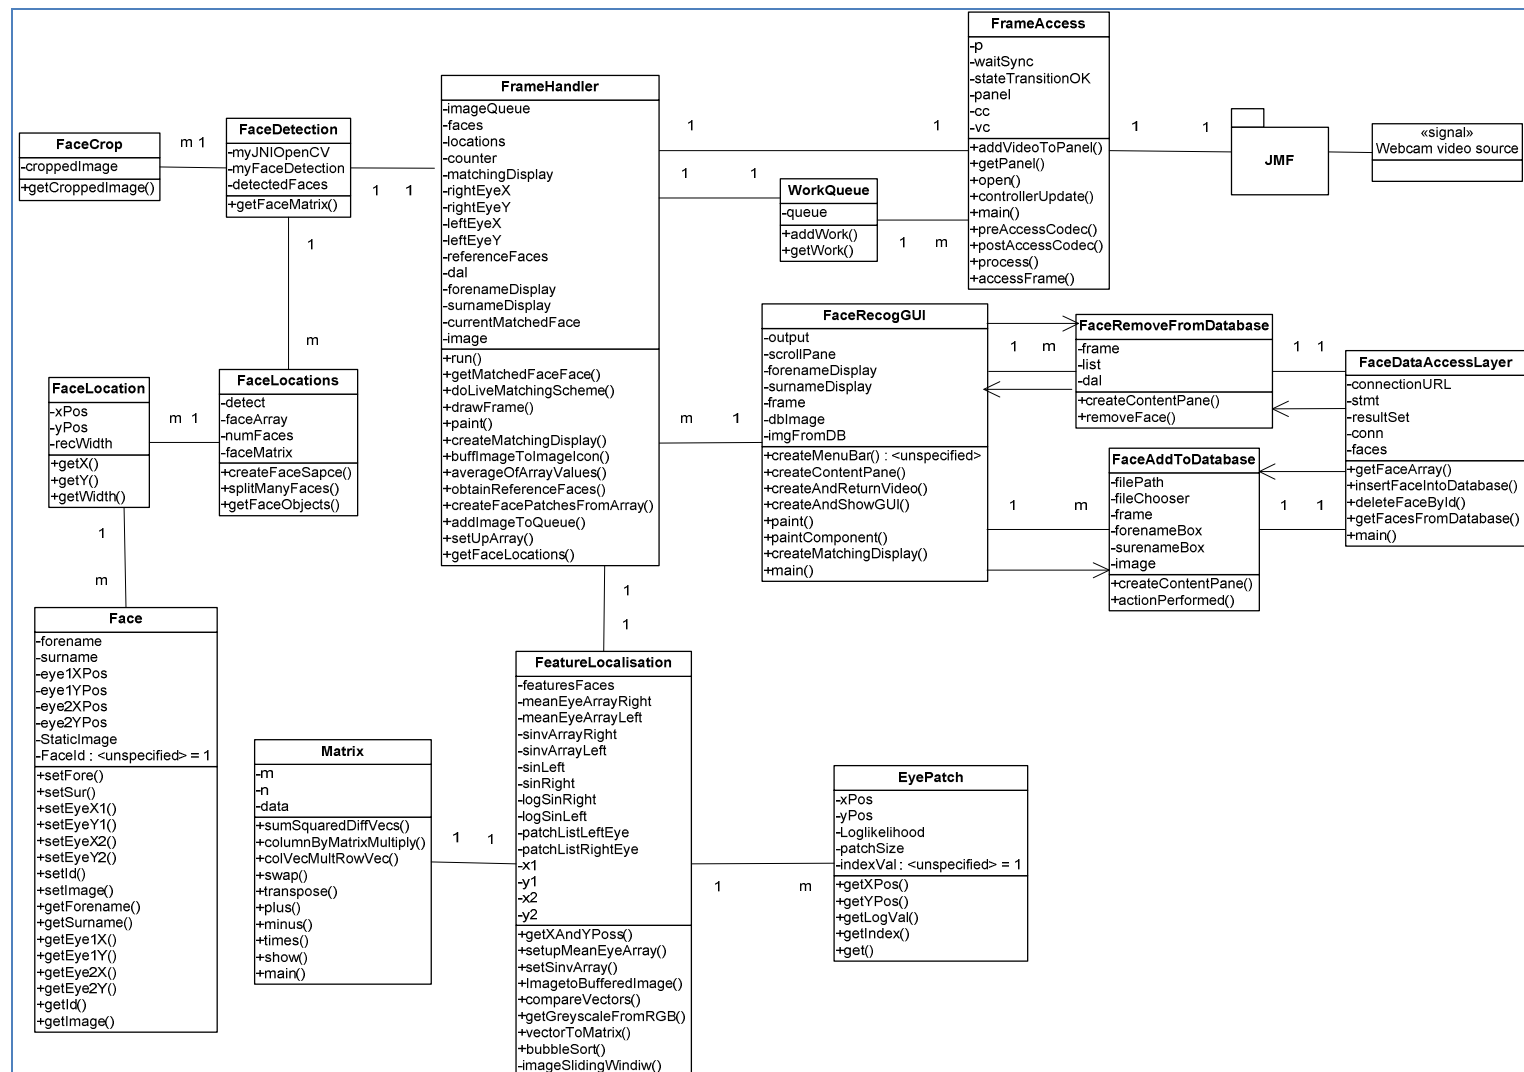


Figure 0-1 - Final Class Diagram

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.1.14 Simple Single-table Database Schema

Face	
PK	<u>FaceId</u>
	Forename Surname EyeOneX EyeOneY EyeTwoX EyeTwoY eucDistBetweenEyes filePath

Figure 0-2 - System Database Schema

8.2 Appendix B – Interface Windows

8.2.1 Main Interface Window




Figure 8.2-1 - Main Interface GUI display

8.2.2 Add Reference Image to database window




Figure 8.2-2 - Add new reference image to database GUI

8.2.3 Remove Reference Image from database window



Figure 8.2-3 - Remove existing image from reference database GUI

8.3 Appendix C – Code Listing

8.3.1 Training Eye Feature Detector (MATLAB)

The below program illustrates the use of Matlab to train the eye features for use in the feature localisation program in Java. Note; the below program shows the detection and model building of the ‘right’ eye, an identical program exists for the extraction of the left eye.

```
addpath(genpath('\\\\smb.csunix.leeds.ac.uk\\vislib\dch_matlab\\lib'));
addpath(genpath('Z:\\work\\Third Year\\PD32'));
load leedsfyp_feret.mat;
n = length(FIMGS);
eyePatchesRight = [];
% K equals regularisation parameter
k = 1;
patchSize = 16;
dim3 = 20;
for i=1:n
    Image = FIMGS(:, :, i);
    eye1 = P(:, 1, i);
    eye2 = P(:, 2, i);
    eye1(1, :) = eye1(2, :); % - dim3/2;
    eye1(2, :) = eye1(1, :) - dim3/2;
    eye2(1, :) = eye2(2, :); % - dim3/2;
    eye2(2, :) = eye2(1, :) - dim3/2;
    eye1(3, :) = dim3;
    eye2(3, :) = dim3;
    eyePatch1 = newextractpatches(Image, eye1, patchSize);
    eyePatch2 = newextractpatches(Image, eye2, patchSize);
    X = normalise(eyePatch1);
    Y = reshape(X, patchSize*patchSize, 1);
    W = normalise(eyePatch2);
    Z = reshape(W, patchSize*patchSize, 1);
    eyePatchesRight = [eyePatchesRight, Y];
end
numEyesRight = length(eyePatchesRight);
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
M          = mean(eyePatchesRight,2);
M
CXRight = eyePatchesRight - repmat(M,1,numEyesRight);
SRight  = CXRight * CXRight' / numEyesRight;
regularisationParamRight = k*eye(size(SRight));
SRight = SRight + regularisationParamRight;
ARight = mkmn(M,SRight);
ARight
M       = ARight.M;
SinvRight = ARight.icovariance;
sRight   = ARight.scaling;
save rightEyeMean.dat M -ascii;
save SinvRight.dat SinvRight -ascii;
save scaleFactor.dat sRight -ascii;
```

Figure 8.3-1- getMeanEyeRight(Left).m Training Module

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.2 Face.java class source code

```
1  import java.awt.*;
2
3  public class Face {
4      private String forename;
5      private String surname;
6      private double eyeXPos1;
7      private double eyeYPos1;
8      private double eyeXPos2;
9      private double eyeYPos2;
10     private Image staticImage;
11     private String faceId;
12
13     public Face()
14     {
15
16     }
17     public Face(String fore, String sur, String eye1X, String eye1Y, String eye2X, String
18     eye2Y, String Id, Image image)
19     {
20         forename = fore;
21         surname = sur;
22         eyeXPos1 = Double.valueOf(eye1X);
23         eyeYPos1 = Double.valueOf(eye1Y);
24         eyeXPos2 = Double.valueOf(eye2X);
25         eyeYPos2 = Double.valueOf(eye2Y);
26         staticImage = image;
27         faceId = Id;
28     }
29     public void setFore(String fore)
30     {
31         forename = fore;
32     }
33     public void setSur(String sur)
34     {
35         surname = sur;
36     }
37     public void setEyeX1(String eye)
38     {
39         eyeXPos1 = Double.valueOf(eye);
40     }
41     public void setEyeY1(String eye)
42     {
43         eyeYPos1 = Double.valueOf(eye);
44     }
45     public void setEyeX2(String eye)
46     {
47         eyeXPos2 = Double.valueOf(eye);
48     }
49     public void setEyeY2(String eye)
50     {
51         eyeYPos2 = Double.valueOf(eye);
52     }
53     public void setId(String Id)
54     {
55         faceId = Id;
56     }
57     public void setImage(Image image)
58     {
59         staticImage = image;
60     }
61     public String getForename()
62     {
63         return forename;
64     }
65     public String getSurname()
66     {
67         return surname;
68     }
69 }
```


FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
69 public Double getEye1X()  
70 {  
71     return eyeXPos1;  
72 }  
73 public Double getEye1Y()  
74 {  
75     return eyeYPos1;  
76 }  
77 public Double getEye2X()  
78 {  
79     return eyeXPos2;  
80 }  
81 public Double getEye2Y()  
82 {  
83     return eyeYPos2;  
84 }  
85 public String getId()  
86 {  
87     return faceId;  
88 }  
89 public Image getFaceImage()  
90 {  
91     return staticImage;  
92 }  
93 }
```

Figure 8.3-2 - Face.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.3 FaceCrop.java class source code

```
1  import java.awt.Image;
2  import java.awt.image.CropImageFilter;
3  import java.awt.image.FilteredImageSource;
4  import java.awt.*;
5  import java.awt.image.BufferedImage;
6
7  public class FaceCrop
8  {
9      Image croppedImage;
10     BufferedImage croppedBugImg;
11
12     public FaceCrop(Image image, int x, int y, int wh)
13     {
14         Toolkit tk = Toolkit.getDefaultToolkit();
15         Image originalImage = tk.createImage(new FilteredImageSource(image.getSource(), new
16 CropImageFilter(x, y, wh, wh)));
17         croppedImage = originalImage;
18     }
19     public FaceCrop(BufferedImage buf, int x, int y, int wh)
20     {
21         croppedBugImg = buf.getSubimage(x, y, wh, wh);
22     }
23     public Image getCroppedImage()
24     {
25         return croppedImage;
26     }
27     public BufferedImage getCroppedBugImg()
28     {
29         return croppedBugImg;
30     }
31 }
```

Figure 8.3-3 - FaceCrop.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.4 EyePatch.java class source code

```
1 public class EyePatch {
2     private int xPos;
3     private int yPos;
4     private double logLikelihood;
5     private int patchSize;
6     private int indexVal;
7
8     public EyePatch(int index,int x, int y, double logLike, int patch)
9     {
10         xPos = x;
11         yPos = y;
12         logLikelihood = logLike;
13         patchSize = patch;
14         indexVal = index;
15     }
16     public int getXPos()
17     {
18         return xPos;
19     }
20     public int getYPos()
21     {
22         return yPos;
23     }
24     public double getLogLikelihoodVal()
25     {
26         return logLikelihood;
27     }
28     public int getIndex()
29     {
30         return indexVal;
31     }
32 }
```

Figure 8.3-4 - EyePatch.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.5 *Matrix.java* class source code

```
1 package FaceRecogPackage;
2 /*****
3  * Compilation:  javac Matrix.java
4  * Execution:    java Matrix
5  * Author:       Robert Sedgewick and Kevin Wayne,
6  http://www.cs.princeton.edu/introcs/95linear/Matrix.java.html
7  * Modified by:  Martyn Booth (scs3mwb)
8  *
9  * A class providing basic matrix calculations sed by some of the other classes
10 *
11 *****/
12
13 final public class Matrix {
14     private final int M;           // number of rows
15     private final int N;           // number of columns
16     private final double[][] data; // M-by-N array
17
18     // create M-by-N matrix of 0's
19     public Matrix(int M, int N) {
20         this.M = M;
21         this.N = N;
22         data = new double[M][N];
23     }
24
25     // create matrix based on 2d array
26     public Matrix(double[][] data) {
27         M = data.length;
28         N = data[0].length;
29         this.data = new double[M][N];
30         for (int i = 0; i < M; i++)
31             for (int j = 0; j < N; j++)
32                 this.data[i][j] = data[i][j];
33     }
34
35     // copy constructor
36     private Matrix(Matrix A) { this(A.data); }
37
38     // create and return a random M-by-N matrix with values between 0 and 1
39     public static Matrix random(int M, int N) {
40         Matrix A = new Matrix(M, N);
41         for (int i = 0; i < M; i++)
42             for (int j = 0; j < N; j++)
43                 A.data[i][j] = Math.random();
44         return A;
45     }
46     public static double sumSquaredDiffVecs(double[] vector1, double[] vector2)
47     {
48         double differenceValue = 0.0;
49         for(int i=0;i<vector1.length;i++)
50         {
51             double tempVal = vector2[i] - vector1[i];
52             differenceValue += tempVal * tempVal;
53         }
54         return differenceValue;
55     }
56     public static double[] columnByMatrixMultiply(Matrix m, double[] columnVector)
57     {
58         double[] newColumn = new double[columnVector.length];
59         for(int i =0; i < m.data[0].length;i++)
60         {
61             double tempValue = 0.0;
62             for(int j=0;j<m.data[0].length;j++)
63             {
64                 tempValue += m.data[i][j] * columnVector[j];
65                 newColumn[i] = tempValue;
66             }
67         }
68     }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
1      }
2      return newColumnn;
3  }
4  public static double colVecMultRowVec(double[] rowVector, double[] columnVector)
5  {
6      double result = 0.0;
7      for(int i=0;i<rowVector.length;i++)
8      {
9          result += rowVector[i] * columnVector[i];
10     }
11     return result;
12 }
13
14 // create and return the transpose of the invoking matrix
15 public Matrix transpose() {
16     Matrix A = new Matrix(N, M);
17     for (int i = 0; i < M; i++)
18         for (int j = 0; j < N; j++)
19             A.data[j][i] = this.data[i][j];
20     return A;
21 }
22
23 // return C = A * B
24 public Matrix times(Matrix B) {
25     Matrix A = this;
26     if (A.N != B.M) throw new RuntimeException("Illegal matrix dimensions.");
27     Matrix C = new Matrix(A.M, B.N);
28     for (int i = 0; i < C.M; i++)
29         for (int j = 0; j < C.N; j++)
30             for (int k = 0; k < A.N; k++)
31                 C.data[i][j] += (A.data[i][k] * B.data[k][j]);
32     return C;
33 }
34
35 // print matrix to standard output
36 public void show() {
37     for (int i = 0; i < M; i++) {
38         for (int j = 0; j < N; j++)
39             System.out.printf("%9.4f ", data[i][j]);
40         System.out.println();
41     }
42 }
43
44
45
46 // test client
47 public static void main(String[] args) {
48     double[][] d = { { 1, 2, 3 }, { 4, 5, 6 }, { 9, 1, 3 } };
49     Matrix D = new Matrix(d);
50     D.show();
51     System.out.println();
52
53     Matrix A = Matrix.random(5, 5);
54     A.show();
55     System.out.println();
56
57     Matrix B = A.transpose();
58     B.show();
59     System.out.println();
60
61     B.times(A).show();
62     System.out.println();
63
64     Matrix b = Matrix.random(5, 1);
65     b.show();
66     System.out.println();
67
68 }}
```

Figure 8.3-5 - Matrix.java class

8.3.6 *OpenFileAction.java class source code*

```
import javax.swing.*;
import java.io.*;
import java.awt.event.*;

// This action creates and shows a modal open-file dialog.
public class OpenFileAction extends AbstractAction {
    JFrame frame;
    JFileChooser chooser;

    OpenFileAction(JFrame frame, JFileChooser chooser) {
        super("Open...");
        this.chooser = chooser;
        this.frame = frame;
    }

    public void actionPerformed(ActionEvent evt) {
        // Show dialog; this method does not return until dialog is closed
        chooser.showOpenDialog(frame);

        // Get the selected file
        File file = chooser.getSelectedFile();
    }
};
```

Figure 8.3-6 - OpenFileAction.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.7 FaceLocation.java class source code

```
1 public class FaceLocation {  
2  
3     private int xPositon;  
4     private int yPositon;  
5     private int recWidth;  
6  
7     public FaceLocation(int x, int y, int width)  
8     {  
9         xPositon = x;  
10        yPositon = y;  
11        recWidth = width;  
12    }  
13    public int getX()  
14    {  
15        return xPositon;  
16    }  
17    public int getY()  
18    {  
19        return yPositon;  
20    }  
21    public int getWidth()  
22    {  
23        return recWidth; }  
24 }  
25
```

Figure 8.3-7 - FaceLocation.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.8 FaceLocations.java class source code

```
import java.util.*;

public class FaceLocations {

    private FaceDetection detect;
    private int[] faceArray;
    private int numFaces = 0;
    private ArrayList faceMatrix;

    public FaceLocations()
    {
    }

    public FaceLocations(String file, int cascade, int height, int width)
    {
        faceMatrix = new ArrayList();
        String fileNameSlash = file;
        String fileName;
        if(fileNameSlash.indexOf("\\") < 0)
        {
            fileName = "trainPhotos/" + file;
        }
        else
        {
            fileName = file;
        }

        createFaceSpace(fileName, cascade, height, width);
    }

    public void createFaceSpace(String file, int cascade, int height, int width)
    {
        detect = new FaceDetection(file, 1, 40, 40);
        faceArray = detect.getFaceMatrix();
        numFaces = faceArray.length / 4;
        splitManySpaces(faceArray);
    }

    public void splitManySpaces(int[] faces)
    {
        for (int i = 0; i < numFaces; i++)
        {
            int x = 0;
            int y = 0;
            int w = 0;
            x = faceArray[4 * i + 0];
            y = faceArray[4 * i + 1];
            w = faceArray[4 * i + 2];
            FaceLocation face = new FaceLocation(x,y,w);
            faceMatrix.add(face);
            // System.out.println("1: " + x + " 2: " + y + " 3: " + w);
        }
    }

    public ArrayList getFaceObjects()
    {
        return faceMatrix;
    }

    public static void main(String args[])
    {
        FaceLocations myFaceDetection = new FaceLocations();
    }
}
```

Figure 8.3-8 - FaceLocations.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.9 FaceDataAccessLayer.java class source code

```
1  import java.awt.Toolkit;
2  import java.sql.*;
3  import java.util.*;
4  import java.io.*;
5  import java.awt.*;
6
7
8  public class FaceDataAccessLayer {
9      String connectionUrl = "jdbc:sqlserver://localhost:1433;" +
10         "databaseName=FaceRecognition;user=FaceUser;password=Password1";
11      Statement stmt = null;
12      ResultSet rs = null;
13      Connection conn;
14      ArrayList<Face> faces;
15
16      public FaceDataAccessLayer()
17      {
18          try
19          {
20              Class.forName("com.microsoft.sqlserver.jdbc.SQLServerDriver");
21              conn = DriverManager.getConnection(connectionUrl);
22          }
23          catch (Exception e)
24          {
25              System.out.println(e.getMessage());
26              System.out.println("Connection error");
27          }
28      }
29
30      public ArrayList<Face> getFaceArray()
31      {
32          return faces;
33      }
34
35      public void insertFaceIntoDatabase(String name, String sur, int eye1X, int eye1Y, int
36      eye2X, int eye2Y, String faceIm)
37      {
38          int len;
39          String query;
40          PreparedStatement pstmt;
41
42          try
43          {
44              File file = new File(faceIm);
45              FileInputStream fis = new FileInputStream(file);
46              len = (int)file.length();
47
48              query = ("insert into Face (Forename,Surname,eyeOneX,eyeOneY,eyeTwoX,eyeTwoY,
49      FaceImage, FileName, FileSize) VALUES(?,?,?,?,?,?,?,?)");
50              pstmt = conn.prepareStatement(query);
51              pstmt.setString(1,name);
52              pstmt.setString(2,sur);
53              pstmt.setString(3,String.valueOf(eye1X));
54              pstmt.setString(4,String.valueOf(eye1Y));
55              pstmt.setString(5,String.valueOf(eye2X));
56              pstmt.setString(6,String.valueOf(eye2Y));
57              pstmt.setString(8,file.getName());
58              pstmt.setInt(9, len);
59
60              // Method used to insert a stream of bytes
61              pstmt.setBinaryStream(7, fis, len);
62              pstmt.executeUpdate();
63          }
64          catch (Exception e)
65          {
66              e.printStackTrace();
67          }
68      }
69  }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
66     }
67   }
68   public void deleteFaceById(String Id)
69   {
70     String query = "DELETE FROM Face WHERE FaceId = ";
71
72     try
73     {
74       stmt = conn.createStatement();
75       query += Id;
76       System.out.println(query);
77       stmt.executeQuery(query);
78       conn.close();
79     }
80     catch (Exception e)
81     {
82       e.printStackTrace();
83     }
84   }
85
86   public void getFacesFromDatabase()
87   {
88     String SQL = "SELECT * FROM Face";
89     faces = new ArrayList();
90     try
91     {
92       stmt = conn.createStatement();
93       rs = stmt.executeQuery(SQL);
94       while (rs.next())
95       {
96
97         //System.out.println(rs.getString(1)+"\t"+rs.getString(2)+"\t"+rs.getString(3)+"\t"+rs.ge
98         tString(4)+"\t"+rs.getString(5)+"\t"+rs.getString(6)+"\t"+rs.getString(7));
99         Face newFace = new Face();
100        newFace.setId(rs.getString(1));
101        newFace.setFore(rs.getString(2));
102        newFace.setSur(rs.getString(3));
103        newFace.setEyeX1(rs.getString(4));
104        newFace.setEyeY1(rs.getString(5));
105        newFace.setEyeX2(rs.getString(6));
106        newFace.setEyeY2(rs.getString(7));
107        //Turn the image capture back on once the image insert is ready
108
109        Toolkit toolkit = Toolkit.getDefaultToolkit();
110        byte[] imageBinData;
111        imageBinData = rs.getBytes(8);
112        Image tempIm = toolkit.createImage(imageBinData);
113        newFace.setImage(tempIm);
114        faces.add(newFace);
115      }
116      //conn.close();
117    }
118    catch (Exception e)
119    {
120      System.out.println(e.getMessage());
121      System.out.println("Get Faces statement failed");
122    }
123  }
124
125  }
126  public static void main(String[] args)
127  {
128    FaceDataAccessLayer dal = new FaceDataAccessLayer();
129  }
130 }
131
```

Figure 8.3-9 - FaceDataAccessLayer.java class

8.3.10 FaceAddToDatabase.java class source code

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
1 import java.awt.*;
2 import java.awt.event.*;
3 import javax.swing.*;
4 import java.io.*;
5 import java.util.*;
6
7
8 public class FaceAddToDatabase extends JFrame
9 {
10     private JTextField filePath = new JTextField(15);
11     private JFileChooser _fileChooser = new JFileChooser();
12     private JFrame frame;
13     private JTextField forenameBox;
14     private JTextField surnameBox;
15     private Image image;
16
17     public FaceAddToDatabase()
18     {
19         frame = new JFrame("Add New Reference Face");
20         frame.setDefaultCloseOperation(JFrame.HIDE_ON_CLOSE);
21         //frame.setSize(width, height);
22         frame.setContentPane(this.createContentPane());
23         frame.setVisible(true);
24         frame.pack();
25         frame.setLocationRelativeTo(null);
26
27     }
28
29     public Container createContentPane() {
30         JLabel forename = new JLabel();
31         forename.setSize(100, 25);
32         forename.setText("Forename:");
33         JLabel surname = new JLabel();
34         surname.setSize(100, 25);
35         surname.setText("Surname:");
36         JLabel image = new JLabel();
37         image.setSize(100,25);
38         image.setText("Face Image to upload:");
39
40         forenameBox = new JTextField();
41         forename.setLabelFor(forenameBox);
42         //forenameBox.setSize(100,25);
43         forenameBox.setColumns(15);
44         surnameBox = new JTextField();
45         surname.setLabelFor(surnameBox);
46         //surnameBox.setSize(100,25);
47         surnameBox.setColumns(15);
48         //Create the content-pane-to-be.
49         //JTextField filePath = new JTextField();
50         image.setLabelFor(filePath);
51         //filePath.setColumns(15);
52         JButton imageBrowse = new JButton();
53         imageBrowse.addActionListener(new OpenAction());
54         imageBrowse.setText("Browse");
55         JPanel contentPane = new JPanel(new BorderLayout());
56         contentPane.setOpaque(true);
57         JPanel labelPanel = new JPanel(new GridLayout(0,1));
58         labelPanel.add(forename);
59         labelPanel.add(surname);
60         labelPanel.add(image);
61         JPanel textPanel = new JPanel(new GridLayout(0,1));
62         textPanel.add(forenameBox);
63         textPanel.add(surnameBox);
64         textPanel.add(filePath);
65         JPanel right = new JPanel(new GridLayout(0,1));
66         JLabel nullLabel = new JLabel();
67         nullLabel.setSize(100,25);
68         nullLabel.setText(" ");
69         JLabel nullLabel2 = new JLabel();
70         nullLabel2.setSize(100,25);
71         nullLabel2.setText(" ");
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
72     right.add(nullLabel);
73     //right.add(nullLabel2);
74     right.add(imageBrowse);
75     JPanel options = new JPanel(new FlowLayout());
76     JButton submit = new JButton("OK");
77     submit.addActionListener(new SubmitFaceToDB());
78     JButton cancel = new JButton("Cancel");
79     cancel.addActionListener(new HideAddFace());
80     options.add(submit);
81     options.add(cancel);
82     contentPane.add(labelPanel, BorderLayout.WEST);
83     contentPane.add(textPanel, BorderLayout.CENTER);
84     contentPane.add(right, BorderLayout.EAST);
85     contentPane.add(options, BorderLayout.SOUTH);
86
87     return contentPane;
88 }
89
90 ////////////////////////////////////////////////// OpenAction
91 class OpenAction implements ActionListener {
92     public void actionPerformed(ActionEvent ae) {
93         //... Open a file dialog.
94         int retval = _fileChooser.showOpenDialog(FaceAddToDatabase.this);
95         if (retval == JFileChooser.APPROVE_OPTION) {
96             //... The user selected a file, get it, use it.
97             File file = _fileChooser.getSelectedFile();
98
99             //... Update user interface.
100             filePath.setText(file.getPath());
101         }
102     }
103 }
104 private class HideAddFace extends AbstractAction
105 {
106     public HideAddFace()
107     {
108         super("Disable frame");
109     }
110
111     public void actionPerformed(ActionEvent e)
112     {
113         frame.setVisible(false);
114     }
115 }
116 private class SubmitFaceToDB extends AbstractAction
117 {
118     public SubmitFaceToDB()
119     {
120     }
121     public void actionPerformed(ActionEvent e)
122     {
123         String forenameVal = forenameBox.getText();
124         String surnameVal = surnameBox.getText();
125         //System.out.println(forenameVal);
126         //System.out.println(surnameVal);
127         String filePathVal = filePath.getText();
128         Toolkit toolkit = Toolkit.getDefaultToolkit();
129         String file = filePathVal;
130         Image faceImage = toolkit.getImage(file);
131         FaceLocations locations = new FaceLocations(file, 1, 40, 40);
132         ArrayList<FaceLocation> faces = new ArrayList<FaceLocation>();
133         faces = locations.getFaceObjects();
134         FaceLocation face = (FaceLocation) faces.get(0);
135         int x = face.getX();
136         int y = face.getY();
137         int w = face.getWidth();
138         FaceCrop crop = new FaceCrop(image, x, y, w);
139         FeatureLocalisation feature = new
140 FeatureLocalisation(crop.getCroppedImage(), file, 16);
141         int eye1X = feature.getXPos1();
142         int eye1Y = feature.getYPos1();
143         int eye2X = feature.getXPos2();
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
143         int eye2Y = feature.getYPos2();
144         /*BufferedImage buffFace = FeatureLocalisation.toBufferedImage(faceImage);
145         buffFace = buffFace.getSubimage(x, y, w, w);*/
146         //filePathVal = filePathVal.replace("\\", "/");
147         //System.out.print(filePath.getText());
148         //Image img = Toolkit.getDefaultToolkit().getImage(filePathVal);
149         //String file = filePath.getText();
150         //System.out.println(filePathVal);
151         //image =
152         toolkit.getImage("http://www.particle.kth.se/~lindsey/JavaCourse/Book/Part1/Java/Chapter06/javalo
153         go52x88.gif");//filePath.getText());
154         FaceDataAccessLayer dal = new FaceDataAccessLayer();
155         /*System.out.println(img.getHeight(null));
156         System.out.println(img.getWidth(null));*/
157         dal.insertFaceIntoDatabase(forenameVal, surnameVal, eye1X, eye1Y, eye2X,
158         eye2Y, filePathVal);
159         forenameBox.setText("");
160         surnameBox.setText("");
161         filePath.setText("");
162     }
163 }
164 }
```

Figure 8.3-10 - FaceAddToDatabase.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.11 FaceRemoveFromDatabase.java class source code

```
1  import java.awt.*;
2  import java.awt.event.*;
3  import javax.swing.*;
4  import java.util.*;
5
6  class FaceRemoveFromDatabase extends JFrame
7  {
8      private JFrame frame;
9      private JList list;
10     private FaceDataAccessLayer dal;
11
12     public FaceRemoveFromDatabase()
13     {
14         dal = new FaceDataAccessLayer();
15         frame = new JFrame("Remove a Reference Face");
16         frame.setDefaultCloseOperation(JFrame.HIDE_ON_CLOSE);
17         //frame.setSize(width, height);
18         frame.setContentPane(this.createContentPane());
19         frame.setVisible(true);
20         frame.pack();
21         frame.setLocationRelativeTo(null);
22     }
23     public Container createContentPane()
24     {
25         JPanel contentPane = new JPanel(new GridLayout(0,1));
26         JPanel jListPane = new JPanel(new FlowLayout());
27         contentPane.setOpaque(true);
28         dal.getFacesFromDatabase();
29         ArrayList<Face> faces = dal.getFaceArray();
30         String[] faceDetails = new String[faces.size()];
31         for(int i = 0; i<faces.size(); i++)
32         {
33             String tempString;
34             Face face = faces.get(i);
35             String forename = face.getForename();
36             String surname = face.getSurname();
37             //String id = face.getId();
38             tempString = forename+"\t"+surname;
39             faceDetails[i] = tempString;
40         }
41         list = new JList(faceDetails);
42         jListPane.add(list);
43         JPanel buttonPane = new JPanel(new FlowLayout());
44         JButton submit = new JButton("OK");
45         submit.addActionListener(new RemoveFace());
46         JButton cancel = new JButton("Cancel");
47         buttonPane.add(submit);
48         buttonPane.add(cancel);
49         cancel.addActionListener(new HideRemFace());
50         contentPane.add(jListPane);
51         contentPane.add(buttonPane);
52
53         return contentPane;
54     }
55     private class RemoveFace extends AbstractAction
56     {
57         public RemoveFace()
58         {
59
60         }
61         public void actionPerformed(ActionEvent e)
62         {
63             //FaceDataAccessLayer dataAL = new FaceDataAccessLayer();
64             int currentItem = list.getSelectedIndex();
65             dal.getFacesFromDatabase();
66             ArrayList<Face> facesArr = dal.getFaceArray();
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
67         String currentId = facesArr.get(currentItem).getId();
68         //System.out.println(currentId);
69         dal.deleteFaceById(currentId);
70         frame.hide();
71     }
72 }
73 private class HideRemFace extends AbstractAction
74 {
75     public HideRemFace()
76     {
77         super("Disable frame");
78     }
79
80     public void actionPerformed(ActionEvent e)
81     {
82         frame.setVisible(false);
83     }
84 }
85 }
```

Figure 8.3-11 - FaceRemoveFromDatabase.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.12 *FrameAccess.java* class source code

```
1  import java.awt.*;
2  import javax.media.*;
3  import javax.media.control.TrackControl;
4  import javax.media.Format;
5  import javax.media.format.*;
6  import javax.swing.JFrame;
7  import javax.swing.JPanel;
8  import javax.media.util.*;
9  import java.io.*;
10 import java.awt.image.BufferedImage;
11 import javax.imageio.ImageIO;
12
13 public class FrameAccess extends JFrame implements ControllerListener {
14
15     Processor p;
16     Object waitSync = new Object();
17     boolean stateTransitionOK = true;
18     private JPanel panel;
19     private Component cc;
20         private Component vc;
21         private int counter=0;
22         private int frameThrottlingCounter=0;
23
24     public FrameAccess()
25     {
26
27     }
28     public void addVideoToPanel(String urlIn)
29     {
30         panel = new JPanel(new BorderLayout());
31         panel.setOpaque(true);
32         String url;
33         if(urlIn.equalsIgnoreCase(""))
34         {
35             url = "vfw:Microsoft WDM Image Capture (Win32):0";
36         }
37         else
38         {
39             url = urlIn;
40         }
41
42         MediaLocator ml;
43
44         if ((ml = new MediaLocator(url)) == null)
45         {
46             System.err.println("Cannot build media locator from: " + url);
47             System.exit(0);
48         }
49         open(ml);
50         if ((vc = p.getVisualComponent()) != null)
51         {
52             panel.add(vc,BorderLayout.CENTER);
53         }
54
55         if ((cc = p.getControlPanelComponent()) != null)
56         {
57             panel.add(cc,BorderLayout.SOUTH);
58         }
59         p.start();
60
61         panel.setOpaque(true);
62         panel.setVisible(true);
63
64         //panel.add(scrollPane, BorderLayout.CENTER);
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
65 }
66 public JPanel getPanel()
67 {
68     return panel;
69 }
70 /**
71  * Given a media locator, create a processor and use that processor
72  * as a player to playback the media.
73  *
74  * During the processor's Configured state, two "pass-thru" codecs,
75  * PreAccessCodec and PostAccessCodec, are set on the video track.
76  * These codecs are used to get access to individual video frames
77  * of the media.
78  *
79  * Much of the code is just standard code to present media in JMF.
80  */
81 public boolean open(MediaLocator ml) {
82
83     try
84     {
85         p = Manager.createProcessor(ml);
86     }
87     catch (Exception e)
88     {
89         System.err.println("Failed to create a processor from the given url: " + e);
90         return false;
91     }
92
93     p.addControllerListener(this);
94
95     // Put the Processor into configured state.
96     p.configure();
97     if (!waitForState(p.Configured))
98     {
99         System.err.println("Failed to configure the processor.");
100        return false;
101    }
102
103    // So I can use it as a player.
104    p.setContentDescriptor(null);
105
106    // Obtain the track controls.
107    TrackControl tc[] = p.getTrackControls();
108
109    if (tc == null)
110    {
111        System.err.println("Failed to obtain track controls from the processor.");
112        return false;
113    }
114
115    // Search for the track control for the video track.
116    TrackControl videoTrack = null;
117
118    for (int i = 0; i < tc.length; i++)
119    {
120        if (tc[i].getFormat() instanceof VideoFormat)
121        {
122            videoTrack = tc[i];
123            break;
124        }
125    }
126
127    if (videoTrack == null)
128    {
129        System.err.println("The input media does not contain a video track.");
130        return false;
131    }
132
133    System.err.println("Video format: " + videoTrack.getFormat());
134
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
135 // Instantiate and set the frame access codec to the data flow path.
136 try
137 {
138     Codec codec[] =
139     {
140         new PreAccessCodec(), new PostAccessCodec();
141         videoTrack.setCodecChain(codec);
142     }
143 catch (UnsupportedPlugInException e)
144 {
145     System.err.println("The process does not support effects.");
146 }
147
148 // Realize the processor.
149 p.prefetch();
150 if (!waitForState(p.Prefetched))
151 {
152     System.err.println("Failed to realize the processor.");
153     return false;
154 }
155
156 // Display the visual & control component if there's one.
157
158 /*setLayout(new BorderLayout());
159
160
161 if ((vc = p.getVisualComponent()) != null) {
162     add("Center", vc);
163     /*panel*///add(vc, BorderLayout.CENTER);
164 /*}
165
166 if ((cc = p.getControlPanelComponent()) != null) {
167     add("South", cc);
168     /*panel.*///add(cc, BorderLayout.SOUTH);
169 //}
170
171 // Start the processor.
172 //p.start();
173
174 //panel.setOpaque(true);
175 //setVisible(true);
176
177
178 return true;
179 }
180
181 public void addNotify() {
182     super.addNotify();
183     pack();
184 }
185
186 /**
187  * Block until the processor has transitioned to the given state.
188  * Return false if the transition failed.
189  */
190 boolean waitForState(int state) {
191     synchronized (waitSync) {
192         try {
193             while (p.getState() != state && stateTransitionOK)
194                 waitSync.wait();
195             } catch (Exception e) {}
196         }
197     return stateTransitionOK;
198 }
199
200
201 /**
202  * Controller Listener.
203  */
204 public void controllerUpdate(ControllerEvent evt) {
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
205
206     if (evt instanceof ConfigureCompleteEvent ||
207         evt instanceof RealizeCompleteEvent ||
208         evt instanceof PrefetchCompleteEvent) {
209         synchronized (waitSync) {
210             stateTransitionOK = true;
211             waitSync.notifyAll();
212         }
213     } else if (evt instanceof ResourceUnavailableEvent) {
214         synchronized (waitSync) {
215             stateTransitionOK = false;
216             waitSync.notifyAll();
217         }
218     } else if (evt instanceof EndOfMediaEvent) {
219         p.close();
220         System.exit(0);
221     }
222 }
223
224 /**
225  * Main program
226  */
227
228 public static void main(String [] args)
229 {
230
231     if (args.length < 0)
232     {
233         prUsage();
234         System.exit(0);
235     }
236
237     String url = "vfw:Microsoft WDM Image Capture (Win32):0";//args[0];
238
239     if (url.indexOf(".") < 0)
240     {
241         prUsage();
242         System.exit(0);
243     }
244
245     MediaLocator ml;
246
247     if ((ml = new MediaLocator(url)) == null)
248     {
249         System.err.println("Cannot build media locator from: " + url);
250         System.exit(0);
251     }
252
253     FrameAccess fa = new FrameAccess();
254
255     if (!fa.open(ml))
256         System.exit(0);
257 }
258
259 static void prUsage()
260 {
261     System.err.println("Usage: java FrameAccess <url>");
262 }
263
264
265
266
267
268 /*****
269  * Inner class.
270  *
271  * A pass-through codec to access to individual frames.
272  *****/
273
274 public class PreAccessCodec implements Codec {
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
275
276 /**
277  * Callback to access individual video frames.
278  */
279  void accessFrame(Buffer frame)
280  {
281
282      // For demo, we'll just print out the frame #, time &
283      // data length.
284
285      long t = (long)(frame.getTimeStamp()/10000000f);
286
287      System.err.println("Pre: frame #: " + frame.getSequenceNumber() +
288          ", time: " + ((float)t)/100f +
289          ", len: " + frame.getLength());
290      if(frameThrottlingCounter % 10 == 0)
291      {
292          BufferToImage stopBuffer = new BufferToImage((VideoFormat)frame.getFormat());
293          Image stopImage = stopBuffer.createImage(frame);
294          //File fOut = new File("VideoFrameBuffer/"+counter+".jpg");
295          counter++;
296          try {
297              Thread.sleep(400);
298              //1000 millisecs
299          }
300          catch(InterruptedException ex) {
301          }
302
303          try
304          {
305              //Possibly no video being obtained here
306              System.out.println("Create buffered image");
307              BufferedImage outImage = new BufferedImage(640,480,BufferedImage.TYPE_INT_RGB);
308              //FrameHandler.addImageToQueue(outImage);
309              //if(outImage!=null)
310              //System.out.println("hi4");
311              System.out.println("Get graphics");
312              Graphics og = outImage.getGraphics();
313              //FrameHandler.addImageToQueue(outImage);
314              System.out.println("write image to Image");
315              og.drawImage(stopImage, 0, 0, 640, 480, null);
316              //System.out.println("Width: "+outImage.getWidth()+"Height: "+outImage.getHeight());
317              System.out.println("Add image to stack");
318              FrameHandler.addImageToQueue(outImage);
319              System.out.println("image added to queue");
320              FrameHandler.getStackSize();
321              //ImageIO.write(outImage,"jpg",fOut);
322          }
323          catch (Exception ie)
324          {
325              System.out.println("get video frame failed");
326              System.out.println("Error :"+ ie);
327              System.out.println(ie.getMessage());
328              System.out.println(ie.getStackTrace());
329          }
330      }
331
332  }
333
334
335  /**
336   * The code for a pass through codec.
337   */
338
339  // We'll advertize as supporting all video formats.
340  protected Format supportedIns[] = new Format []
341  {
342      new VideoFormat(null)
343  };
344  public boolean clearFrameBuffer()
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
345     {
346         boolean buffClear = false;
347         return buffClear;
348     }
349
350     // We'll advertize as supporting all video formats.
351     protected Format supportedOuts[] = new Format []
352     {
353         new VideoFormat(null)
354     };
355
356     Format input = null, output = null;
357
358     public String getName()
359     {
360         return "Pre-Access Codec";
361     }
362
363     // No op.
364     public void open()
365     {
366     }
367
368     // No op.
369     public void close()
370     {
371     }
372
373     // No op.
374     public void reset()
375     {
376     }
377
378     public Format [] getSupportedInputFormats()
379     {
380         return supportedIns;
381     }
382
383     public Format [] getSupportedOutputFormats(Format in)
384     {
385         if (in == null)
386             return supportedOuts;
387         else
388         {
389             // If an input format is given, we use that input format
390             // as the output since we are not modifying the bit stream
391             // at all.
392             Format outs[] = new Format[1];
393             outs[0] = in;
394             return outs;
395         }
396     }
397
398     public Format setInputFormat(Format format)
399     {
400         input = format;
401         return input;
402     }
403
404     public Format setOutputFormat(Format format)
405     {
406         output = format;
407         return output;
408     }
409
410     public int process(Buffer in, Buffer out)
411     {
412
413         // This is the "Callback" to access individual frames.
414         accessFrame(in);
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
415
416 // Swap the data between the input & output.
417 Object data = in.getData();
418 in.setData(out.getData());
419 out.setData(data);
420
421 // Copy the input attributes to the output
422 out.setFormat(in.getFormat());
423 out.setLength(in.getLength());
424 out.setOffset(in.getOffset());
425
426 return BUFFER_PROCESSED_OK;
427 }
428
429 public Object[] getControls()
430 {
431     return new Object[0];
432 }
433
434 public Object getControl(String type)
435 {
436     return null;
437 }
438 }
439
440 public class PostAccessCodec extends PreAccessCodec
441 {
442     // We'll advertize as supporting all video formats.
443     public PostAccessCodec()
444     {
445         supportedIns = new Format []
446         {
447             new RGBFormat()
448         };
449     }
450
451     /**
452     * Callback to access individual video frames.
453     */
454     void accessFrame(Buffer frame)
455     {
456
457         // For demo, we'll just print out the frame #, time &
458         // data length.
459
460         long t = (long)(frame.getTimeStamp()/100000000f);
461
462         System.err.println("Post: frame #: " + frame.getSequenceNumber() +
463             ", time: " + ((float)t)/100f +
464             ", len: " + frame.getLength());
465     }
466
467     public String getName() {
468         return "Post-Access Codec";
469     }
470 }
471 }
```

Figure 8.3-12 - FrameAccess.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.13 FaceDetection.java class source code

```
1  class JNIOpenCV {
2      static {
3          System.loadLibrary("JNI2OpenCV");
4      }
5      public native int[] detectFace(int minFaceWidth, int minFaceHeight, String cascade, String
6  filename);
7  }
8
9  public class FaceDetection {
10     private JNIOpenCV myJNIOpenCV;
11     private FaceDetection myFaceDetection;
12     private int[] detectedFaces;
13
14     public FaceDetection(String file, int cascadeXML, int faceHeight, int faceWidth) {
15         myJNIOpenCV = new JNIOpenCV();
16         String filename = file;
17         String cascade = "";
18         switch (cascadeXML)
19         {
20             case 1:
21                 cascade = "Cascades/haarcascade_frontalface_alt.xml";
22                 break;
23             case 2:
24                 cascade = "Cascades/haarcascade_frontalface_alt1.xml";
25                 break;
26             case 3:
27                 cascade = "Cascades/haarcascade_frontalface_alt2.xml";
28                 break;
29             case 4:
30                 cascade = "Cascades/haarcascade_frontalface_alt_tree.xml";
31                 break;
32             case 5:
33                 cascade = "Cascades/haarcascade_frontalface_default.xml";
34                 break;
35             case 6:
36                 cascade = "Cascades/haarcascade_fullbody.xml";
37                 break;
38             case 7:
39                 cascade = "Cascades/haarcascade_lowerbody.xml";
40                 break;
41             case 8:
42                 cascade = "Cascades/haarcascade_profileface.xml";
43                 break;
44             case 9:
45                 cascade = "Cascades/haarcascade_upperbody.xml";
46                 break;
47             default:
48                 cascade = "Cascades/haarcascade_frontalface_alt.xml";
49                 break; }
50         detectedFaces = myJNIOpenCV.detectFace(faceHeight, faceWidth, cascade, filename);
51         int numFaces = detectedFaces.length / 4;
52
53         System.out.println("numFaces = " + numFaces);
54         for (int i = 0; i < numFaces; i++) {
55             System.out.println("Face " + i + ": " + detectedFaces[4 * i + 0] + " " +
56 detectedFaces[4 * i + 1] + " " + detectedFaces[4 * i + 2] + " " + detectedFaces[4 * i + 3]);
57         } }
58     public int[] getFaceMatrix()
59     {
60         return detectedFaces; }
61     public static void main(String args[]) {
62         String fileIs = "C:/Documents and
63 Settings/Martyn/MyDocuments/Eclipse/FaceRecognitionProj/trainPhotos/lena.jpg";
64         FaceDetection myFaceDetection = new FaceDetection(fileIs,40,40,1);
65
66     }
```


67

```
}
```

Figure 8.3-13 - FaceDetection.java class

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.14 FeatureLocalisation.java original class source code

```
1  import java.awt.*;
2  import java.awt.event.*;
3  import java.util.*;
4  import java.awt.image.*;
5  import javax.imageio.*;
6  import java.io.*;
7  import java.lang.Math;
8  import javax.swing.JFrame;
9
10 public class FeatureLocalisation extends JFrame
11 {
12     private Image featuresFace;
13     private BufferedImage featuresFaceBuff;
14     private double[] meanEyeArrayRight;
15     private double[] meanEyeArrayLeft;
16     private double[][] sinvArrayRight;
17     private double[][] sinvArrayLeft;
18     private static double sinRight = 5.6276e-103;
19     private static double sinLeft = 5.2953e-103;
20     private static double logSinRight = Math.log(sinRight);
21     private static double logSinLeft = Math.log(sinLeft);
22     private ArrayList<EyePatch> patchListLeftEye;
23     private ArrayList<EyePatch> patchListRightEye;
24     private int x1;
25     private int y1;
26     private int x2;
27     private int y2;
28
29     /**
30      * setMeanEyeArray
31      *
32      * File from Matlab containing mean eye vector is read into Java and stored in global
33 variable patchEyeList Left or Right
34      *
35      * @param filename .dat file containing mean eye values (one per line)
36      * @param array array to read the values into
37      */
38     public int getXPos1()
39     {
40         return x1;
41     }
42     public int getYPos1()
43     {
44         return y1;
45     }
46     public int getXPos2()
47     {
48         return x2;
49     }
50     public int getYPos2()
51     {
52         return y2;
53     }
54     public void setMeanEyeArray(String filename, double[] array)
55     {
56         //array = new double[256];
57         try
58         {
59             BufferedReader in = new BufferedReader(new FileReader("trainData\\"+filename));
60             String str;
61             int counter = 0;
62             while ((str = in.readLine()) != null)
63             {
64                 double d = Double.valueOf(str.trim()).doubleValue();
65                 array[counter] = d;
66                 counter++;
67             }
68             in.close();
69         }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
69     }
70     catch (IOException e)
71     {
72         System.out.println(e.getMessage());
73     }
74 }
75
76 public void setSinvArray(String filename, int patchSize, double[][] sinvArray)
77 {
78     sinvArray = new double[patchSize*patchSize][patchSize*patchSize];
79     try
80     {
81         //int resultLength = 0;
82         BufferedReader in = new BufferedReader(new FileReader("trainData\\"+filename));
83         String str;
84         int counter = 0;
85         while ((str = in.readLine()) != null)
86         {
87             String[] result = str.split(" ");
88             for(int token=0; token<result.length; token++)
89             {
90                 int lineCounter = 0;
91                 if(!result[token].equals(""))
92                 {
93                     //System.out.println(result[token]);
94                     double d =
95 Double.valueOf(result[token].trim()).doubleValue();
96                     sinvArray[counter][lineCounter] = d;
97                     lineCounter++;
98                 }
99             }
100             counter++;
101         }
102         //System.out.println(resultLength);
103         in.close();
104     }
105     catch (IOException e)
106     {
107         System.out.println(e.getMessage());
108     }
109 }
110 public static BufferedImage toBufferedImage(Image image) {
111     if (image instanceof BufferedImage)
112         return (BufferedImage) image;
113     ColorModel cm = getColorModel(image);
114     int width = image.getWidth(null);
115     int height = image.getHeight(null);
116     return copy(createBufferedImage(cm, width, height), image);
117 }
118 public static BufferedImage createBufferedImage(ColorModel cm, int w, int h) {
119     WritableRaster raster = cm.createCompatibleWritableRaster(w, h);
120     boolean isRasterPremultiplied = cm.isAlphaPremultiplied();
121     return new BufferedImage(cm, raster, isRasterPremultiplied, null);
122 }
123 public static BufferedImage copy(BufferedImage target, Image source) {
124     Graphics2D g = target.createGraphics();
125     g.drawImage(source, 0, 0, null);
126     g.dispose();
127     return target;
128 }
129 public static ColorModel getColorModel(Image image) {
130     try {
131         PixelGrabber pg = new PixelGrabber(image, 0,0,1,1, false);
132         pg.grabPixels();
133         return pg.getColorModel();
134     } catch (InterruptedException e) {
135         throw new RuntimeException("Unexpected interruption", e);
136     }
137 }
138 public double[] compareVectors(double[] meanEyeV, double[] testPatchVector)
139 {
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
140     double[] compareVector = new double[meanEyeV.length];
141     for(int i=0;i<meanEyeV.length;i++)
142     {
143         compareVector[i] = testPatchVector[i] - meanEyeV[i];
144     }
145     return compareVector;
146 }
147
148
149 public FeatureLocalisation(Image image, String fileName, int patchSize)
150 {
151     patchListLeftEye = new ArrayList<EyePatch>();
152     patchListRightEye= new ArrayList<EyePatch>();
153     meanEyeArrayLeft = new double[patchSize*patchSize];
154     meanEyeArrayRight = new double[patchSize*patchSize];
155     sinvArrayRight = new double[patchSize*patchSize][patchSize*patchSize];
156     sinvArrayLeft = new double[patchSize*patchSize][patchSize*patchSize];
157     setMeanEyeArray("leftEyeMean.dat",meanEyeArrayLeft);
158     setMeanEyeArray("RightEyeMean.dat",meanEyeArrayRight);
159     setSinvArray("SinvRight.dat", patchSize, sinvArrayRight);
160     setSinvArray("SinvLeft.dat", patchSize, sinvArrayLeft);
161     featuresFace = image;
162     featuresFace = image.getScaledInstance(100,100,5);
163
164     Toolkit toolkit = Toolkit.getDefaultToolkit();
165     String file = fileName;
166     if (file.indexOf("\\") < 0)
167     {
168         file = "trainPhotos/" + fileName;
169     }
170     else
171     {
172         file = fileName;
173     }
174     image = toolkit.getImage(file);
175     MediaTracker mediaTracker = new MediaTracker(this);
176     mediaTracker.addImage(featuresFace, 0);
177     try
178     {
179         mediaTracker.waitForID(0);
180     }
181     catch (InterruptedException ie)
182     {
183         System.err.println(ie);
184         System.exit(1);
185     }
186     addWindowListener(new WindowAdapter() { public void windowClosing(WindowEvent e)
187     {
188         //System.exit(0);
189         hide();
190     }});
191     setSize(featuresFace.getWidth(null), featuresFace.getHeight(null));
192     setTitle("Face");
193     show();
194     BufferedImage imgToBuf = toBufferedImage(featuresFace);
195     imageSlidingWindow(imgToBuf,patchSize);
196 }
197 public FeatureLocalisation(BufferedImage image, int patchSize)
198 {
199     System.out.println("Running feature localisation");
200     patchListLeftEye = new ArrayList<EyePatch>();
201     System.out.println("1");
202     patchListRightEye= new ArrayList<EyePatch>();
203     System.out.println("2");
204     meanEyeArrayLeft = new double[patchSize*patchSize];
205     System.out.println("3");
206     meanEyeArrayRight = new double[patchSize*patchSize];
207     System.out.println("4");
208     sinvArrayRight = new double[patchSize*patchSize][patchSize*patchSize];
209     System.out.println("5");
210     sinvArrayLeft = new double[patchSize*patchSize][patchSize*patchSize];
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
211 System.out.println("6");
212 setMeanEyeArray("leftEyeMean.dat",meanEyeArrayLeft);
213 System.out.println("7");
214 setMeanEyeArray("RightEyeMean.dat",meanEyeArrayRight);
215 System.out.println("8");
216 setSinvArray("SinvRight.dat", patchSize, sinvArrayRight);
217 System.out.println("9");
218 setSinvArray("SinvLeft.dat", patchSize, sinvArrayLeft);
219 System.out.println("Arrays setup and running");
220 featuresFaceBuff = image;
221 //feturesFaceBuff = image.getScaledInstance(100,100,5);
222 GraphicsConfiguration gc = getDefaultConfiguration();
223 System.out.println("Running sliding window");
224
225 imageSlidingWindow(getScaledInstance(featuresFaceBuff,100,100,null),patchSize);
226 }
227 public static BufferedImage getScaledInstance(BufferedImage image, int width, int height,
228 GraphicsConfiguration gc) {
229     if (gc == null)
230         gc = getDefaultConfiguration();
231     int transparency = image.getColorModel().getTransparency();
232     return copy(image, gc.createCompatibleImage(width, height, transparency));
233 }
234 public static GraphicsConfiguration getDefaultConfiguration() {
235     GraphicsEnvironment ge = GraphicsEnvironment.getLocalGraphicsEnvironment();
236     GraphicsDevice gd = ge.getDefaultScreenDevice();
237     return gd.getDefaultConfiguration();
238 }
239 public void paint(Graphics graphics) {
240     graphics.drawImage(featuresFace, 0, 0, null);
241     getRectanglesAndDraw(graphics);
242 }
243
244 public double getGreyScaleFromRGBScaled(int red, int green, int blue)
245 {
246     double grayValue = 0.0;
247     double redValue = 0.3;
248     double greenValue = 0.59;
249     double blueValue = 0.11;
250     grayValue = (redValue*red) + (greenValue*green) + (blueValue*blue);
251     return (grayValue/256);
252 }
253 public double[][] vectorToMatrix(double[] vector)
254 {
255     int matrixDimension = (int)Math.sqrt(vector.length);
256     double[][] reconstructedMatrix = new double[matrixDimension][matrixDimension];
257     int columnCounter = 0;
258     int rowCounter = 0;
259     for(int i=0;i<vector.length;i++)
260     {
261         reconstructedMatrix[rowCounter][columnCounter] = vector[i];
262         if(rowCounter == matrixDimension-1)
263         {
264             rowCounter = 0;
265             columnCounter++;
266         }
267         else
268         {
269             rowCounter++;
270         }
271     }
272     return reconstructedMatrix;
273 }
274 public void drawRectangles(Graphics graphics)
275 {
276     getRectanglesAndDraw(graphics);
277 }
278 public void drawRectangle(Graphics graphic, int x, int y, int z)
279 {
280     graphic.drawRect(x, y, z, z);
281 }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
282 public void getRectanglesAndDraw(Graphics graphicPass)
283 {
284     drawRectangle(graphicPass, y1, x1, 16);
285 }
286
287
288 // bubbleSort - sorts, in place, the items in the ArrayList 'entries'
289 public void bubbleSort(ArrayList<EyePatch> faceArrayUnsorted)
290 {
291     boolean swapped=true;
292     for (int i = faceArrayUnsorted.size()-1; i >= 0 && swapped; i--)
293     {
294         swapped = false;
295         for (int j = 0; j < i; j++)
296         {
297             if (faceArrayUnsorted.get(j+1).getLogLikelihoodVal() <
298 (faceArrayUnsorted.get(j).getLogLikelihoodVal()))
299             {
300                 EyePatch tmp = faceArrayUnsorted.get(j+1);
301                 faceArrayUnsorted.set(j+1, faceArrayUnsorted.get(j));
302                 faceArrayUnsorted.set(j, tmp);
303                 swapped = true;
304             }
305         }
306     }
307 }
308 public void imageSlidingWindow(BufferedImage bufferedFace, int patchSize)
309 {
310
311     int height = bufferedFace.getHeight(null);
312     int width = bufferedFace.getWidth(null);
313     int lastValueX = width - patchSize;
314     System.out.println("Maxiumu X Value of image:\t"+lastValueX);
315     System.out.println("Running sliding window");
316     //Value here is divided by two to remove the bottom half of the image from the
317 search space
318     int startValueY = height/2;//84;//height;
319     //The last value is in effect close to the top because of the x,y coordinates
320     int lastValueY = height-patchSize-1;//(height/2) - patchSize;
321     //Load face into buffer from disk
322     int patchIndexCounter = 0;
323     for(int y=startValueY;y<lastValueY;y++)
324     {
325         for(int x=0;x<lastValueX;x++)
326         {
327             //Array of pixels to represent image patch in colour
328             int[][] pixels = new int[patchSize][patchSize];
329             double[][] grayPixels = new double[patchSize][patchSize];
330             for(int row=0;row<patchSize;row++)
331             {
332                 //Extract rows from the image patch and load them into the RGB
333 array above
334                 bufferedFace.getRGB(x, /*row+(y-patchSize)*/y+row,
335 patchSize, 1, pixels[row], 0, patchSize);
336             }
337             for(int i=0;i<patchSize;i++)
338             {
339                 //System.out.println();
340                 for(int j=0;j<patchSize;j++)
341                 {
342                     //int alpha = 0;
343                     int red = 0;
344                     int green = 0;
345                     int blue = 0;
346                     //alpha = (pixels[i][j]>>24)&255;
347                     red = (pixels[i][j]>>16)&255;
348                     green = (pixels[i][j]>>8) &255;
349                     blue = pixels[i][j] &255;
350                     double grayScaleValue =
351 getGreyScaleFromRGBScaled(red,green,blue);
352                     grayPixels[i][j] = grayScaleValue;
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
353                                     //System.out.println("Grayscale value:
354 "+grayScaleValue);
355
356                                     }
357                                     }
358                                     double[] testEyePatchVector = new double[patchSize*patchSize];
359                                     //Counter to create the vector
360                                     int vectorCounter = 0;
361                                     for(int colV =0;colV<patchSize;colV++)
362                                     {
363                                         for(int rowV=patchSize-1;rowV>-1;rowV--)
364                                         {
365                                             testEyePatchVector[vectorCounter] =
366                                             grayPixels[rowV][colV];
367                                             vectorCounter++;
368                                         }
369                                     }
370                                     double[] eyePatchCompVectorRight =
371                                     compareVectors(meanEyeArrayRight,testEyePatchVector);
372                                     double[] eyePatchCompVectorLeft = compareVectors(meanEyeArrayLeft,
373                                     testEyePatchVector);
374                                     Matrix SinvMatrixRight = new Matrix(sinvArrayRight);
375                                     Matrix SinvMatrixLeft = new Matrix(sinvArrayLeft);
376                                     double[] newColumnRight =
377                                     Matrix.columnByMatrixMultiply(SinvMatrixRight,eyePatchCompVectorRight);
378                                     double[] newColumnLeft =
379                                     Matrix.columnByMatrixMultiply(SinvMatrixLeft, eyePatchCompVectorLeft);
380                                     double colRowVecRight =
381                                     Matrix.colVecMultRowVec(eyePatchCompVectorRight, newColumnRight);
382                                     double colRowVecLeft =
383                                     Matrix.colVecMultRowVec(eyePatchCompVectorLeft, newColumnLeft);
384
385                                     EyePatch thisPatchLeft = new
386                                     EyePatch(patchIndexCounter,x,y,colRowVecRight,patchSize);
387                                     EyePatch thisPatchRight= new
388                                     EyePatch(patchIndexCounter,x,y,colRowVecLeft,patchSize);
389                                     patchIndexCounter++;
390                                     patchListLeftEye.add(thisPatchLeft);
391                                     patchListRightEye.add(thisPatchRight);
392                                     double[][] compMatrixRight =
393                                     vectorToMatrix(eyePatchCompVectorRight);
394                                     double[][] compMatrixLeft =
395                                     vectorToMatrix(eyePatchCompVectorLeft);
396                                     Matrix meanCompRight = new Matrix(compMatrixRight);
397                                     Matrix meanCompLeft = new Matrix(compMatrixLeft);
398                                     Matrix compMatrixTransposeRight = meanCompRight.transpose();
399                                     Matrix compMatrixTransposeLeft = meanCompLeft.transpose();
400                                     Matrix meanCompMultSinvRight =
401                                     SinvMatrixRight.times(meanCompRight);
402                                     Matrix meanCompMultSinvLeft = SinvMatrixLeft.times(meanCompLeft);
403                                     }
404                                     }
405
406                                     int highIndex = 0;
407                                     int highIndex2= 0;
408                                     double logVal = 100.0;
409
410                                     int xPos = 0;
411                                     int yPos = 0;
412                                     int xPos2 = 0;
413                                     int yPos2 = 0;
414
415                                     System.out.println("Right eye position");
416                                     logVal = 100.0;
417                                     for(int i=0;i<patchListRightEye.size();i++)
418                                     {
419                                         EyePatch testPatch2 = patchListRightEye.get(i);
420                                         //System.out.println("XPosition of patch:\t"+Math.abs(testPatch2.getXPos()
421                                         - x1));
422                                         if(testPatch2.getLogLikelihoodVal() < logVal/*
423                                         &&(Math.abs(testPatch2.getXPos() - x1) > (patchSize))*/)

```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
424         {
425             //if(Math.abs(testPatch2.getXPos() - x1) > patchSize)
426             {
427                 logVal = testPatch2.getLogLikelihoodVal();
428                 highIndex2 = testPatch2.getIndex();
429                 xPos = testPatch2.getXPos();
430                 yPos = testPatch2.getYPos();
431                 System.out.print(testPatch2.getXPos());
432             }
433         }
434     }
435     x1 = xPos;
436     y1 = yPos;
437     for(int i=0;i<patchListLeftEye.size();i++)
438     {
439         EyePatch testPatch = patchListLeftEye.get(i);
440         if(testPatch.getLogLikelihoodVal() < logVal)
441         {
442             if(Math.abs(testPatch.getXPos() - x1) > patchSize)
443             {
444                 logVal = testPatch.getLogLikelihoodVal();
445                 highIndex = testPatch.getIndex();
446                 xPos2 = testPatch.getXPos();
447                 yPos2 = testPatch.getYPos();
448                 System.out.print(xPos+"\t");
449             }
450         }
451     }
452     x2 = xPos2;
453     y2 = yPos2;
454     System.out.println("X1: "+x1+"\tY1: "+y1+"\tX2: "+x2+"\tY2: "+y2);
455     System.out.println("Sliding window applied!");
456 }
457 }
```

Figure 8.3-14 - FeatureLocalisation.java Source code (original)

8.3.15 *FrameHandler.java* class source code

```
1  import java.awt.FlowLayout;
2  import java.awt.Graphics;
3  import java.awt.GridLayout;
4  import java.awt.MediaTracker;
5  import java.awt.Toolkit;
6  import java.awt.image.BufferedImage;
7  import java.util.ArrayList;
8  import java.io.*;
9  import javax.imageio.ImageIO;
10 import javax.swing.JButton;
11 import javax.swing.JLabel;
12 import javax.swing.JPanel;
13 import javax.swing.ImageIcon;
14
15 import java.awt.event.WindowAdapter;
16 import java.awt.event.WindowEvent;
17 import java.awt.geom.*;
18 import java.util.TimerTask;
19
20
21 /**
22  * FrameHandler.java
23  * @author Martyn Booth
24  * Implements a new thread for that the infinite loop used here does not
25  */
26
27 public class FrameHandler extends Thread//implements Runnable
28 {
29     private static WorkQueue imageQueue;
30     private ArrayList faces;
31     private FaceLocations locations;
32     private int counter = 0;
33     private JPanel matchingDisplay;
34     private int[] rightEyeX;
35     private int[] rightEyeY;
36     private int[] leftEyeX;
37     private int[] leftEyeY;
38     private ArrayList<Face> referenceFaces;
39     private FaceDataAccessLayer dal;
40     private JLabel forenameDisplay;
41     private JLabel surnameDisplay;
42     private JLabel imageIconHost;
43     private JPanel dbImage;
44     private BufferedImage imgFromDB;
45     private static Face currentMatchedFace;
46     private java.awt.Image image;
47
48     public FrameHandler()
49     {
50         imageIconHost = new JLabel();
51         imageQueue = new WorkQueue();
52     }
53     public void run()
54     {
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
55         doLiveMatchingScheme();
56     }
57     public static Face getMatchedFaceFace()
58     {
59         return currentMatchedFace;
60     }
61     public void doLiveMatchingScheme()
62     {
63         for(;;)
64         {
65             dal = new FaceDataAccessLayer();
66             obtainReferenceFaces();
67             int numberOfFrames = 5;
68             rightEyeX = new int[numberOfFrames];
69             rightEyeY = new int[numberOfFrames];
70             leftEyeX = new int[numberOfFrames];
71             leftEyeY = new int[numberOfFrames];
72             FeatureLocalisation feature;
73             int counter2 = 0;
74             for(int i = 0; i < numberOfFrames; i++)
75             {
76                 String path = "";
77                 BufferedImage vidFrame = new BufferedImage(20,20,BufferedImage.TYPE_INT_RGB);
78                 File fOut = new File("Dummy");
79                 try
80                 {
81
82                     vidFrame = (BufferedImage)imageQueue.getWork();
83                     System.out.println("Got work from queue");
84                     Graphics og = vidFrame.getGraphics();
85                     path = "VideoFrameBuffer/"+counter+".jpg";
86
87                     fOut = new File(path);
88                     ImageIO.write(vidFrame, "jpg", fOut);
89                     String newPath = "C:\\Documents and Settings\\Martyn\\My
90 Documents\\Eclipse\\FaceRecognitionProj\\VideoFrameBuffer\\"+counter+".jpg";
91                     System.out.println("Get Face locations");
92                     getFaceLocations(newPath);
93                     System.out.println("intialise array");
94                     setUpArray();
95                     System.out.println("Get Features");
96                     feature = createFacePatchesFromArray(vidFrame,16);
97                     System.out.println("Get face patches complete");
98                     rightEyeX[counter2] = feature.getXPos1();
99                     rightEyeY[counter2] = feature.getYPos1();
100                     leftEyeX[counter2] = feature.getXPos2();
101                     leftEyeY[counter2] = feature.getYPos2();
102                     boolean success = fOut.delete();
103                     if(!success)
104                     {
105                         System.out.println("Temporary file not deleted correctly");
106                     }
107
108                 }
109             }
110             catch(Exception e)
111             {
112                 System.out.println("Remove video frame from work queue failed: "+e.getMessage());
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
113         }
114         System.out.println("Continuing image access");
115         counter++;
116
117         }System.out.println("iteration complete");
118         int rightEyeXInt = averageOfArrayValues(rightEyeX);
119         int rightEyeYInt = averageOfArrayValues(rightEyeY);
120         int leftEyeXInt = averageOfArrayValues(leftEyeX);
121         int leftEyeYInt = averageOfArrayValues(leftEyeY);
122         System.out.println("arrays averaged");
123         Point2D testRightEye = new Point2D.Double(rightEyeXInt, rightEyeYInt);
124         Point2D testLeftEye = new Point2D.Double(leftEyeXInt, leftEyeYInt);
125         double testImgEucDist = testRightEye.distance(testLeftEye);
126         int matchedFace = 0;
127         for(int i = 0; i < faces.size(); i++)
128     {
129         Face face = (Face)faces.get(i);
130         double faceEucDistTest = face.getEuclideanDistance();
131         if(Math.abs(testImgEucDist - faceEucDistTest) < Math.abs(testImgEucDist - face.getEuclideanDistance()))
132         {
133             matchedFace = i;
134         }
135     }
136     currentMatchedFace = (Face)faces.get(matchedFace);
137     System.out.println(currentMatchedFace.getForename());
138     Face face = FrameHandler.getMatchedFaceFace();
139     FaceRecogGUI.forenameDisplay.setText(face.getForename());
140     FaceRecogGUI.surnameDisplay.setText(face.getSurname());
141     BufferedImage buffImg = FeatureLocalisation.toBufferedImage(face.getFaceImage());
142     FaceRecogGUI.imageLabel = buffImageToImageIcon(buffImg);
143     FaceRecogGUI.matchDetails.repaint();
144
145     }
146 }
147 /**
148  * Test method to check on the presence of images from the frame splitter
149  * @param fileName name of the image being tested
150  */
151 public void drawFrame(String fileName)
152 {
153     javax.swing.JFrame frame = new javax.swing.JFrame("Frame image");
154     Toolkit toolkit = Toolkit.getDefaultToolkit();
155     String file = "*/trainPhotos/" + fileName;
156     file = "C:\\Documents and Settings\\Martyn\\My Documents\\Eclipse\\FaceRecognitionProj\\VideoFrameBuffer\\0.jpg";
157     image = toolkit.getImage(file);
158     MediaTracker mediaTracker = new MediaTracker(frame);
159     mediaTracker.addImage(image, 0);
160
161     try
162     {
163         mediaTracker.waitForID(0);
164     }
165     catch (InterruptedException ie)
166     {
167         System.err.println(ie);
168         System.exit(1);
169     }
170     frame.addWindowListener(new WindowAdapter() { public void windowClosing(WindowEvent e)
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
171         {
172             System.exit(0);
173         }
174     });
175
176     frame.setSize(image.getWidth(null), image.getHeight(null));
177     frame.setTitle(fileName);
178     frame.setVisible(true);
179
180 }
181 public void paint(Graphics graphics)
182 {
183     graphics.drawImage(image, 0, 0, null);
184 }
185 /**
186  * Create the panel for displaying the matching face details and return to main GUI
187  * @return JPanel containing relevant matching details
188  */
189 public JPanel createMatchingDisplay()
190 {
191     //Create the display for showing the face that has been matched to the video
192     JPanel matchContainer = new JPanel(new GridLayout(0,1));
193     dbImage = new JPanel(new FlowLayout());
194     JPanel matchFieldsContainer = new JPanel(new FlowLayout());
195     JPanel labelPanel = new JPanel(new GridLayout(0,1));
196     JPanel answerPanel = new JPanel(new GridLayout(0,1));
197     JLabel forenameLabel = new JLabel("Forename: ");
198     JLabel surnameLabel = new JLabel("Surname: ");
199     //dbImage panel for showing image from database
200     String name = "";
201     String surname = "";
202     BufferedImage buffImg = new BufferedImage(20,20,BufferedImage.TYPE_INT_RGB);
203     if(faces!=null)
204     {
205         Face face = currentMatchedFace;
206         name = face.getForename();
207         surname = face.getSurname();
208         buffImg = FeatureLocalisation.toBufferedImage(face.getFaceImage());
209     }
210     //dbImage.add(imageIconHost);
211     labelPanel.add(forenameLabel);
212     labelPanel.add(surnameLabel);
213     forenameDisplay = new JLabel(name);
214     surnameDisplay = new JLabel(surname);
215     answerPanel.add(forenameDisplay);
216     answerPanel.add(surnameDisplay);
217     matchFieldsContainer.add(answerPanel);
218     matchFieldsContainer.add(labelPanel);
219     JLabel imageLabel = buffImageToImageIcon(buffImg);
220     dbImage.add(imageLabel);
221     matchContainer.add(dbImage);
222     matchContainer.add(matchFieldsContainer);
223     return matchContainer;
224 }
225 public JLabel buffImageToImageIcon(BufferedImage buffImg)
226 {
227     JLabel imageHost = new JLabel(new ImageIcon(buffImg));
228 }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
229         return imageHost;
230     }
231     private int averageOfArrayValues(int[] array)
232     {
233         int average = 0;
234         int temp = 0;
235         for(int i = 0; i<array.length;i++)
236         {
237             temp += array[i];
238         }
239         double averageNum = (double)temp/array.length;
240         average = (int)averageNum;
241         return average;
242     }
243     public boolean obtainReferenceFaces()
244     {
245         boolean isSuccess = false;
246         referenceFaces = dal.getFaceArray();
247         isSuccess = true;
248         return isSuccess;
249     }
250 }
251 public FeatureLocalisation createFacePatchesFromArray(BufferedImage image, int patchSize)
252 {
253     FeatureLocalisation feature;
254     //for(int i=0; i < 1;i++)//faces.size();i++)
255     //{
256         System.out.println("Get face location");
257         FaceLocation face = (FaceLocation)faces.get(0);//i);
258         System.out.println("Got face locaiton 1");
259         System.out.println(face.toString());
260         int x = face.getX();
261         System.out.println(x);
262
263         int y = face.getY();
264         System.out.println(y);
265         int w = face.getWidth();
266         System.out.println(w);
267         System.out.println("crop image");
268         FaceCrop crop = new FaceCrop(image,x,y,w);
269         System.out.println("get features");
270         feature = new FeatureLocalisation(crop.getCroppedBuflmg(),patchSize);
271         System.out.println("return features");
272     //}
273     return feature;
274 }
275 public static synchronized void addImageToQueue(BufferedImage img)
276 {
277     imageQueue.addWork(img);
278 }
279 public static synchronized void getStackSize()
280 {
281     System.out.println(imageQueue.getSize());
282 }
283 public void setUpArray()
284 {
285     faces = locations.getFaceObjects();
286 }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
287 public void getFaceLocations(String fileName)
288 {
289     System.out.println("Getting face locations");
290     locations = new FaceLocations(fileName,1,40,40);
291     System.out.println(locations.getFaceObjects().get(0));
292 }
293 }
```

Figure 8.3-15 - FrameHandler.java source

8.3.16 *WorkQueue.java* class

```
1 import java.util.LinkedList;
2 import java.awt.image.BufferedImage;
3
4 public class WorkQueue {
5     LinkedList<BufferedImage> queue = new LinkedList<BufferedImage>();
6
7     // Add work to the work queue
8     public synchronized void addWork(BufferedImage img) {
9         queue.addLast(img);
10        notifyAll();
11    }
12
13    // Retrieve work from the work queue; block if the queue is empty
14    public synchronized Object getWork() throws InterruptedException {
15        while (queue.isEmpty()) {
16            System.out.println("Getting work from queue");
17            wait();
18        }
19        return queue.removeFirst();
20    }
21    public int getSize()
22    {
23        return queue.size();
24    }
25 }
```

Figure 8.3-16 - *WorkQueue.java* source

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.17 FaceRecogGUI.java class (Main program entry point)

```
1  import java.awt.*;
2  import java.awt.event.*;
3
4  import javax.swing.AbstractAction;
5
6  import javax.swing.ImageIcon;
7  import javax.swing.JPanel;
8  import javax.swing.JTextArea;
9  import javax.swing.JScrollPane;
10 import javax.swing.JFrame;
11 import javax.swing.JLabel;
12 import javax.swing.JButton;
13 import java.awt.image.*;
14 import java.io.*;
15 import java.util.Timer;
16 import javax.security.auth.Refreshable;
17
18 /**
19  * FaceRecogGUI.java
20  * Author: Martyn Booth
21  * */
22
23 /*
24  * Main class, provides GUI functionality for all of the other classes and acts as a glue
25  */
26 public class FaceRecogGUI implements ActionListener, ItemListener
27 {
28     JTextArea output;
29     JScrollPane scrollPane;
30     private static JFrame frame;
31     private static JLabel forenameDisplay;
32     private static JLabel surnameDisplay;
33     private static JLabel imageIconHost;
34     private static JLabel imageLabel;
35     private static ImageIcon imgIcon;
36     private JPanel dbImage;
37     private BufferedImage imgFromDB;
38     private static JPanel matchDetails;
39
40     public FaceRecogGUI ()
41     {
42         imageLabel = new JLabel();
43     }
44     public MenuBar createMenuBar() {
45
46         MenuBar menuBar;
47         Menu menu;
48         MenuItem menuItem;
49         menuBar = new MenuBar();
50         menu = new Menu("File");
51         menuBar.add(menu);
52         menuItem = new MenuItem("Add new Reference Face");
53         menu.add(menuItem);
54         // ...for each MenuItem instance:
55         menuItem.addActionListener(new ActionListener()
56         {
57             public void actionPerformed(ActionEvent event)
58             {
59                 FaceAddToDatabase add = new FaceAddToDatabase();
60             }
61         });
62     }
63
64
65 }
```


FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
66     menuItem = new JMenuItem("Remove Reference Face");
67     menu.add(menuItem);
68     // ...for each JMenuItem instance:
69     menuItem.addActionListener(new ActionListener()
70     {
71         public void actionPerformed(ActionEvent event)
72         {
73             FaceRemoveFromDatabase add = new FaceRemoveFromDatabase();
74             }
75         }
76     });
77
78     return menuBar;
79
80 }
81
82
83 public Container createContentPane() {
84     //Create the content-pane-to-be.
85     JPanel contentPane = new JPanel(new BorderLayout());
86     contentPane.setOpaque(true);
87
88     //Create a selection of panes and layouts to support the visual components
89     JPanel vidPane = new JPanel(new BorderLayout());
90     JPanel vidPaneContainer = new JPanel(new BorderLayout());
91     vidPane = createAndReturnVideo();
92     vidPane.setSize(450, 400);
93     JPanel vidPaneName = new JPanel(new BorderLayout());
94     vidPaneName.setSize(450, 200);
95     JLabel vidLabel = new JLabel();
96     vidLabel.setText("Live Video Feed:");
97     vidPaneName.add(vidLabel, BorderLayout.CENTER);
98     vidPaneContainer.add(vidPaneName, BorderLayout.NORTH);
99     vidPaneContainer.add(vidPane, BorderLayout.SOUTH);
100
101     matchDetails = new JPanel(new BorderLayout(), true);
102     matchDetails = createMatchingDisplay();
103
104
105     //Add the panels to the content pane.
106     contentPane.add(vidPaneContainer);
107
108     contentPane.add(matchDetails);
109
110     return contentPane;
111 }
112
113
114 public JPanel createAndReturnVideo()
115 {
116     //Get the video source from the webcam and display on screen
117     JPanel videoPanel = new JPanel(new BorderLayout());
118     videoPanel.setSize(450, 400);
119     FrameAccess access = new FrameAccess();
120     access.addVideoToPanel("");
121     videoPanel = access.getPanel();
122     videoPanel.setOpaque(true);
123     videoPanel.setVisible(true);
124
125     return videoPanel;
126 }
127 public void paint(Graphics graphics)
128 {
129 }
130 public void actionPerformed(ActionEvent e){}
131 public void itemStateChanged(ItemEvent e){}
132
133 /**
134  * Create the GUI and show it. For thread safety,
135  * this method should be invoked from the
136  * event-dispatching thread.
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
137  */
138  /*public JPanel createMatchingDisplay()
139  {
140      FrameHandler handler = new FrameHandler();
141      handler.start();
142      matchDetails = handler.createMatchingDisplay();
143      JButton exitButton = new JButton("Exit");
144      exitButton.addActionListener(new ExitApplicationClick());
145      matchDetails.add(exitButton);
146      return matchDetails;
147  }*/
148  private static void createAndShowGUI() {
149      //Create and set up the window.
150      JFrame frame = new JFrame("Face Recognition GUI");
151      frame.setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
152
153      //Create and set up the content pane.
154      FaceRecogGUI demo = new FaceRecogGUI();
155
156      frame.setContentPane(demo.createContentPane());
157      frame.setMenuBar(demo.createMenuBar());
158
159      //Display the window.
160      frame.pack();
161      frame.setVisible(true);
162  }
163  private class ExitApplicationClick extends AbstractAction
164  {
165      public ExitApplicationClick()
166      {
167          super("Disable frame");
168      }
169
170      public void actionPerformed(ActionEvent e)
171      {
172          frame.setVisible(false);
173          System.exit(0);
174      }
175  }
176  public JPanel createMatchingDisplay()
177  {
178      FrameHandler handler = new FrameHandler();
179      handler.start();
180      //Create the display for showing the face that has been matched to the video
181      JPanel matchContainer = new JPanel(new GridLayout(0,1));
182      dbImage = new JPanel(new FlowLayout());
183      JPanel matchFieldsContainer = new JPanel(new FlowLayout());
184      JPanel labelPanel = new JPanel(new GridLayout(0,1));
185      JPanel answerPanel = new JPanel(new GridLayout(0,1));
186      JLabel forenameLabel = new JLabel("Forename: ");
187      JLabel surnameLabel = new JLabel("Surname: ");
188      //dbImage panel for showing image from database
189      String name = "";
190      String surname = "";
191      //BufferedImage buffImg = new BufferedImage(20,20,BufferedImage.TYPE_INT_RGB);
192      //dbImage.add(imageIconHost);
193      labelPanel.add(forenameLabel);
194      labelPanel.add(surnameLabel);
195      forenameDisplay = new JLabel(name);
196      surnameDisplay = new JLabel(surname);
197      imgIcon = new ImageIcon();
198      imageLabel = new JLabel(imgIcon);
199      dbImage.add(imageLabel);
200      answerPanel.add(forenameDisplay);
201      answerPanel.add(surnameDisplay);
202      matchFieldsContainer.add(labelPanel);
203      matchFieldsContainer.add(answerPanel);
204      JButton exitButton = new JButton("Exit");
205      exitButton.setSize(50, 25);
206      exitButton.addActionListener(new java.awt.event.ActionListener() {
207          public void actionPerformed(java.awt.event.ActionEvent evt) {
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
208         System.exit(0);
209     }
210 });
211
212     //imageLabel = buffImageToImageIcon(buffImg);
213     dbImage.add(imageLabel);
214     matchContainer.add(dbImage);
215     matchContainer.add(matchFieldsContainer);
216     matchContainer.add(exitButton);
217     return matchContainer;
218 }
219
220 public static void main(String[] args) {
221     //Schedule a job for the event-dispatching thread:
222     //creating and showing this application's GUI.
223
224     javax.swing.SwingUtilities.invokeLater(new Runnable() {
225         public void run() {
226             createAndShowGUI();
227         }
228     });
229 }
230 public static void reinitaliseMatchDisplay(String forename, String surname, BufferedImage
231 buffImg)
232 {
233     forenameDisplay.setText(forename);
234     surnameDisplay.setText(surname);
235     Image image = buffImg.getScaledInstance(200, 150, 1);
236     imgIcon.setImage(image); // = new JLabel(new
237 ImageIcon(buffImg)); //buffImageToImageIcon(buffImg);
238     matchDetails.repaint();
239     frame.pack();
240 }
241 public static JLabel buffImageToImageIcon(BufferedImage buffImg)
242 {
243     JLabel imageHost = new JLabel(new ImageIcon(buffImg));
244
245     return imageHost;
246 }
247
248
249
250 }
```

Figure 8.3-17 - FaceRecogGUI.java source code, main program entry point

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.3.18 Java test classes:

```
1 import java.awt.*;
2 import java.awt.event.*;
3 import java.util.*;
4 import javax.swing.JFrame;
5
6 public class ImageDisplay extends JFrame {
7     private Image image;
8     private ArrayList faces;
9     FaceLocations locations;
10    private String fileLoaded;
11
12    public ImageDisplay(String fileName, int patchSize) {
13
14        Toolkit toolkit = Toolkit.getDefaultToolkit();
15        String file = /*"trainPhotos/" +*/ fileName;
16        fileLoaded = fileName;
17        image = toolkit.getImage(file);
18        MediaTracker mediaTracker = new MediaTracker(this);
19        mediaTracker.addImage(image, 0);
20
21        try
22        {
23            mediaTracker.waitForID(0);
24        }
25        catch (InterruptedException ie)
26        {
27            System.err.println(ie);
28            System.exit(1);
29        }
30        addWindowListener(new WindowAdapter() { public void windowClosing(WindowEvent e)
31            {
32                System.exit(0);
33            }
34        });
35
36        setSize(image.getWidth(null), image.getHeight(null));
37        setTitle(fileName);
38        getFaceLocations(fileName);
39        setUpArray();
40        createFacePatchesFromArray(image, fileName, patchSize);
41        show();
42
43    }
44    public void getRectanglesAndDraw(Graphics graphicPass)
45    {
46        getFaceLocations();
47        setUpArray();
48        for(int i=0; i < faces.size(); i++)
49        {
50            FaceLocation face = (FaceLocation)faces.get(i);
51            int x = face.getX();
52            int y = face.getY();
53            int w = face.getWidth();
54            drawRectangle(graphicPass, x, y, w);
55        }
56    }
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
57 public void createFacePatchesFromArray(Image image, String fileName,int patchSize)
58 {
59     for(int i=0; i < faces.size();i++)
60     {
61         FaceLocation face = (FaceLocation)faces.get(i);
62         int x = face.getX();
63         int y = face.getY();
64         int w = face.getWidth();
65         FaceCrop crop = new FaceCrop(image,x,y,w);
66         FeatureLocalisation feature = new FeatureLocalisation(crop.getCroppedImage(),fileName,patchSize);
67     }
68 }
69 public void getFaceLocations()
70 {
71     System.out.println(getFileName());
72     locations = new FaceLocations(getFileName(),1,40,40);
73 }
74 public void getFaceLocations(String fileName2)
75 {
76     locations = new FaceLocations(fileName2,1,40,40);
77 }
78 public void setUpArray()
79 {
80     faces = locations.getFaceObjects();
81 }
82 public void setFileName(String name)
83 {
84     fileLoaded = name;
85 }
86 public String getFileName()
87 {
88     return fileLoaded;
89 }
90 public void paint(Graphics graphics) {
91     graphics.drawImage(image, 0, 0, null);
92     // Rectangle2D.Float rc = new Rectangle2D.Float(215F,203F,174F,174F);
93     //graphics.drawRect(215, 203, 174, 174);
94     getRectanglesAndDraw(graphics);
95     /*Date now2 = new Date();
96     now2.setTime(System.currentTimeMillis());
97     System.out.println(now2);*/
98 }
99 }
100 public void drawRectangles(Graphics graphics)
101 {
102     getRectanglesAndDraw(graphics);
103 }
104 public void drawRectangle(Graphics graphic, int x, int y, int z)
105 {
106     graphic.drawRect(x, y, z, z);
107 }
108
109 public static void main(String[] args)
110 {
111     /*Date now = new Date();
112     now.setTime(System.currentTimeMillis());
113     System.out.println(now);*/
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
114      String      fileTemp      =      "C:\\Documents      and      Settings\\Martyn\\My
115 Documents\\Eclipse\\FaceRecognitionProj\\VideoFrameBuffer\\0.jpg";
116      ImageDisplay display = new ImageDisplay(fileTemp,16);
117      display.setFileName(fileTemp);
118
119      //FaceDetection detect = new FaceDetection();
120
121      }
```

Figure 8.3-18 - ImageDisplay.java test class source code

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
1 import java.awt.*;
2 import java.awt.event.*;
3 import java.awt.geom.*;
4 import java.awt.image.*;
5 import java.io.*;
6 import java.net.*;
7 import java.util.ArrayList;
8
9 import javax.imageio.*;
10 import javax.swing.*;
11
12 public class ImageDisTest {
13     private FaceDataAccessLayer dal;
14     private Image image;
15
16     public ImageDisTest()
17     {
18         dal = new FaceDataAccessLayer();
19         ArrayList<Face> faces = new ArrayList<Face>();
20         dal.getFacesFromDatabase();
21
22         faces = dal.getFaceArray();
23         Face face = (Face)faces.get(0);
24         image = face.getFaceImage();
25     }
26     public static void main(String[] args) throws IOException {
27         ImageDisTest test = new ImageDisTest();
28         GraphicsConfiguration gc = getDefaultConfiguration();
29         Image thisImage = test.image;
30         BufferedImage buff = FeatureLocalisation.toBufferedImage(thisImage);
31         BufferedImage image = toCompatibleImage(buff, gc);
32         final double SCALE = 1.75;
33         int w = (int) (SCALE * image.getWidth());
34         int h = (int) (SCALE * image.getHeight());
35         final BufferedImage resize = getScaledInstance(image, w, h, gc);
36
37         JPanel p = new JPanel(new GridLayout(1,2));
38         p.add(new JLabel(new ImageIcon(image)));
39         p.add(new JLabel(new ImageIcon(resize)));
40
41         JPanel south = new JPanel();
42         JButton save = new JButton("save");
43         save.addActionListener(new ActionListener(){
44             public void actionPerformed(ActionEvent evt) {
45                 try {
46                     ImageIO.write(resize, "jpeg", new File("test.jpeg"));
47                 } catch (IOException e) {
48                     e.printStackTrace();
49                 }
50             }
51         });
52         south.add(save);
53
54         JFrame f = new JFrame("Example");
55         f.setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
56         Container cp = f.getContentPane();
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
57     cp.add(p, BorderLayout.CENTER);
58     cp.add(south, BorderLayout.SOUTH);
59     f.pack();
60     f.setLocationRelativeTo(null);
61     f.setVisible(true);
62 }
63
64 public static GraphicsConfiguration getDefaultConfiguration() {
65     GraphicsEnvironment ge = GraphicsEnvironment.getLocalGraphicsEnvironment();
66     GraphicsDevice gd = ge.getDefaultScreenDevice();
67     return gd.getDefaultConfiguration();
68 }
69
70 public static BufferedImage toCompatibleImage(BufferedImage image, GraphicsConfiguration gc) {
71     if (gc == null)
72         gc = getDefaultConfiguration();
73     int w = image.getWidth();
74     int h = image.getHeight();
75     int transparency = image.getColorModel().getTransparency();
76     BufferedImage result = gc.createCompatibleImage(w, h, transparency);
77     Graphics2D g2 = result.createGraphics();
78     g2.drawImage(image, null);
79     g2.dispose();
80     return result;
81 }
82
83 //returns target
84 public static BufferedImage copy(BufferedImage source, BufferedImage target) {
85     Graphics2D g2 = target.createGraphics();
86     g2.setRenderingHint(RenderingHints.KEY_INTERPOLATION, RenderingHints.VALUE_INTERPOLATION_BICUBIC);
87     double scalex = (double) target.getWidth() / source.getWidth();
88     double scaley = (double) target.getHeight() / source.getHeight();
89     AffineTransform xform = AffineTransform.getScaleInstance(scalex, scaley);
90     g2.drawImage(source, xform);
91     g2.dispose();
92     return target;
93 }
94
95 public static BufferedImage getScaledInstance(BufferedImage image, int width, int height, GraphicsConfiguration gc) {
96     if (gc == null)
97         gc = getDefaultConfiguration();
98     int transparency = image.getColorModel().getTransparency();
99     return copy(image, gc.createCompatibleImage(width, height, transparency));
100 }
101 }
```

Figure 8.3-19 - ImageDisTest.java source code

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
1  import java.awt.FlowLayout;
2  import java.awt.Graphics;
3  import java.awt.GridLayout;
4  import java.awt.MediaTracker;
5  import java.awt.Toolkit;
6  import java.awt.image.BufferedImage;
7  import java.util.ArrayList;
8  import java.io.*;
9
10 import javax.imageio.ImageIO;
11 import javax.swing.JButton;
12 import javax.swing.JLabel;
13 import javax.swing.JPanel;
14 import javax.swing.JFrame;
15 import javax.swing.ImageIcon;
16 import java.awt.Image;
17
18 import java.awt.event.WindowAdapter;
19 import java.awt.event.WindowEvent;
20 import java.awt.geom.*;
21 import java.util.TimerTask;
22 import java.awt.*;
23
24 public class DBTestClass extends JFrame
25 {
26
27     private FaceDataAccessLayer dal;
28     private Image image;
29
30     public DBTestClass()
31     {
32         dal = new FaceDataAccessLayer();
33         ArrayList<Face> faces = new ArrayList<Face>();
34         dal.getFacesFromDatabase();
35
36         faces = dal.getFaceArray();
37         Face face = (Face)faces.get(0);
38         image = face.getFaceImage();
39         Toolkit toolkit = Toolkit.getDefaultToolkit();
40         //toolkit.cre
41         File fOut = new File("file");
42         BufferedImage buff = FeatureLocalisation.toBufferedImage(image);
43         try
44         {
45             ImageIO.write(buff, "jpg", fOut);
46         }
47         catch(Exception e){}
48
49         //image = toolkit.getImage(file);
50         MediaTracker mediaTracker = new MediaTracker(this);
51         mediaTracker.addImage(image, 0);
52
53         try
54         {
55             mediaTracker.waitForID(0);
56         }
57         catch (InterruptedException ie)
```

FACIAL RECOGNITION FROM VIDEO SEQUENCES

```
59         {
60             System.err.println(ie);
61             System.exit(1);
62         }
63         addWindowListener(new WindowAdapter() { public void windowClosing(WindowEvent e)
64             {
65                 System.exit(0);
66             }
67         });
68
69         setSize(image.getWidth(null), image.getHeight(null));
70         setTitle("Test");
71         show();
72
73     }
74     public void paint(Graphics graphics) {
75         graphics.drawImage(image, 0, 0, null);}
76     public static void main(String[] args)
77     {
78         System.out.println("Going");
79         DBTestClass test = new DBTestClass();
80     }
81     public static GraphicsConfiguration getDefaultConfiguration() {
82         GraphicsEnvironment ge = GraphicsEnvironment.getLocalGraphicsEnvironment();
83         GraphicsDevice gd = ge.getDefaultScreenDevice();
84         return gd.getDefaultConfiguration();
85     }
86
87     public static BufferedImage toCompatibleImage(BufferedImage image, GraphicsConfiguration gc) {
88         if (gc == null)
89             gc = getDefaultConfiguration();
90         int w = image.getWidth();
91         int h = image.getHeight();
92         int transparency = image.getColorModel().getTransparency();
93         BufferedImage result = gc.createCompatibleImage(w, h, transparency);
94         Graphics2D g2 = result.createGraphics();
95         g2.drawImage(image, null);
96         g2.dispose();
97         return result;
98     }
99 }
100 }
```

Figure 8.3-20 - DBTestClass.java test class source code

FACIAL RECOGNITION FROM VIDEO SEQUENCES

8.4 Appendix D – System Testing

8.4.1 System Testing Results

<i>Test No.</i>	<i>Description</i>	<i>Successful?</i>	<i>Reference</i>	<i>Additional</i>
	Interfaces correctly displayed			
1	Main interface is correctly displayed	Yes	Figure 8.2-1	
2	Add reference face interface is correctly displayed	Yes	Figure 8.2-2	
3	Remove reference face interface is correctly displayed	Yes	Figure 8.2-3	
	Testing Interface			
4	Menu items launch relevant sub-interfaces	Yes	Figure 8.2-2	
5	Main interface begins with no matching face displayed	Yes	Figure 8.2-3	
6	Test video is displayed constantly	Yes	Figure 8.2-1	
7	Connection to database failed message shown if connection fails	Yes	Figure 8.4-1	
8	Console shows feature information when adding reference face then clears screen	Yes	N/A	
9	Message is displayed if reference image is missing	Yes	Figure 8.4-2	The reference image is stored in a file store and the path is in the database. This checks for manual tampering
10	Message is shown if .dat file (Covariance) from Matlab are missing	Yes	Figure 8.4-4	

FACIAL RECOGNITION FROM VIDEO SEQUENCES

<i>Test No.</i>	<i>Description</i>	<i>Successful?</i>	<i>Reference</i>	<i>Additional</i>
11	Message is shown if .dat file (Mean) from Matlab are missing	Yes	Figure 8.4-3	
12	Message s shown if Video capture device is unavailable	Yes	Figure 8.4-5	
13	Reference face is added to filestore	Yes	N/A	Added on disk
14	Matched person details appear on screen and are updated periodically	Yes	Figure 8.4-6	

Table 8.4-1 - System tests



Figure 8.4-1 - Database unavailable window




Figure 8.4-2 - Reference Image missing warning




Figure 8.4-3 - Mean .dat file missing error message



Figure 8.4-4 - Covariance .dat file missing error message

FACIAL RECOGNITION FROM VIDEO SEQUENCES




Figure 8.4-5 - Video source unavailable error message



Figure 8.4-6 - Final running main interface




Figure 8.4-7 - Results Graph

8.4.2 System Requirements Testing Results

Number	Successful?	Comments
1	Yes	The video feed works correctly
2	Yes	Reference image are cropped about the face but not normalised anymore than that (they are, however, normalised as part of the feature extraction process.)
3	Yes	Bayesian feature localisation utilised
4	Yes	Weak Classifier face detector used
5	Yes	Provided by specialist interface
6	Yes	Image is updated periodically with closest match
7	No	No validation is required because the data is string-like (for the name) and the face is validated automatically by the face detector
8	Yes	Error messages are used extensively and are informative

Table 8.4-2 - System requirements results ref. (Table 4.2-1)

Number	Successful?	Comments
1	Yes	Figure 8.2-1, although this is subjective to user preferences
2	Yes	Figure 8.2-2
3	Yes	Although, real-time here is 4 frames per

FACIAL RECOGNITION FROM VIDEO SEQUENCES

Number	Successful?	Comments
		second
4	Yes	Figure 8.2-1
5	Yes	Figure 8.4-1 through Figure 8.4-5

Table 8.4-3 - User requirements results ref. (Table 4.2-2)

Number	Successful?	Comments
1	Yes	Java, Figure 8.3-17
2	Yes	Figure 8.4-1 through Figure 8.4-5
3	Yes	Figure 8.2-2 and Figure 8.2-3
4	Yes (theoretically)	Although, all reference images were in .jpeg format
5	Yes	Figure 8.2-2
6	Yes	Figure 8.2-3

Table 8.4-4 - System specification results ref. (Table 4.2-3)

Number	Successful?	Comments
1	Yes	Although, real-time here is 4 frames per second
2	Yes	Figure 8.4-7
3	Yes	N/A

Table 8.4-5 - Face Detector Specification results ref. (Table 4.3-1)

8.4.3 Evaluation Framework

The following evaluation framework diagram (Figure 8.4-8) is a useful diagram used to highlight certain issues or successes with an application framework. The right hand column classifies the results as; Advanced (A), Intermediate (I) or Basic (B). These refer to the system implementation as compared to the original design specification. If no 'tick' has been placed in the 'category' section, it is not relevant in this scenario.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

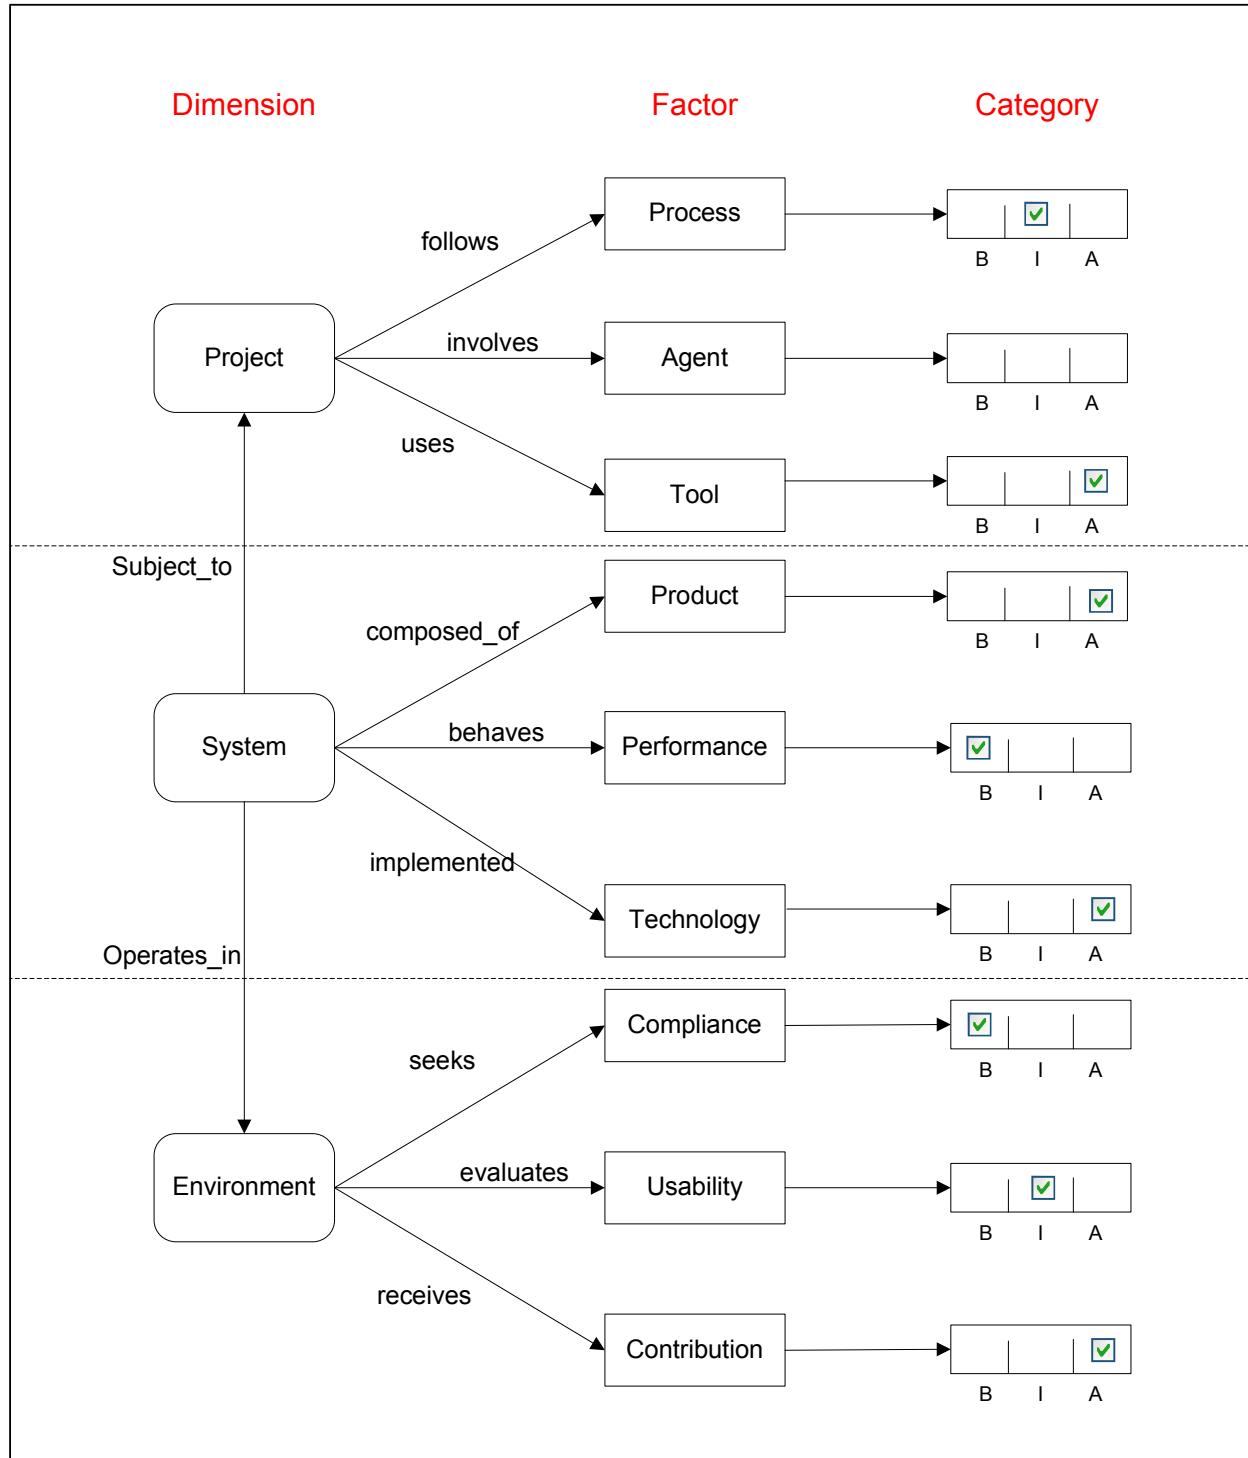


Figure 8.4-8 - Evaluation Framework Diagram

8.5 Appendix E – Personal Reflection

When I began this project, I wanted to pick something that satisfied a number of criteria;

- For the project to be suitably complex in order to challenge me to a suitable extent
- For the project to teach me about complex programming paradigms for use in industry

FACIAL RECOGNITION FROM VIDEO SEQUENCES

- For the project to be loosely linked to the area of computer security and biometrics because this is an area of particular interest to me.

I believe the project realistically started in the worst possible way. I had an uneven module split; weighted toward the first semester and therefore left the project mainly to be done in the second semester. This was a mistake because travelling to and from London in order to conduct job interviews took more time than anticipated. Finally, coming into Easter, I was significantly behind schedule and the project was starting to look in trouble.

However, having worked all hours in the day for a couple of months, I managed to produce something that I was reasonably happy with. The program was becoming extremely complex but was producing excellent results and was nearing completion. The final push of Easter (and the post-Easter coursework lull) gave me time to write up the project properly and finalise the code by ‘gluing’ the individual components together.

I believe that it would have been immensely beneficial for me to conduct much more background research earlier. This means not only looking aimlessly at the mounds of research available, but spending time with the library staff in order to gain a better understanding of how to find the information that I needed. I must have wasted days aimlessly searching for relevant information.

I then found Google scholar and from there, located the IEEE website (to which access credentials are provided by the University library). This website contains hundreds of papers from around the world in the areas of computer vision and image processing. This is truly a ‘golden’ resource.

I would also strongly suggest that people delay job interviews until after the dissertations are in. This really was a significant strain on my time and I would have been much more comfortable in completing the project and maybe even adding additional feature detectors to the software itself.

However, I believe that the programming went very well and am extremely pleased with the results of that. Given more time, as mentioned above, I would have liked to extend the project to use more feature detectors and therefore make it a more reliable systems. Time constraints, which should really have been managed better, made this impossible.

If I were to redo this project, I would begin by looking through a large number of previous projects on this topic because that would have given me a better idea of exactly what was expected of me. This would save me time later on, during the report writing phase, on time spent editing and adding missed sections.

To summarise, planning and time keeping is paramount and I would have liked to do this better. Relying on a supervisor’s knowledge and using the time in meetings to maximum effect was extremely beneficial to my project and I wouldn’t have been able to complete it without that help. Finally, spend some time with library staff to understand how to use the search system. It’s more complex than it appears to be.

As a side note; plan in advance to have somebody read over the report well before submission. Unfortunately, I failed to do this and couldn’t find anybody that would do it at short notice. Therefore, I had to rely on my own, extremely tired, eyes.

8.6 Bibliography

- 1) Aarsten, A., Brugali, D., & Menga, G. (1996). *Patterns for Three-Tier Client/Server Applications*. Torino, Italy: Dept. of Automatica e Informatica.
- 2) Belhumeur, P. N., Hespanha, J. P., & Kriegman, D. J. (1997). Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 19, pp. 711-720. IEEE Computer Society.
- 3) Bing-Bing, C., Vass, J., & Zhuang, X. (1999). Significance-Linked Connected Component Analysis for Wavelet Image Coding. *IEEE TRANSACTIONS ON IMAGE PROCESSING*. 8, pp. 774-785. Princeton: IEEE Computer Society.
- 4) Boloix, G., & Robillard, P. N. (1995). *A Software System Evaluation Framework*. Montreal, Canada: IEEE Computer Society.
- 5) Bornet, O. (2005, May 19). *Learning Based Computer Vision with Intel's Open Source Computer Vision Library*. Retrieved April 2007, 2007, from Intel.com Web site: http://www.intel.com/technology/itj/2005/volume09issue02/art03_learning_vision/p04_face_detection.htm
- 6) Brunelli, R., & Poggio, T. (1993). Face Recognition: Features versus templates. *IEEE Transaction on Pattern Analysis and Machine Intelligence* , 15 (10), 1042-1052.
- 7) Choi, J., Lee, S., Lee, C., & Yi, J. (2001). A Real-time Face Recognition System using Multiple Mean Faces and Dual Mode Fisherfaces. *IEEE Transactions on Pattern Analysis and Machine Intelligence* (pp. 1686-1689). Suwon, Korea: IEEE Computer Society.
- 8) Dr. Hu Software. (2002). *JNI Face Detection (OpenCV wrapper in Java)*. Retrieved February 21, 2007, from Drhu.org: <http://www.drhu.org/freeDownload.asp>
- 9) Dyer, C. R. (2004). *Chuck Dyer on Facial Recognition*. Retrieved April 6, 2007, from The University of Wisconsin: <http://www.cs.wisc.edu/~gdguo/faceExpression.htm>
- 10) Efford, N. (2006). *SY32 Lecture 14*. Retrieved April 07, 2007, from [www.comp.leeds.ac.uk/sy32:](http://www.comp.leeds.ac.uk/sy32/) <http://www.comp.leeds.ac.uk/sy32/lectures/lect14b.pdf>
- 11) Everingham, M., & Zisserman, A. (2006). Regression and Classification Approaches to Eye Localization in Face Images. *7th International Conference on Automatic Face and Gesture Recognition* , 441-448.
- 12) Heisele, B. (2006). A Component-based Framework for Face Detection and Identification. *International Journal of Computer Vision* .
- 13) Heusch, G., Cardinaux, F., & Marcel. (2003). Lighting Normalization Algorithms for Face Verification. *IDIAP Research Institute* .

FACIAL RECOGNITION FROM VIDEO SEQUENCES

- 14) Hjelmås, E., & Wroldsen, J. (1999). Recognizing Faces from Eyes Only. *In Proceedings of the 11th Scandinavian Conference on Image Analysis*. Oslo.
- 15) Horn, C., Gatune, J., Thomas, A., & Woodward, J. (. (2003). *Biometrics; A Look at Face Recognition*. Virginia State Crime Commission: RAND Public Safety and Justice.
- 16) Jiao, F., Gao, W., Chen, X., Cui, G., & Shan, S. (2002). A Face Recognition Method Based on Local Feature Analysis. *The 5th Asian Conference on Computer Vision*. Melbourne, Australia.
- 17) Kawaguchi, T., Rizon, M., & Hidaka, D. (2005). Detection of Eyes from Human Faces by Hough Transform and Separability Filter. *Electronics and Communications in Japan, Part 2* , 88 (5), 2190–2200.
- 18) Kirby, M., & Sirovich, L. (1990). Application of the Karhunen-Loève Procedure for the Characterization of Human Faces. *IEEE Transactions on Pattern Analysis and Machine Intelligence* (pp. 103-108). IEEE.
- 19) Kittler, J., & Nixon, M. S. (2003). Expression-Invariant 3D Face Recognition. *Audio & Video-based Biometric Person Authentication, LNCS 2688* , 62-70.
- 20) Low, A. (1991). *Introductory Computer Vision and Image Processing*. McGraw-Hill Book Company.
- 21) Meek, J. (2002, June 13). *Guardian.co.uk | Newham Facial Recognition*. Retrieved March 22, 2007, from Guardian.co.uk: <http://www.guardian.co.uk/Archive/Article/0,4273,4432506,00.html>
- 22) Moon, H., Chellappa, R., & Rosenfeld, A. (2002). Optimal Edge-Based Shape Detection. *IEEE Transactions on Image Processing*. 11, pp. 1209-1226. Maryland: IEEE Computer Society.
- 23) Osuna, E., Freund, R., & Girosit, F. (1997). Training Support Vector Machines: an Application to Face Detection. *IEEE Transactions for Pattern Recognition and Machine Learning* (pp. 130-136). Cambridge, MA: IEEE Computer Society.
- 24) Pezdek, K., Blandon-Gitlin, I., & Moore, C. (2003). Children's Face Recognition Memory: More Evidence for the Cross-race effect. *Journal of Applied Psychology* , 88 (4), 760-763.
- 25) Roth, D., Yang, M.-H., & Ahuja, N. (2000). *A SNoW-Based Face Detector*. Urbana, IL: Department of Computer Science and the Beckham Institute.
- 26) Rowley, H. A., Baluja, S., & Kanade, T. (1998). Neural Network-Based Face Detection. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 20. IEEE.
- 27) Rowley, H. A., Baluja, S., & Kanade, T. (1998). Rotation Invariant Neural Network-based Face Detection. *IEEE Computer Society Conference on Computer Vision and Pattern Recognition* (pp. 38-44). IEEE.

FACIAL RECOGNITION FROM VIDEO SEQUENCES

- 28) Samal, A., & Iyengar, P. A. (1992). Automatic recognition and analysis of human faces and facial expressions: a survey. *Pattern Recognition*, 25 (1), 65-77.
- 29) Smith, I. L. (2002). *A tutorial on Principal Components Analysis*. Otago, NZ.
- 30) Statsoft. (2004). *Support Vector Machines (SVM)*. Retrieved April 12, 2007, from Statsoft.com: <http://www.statsoft.com/textbook/stsvm.html>
- 31) Torres, L. (2004). *Is there any hope for face recognition?* Barcelona, Spain: Technical University of Catalonia.
- 32) Turk, A. M., & Pentland, A. P. (1991). Face Recognition using Eigenfaces. *Computer Vision and Pattern Recognition* (pp. 586-591). Minnesota: IEEE Computer Society.
- 33) University College London; Oncology. (2006). *Principle Components Analysis (PCA)*. Retrieved April 11, 2007, from UCL.ac.uk: http://www.ucl.ac.uk/oncology/MicroCore/HTML_resource/PCA_1.htm
- 34) Vapnik, V. (1995). *The Nature of Statistical Learning Theory*. New York: Springer.
- 35) Vapnik, V., & Cortes, C. (1995). Support-Vector Networks. In L. Saitta (Ed.), *Machine Learning*. 20, pp. 273-297. Boston, MA: Kluwer Academic Publishers.
- 36) Vernon, D. (1991). *Machine Vision. Automated Visual Inspection and Robot Vision*. Prentice Hall.
- 37) Viola, P., & Jones, M. J. (2004). Robust Real-Time Face Detection. *International Journal of Computer Vision* 57(2) (pp. 137-154). Netherlands: Kluwer Academic Publishers.
- 38) Yang, J., Zhang, D., Frangi, A. F., & Yang, J.-y. (2004). Two-dimensional PCA: A New Approach to Appearance-Based Face Representation and Recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 26, pp. 131-137. Hong Kong: IEEE Computer Society.
- 39) Yang, M.-H. (2002). Kernel Eigenfaces vs. Kernel Fisherfaces: Face Recognition Using Kernel Methods. *Fifth IEEE International Conference on Automatic Face and Gesture Recognition*. Mountain View, CA: IEEE Computer Society.
- 40) Zelek, J. S., & Ersi, F. E. (2006). Local Feature Matching For Face Recognition. *Proceedings of the 3rd Canadian Conference on Computer and Robot Vision* (pp. 4 - 4). Waterloo: University of Waterloo.
- 41) Zhao, S., & Grigat, R.-R. (2006). *An Automatic Face Recognition System in the near Infrared Spectrum*. Hamburg: Technical University Hamburg Harburg, Vision Systems.
- 42) Zhao, W., Chellappa, R., Rosenfeld, A., & Phillips, P. J. (2007). *Face Recognition: A Literature Survey*. Maryland: National Institute of standards and Technology.

